

Working with Static Sites

BRINGING THE POWER OF SIMPLE
TO MODERN WEBSITES



Early Release

RAW & UNEDITED

Raymond Camden
& Brian Rinaldi

Working with Static Sites

Bringing the Power of Simple to Modern Websites

Raymond Camden and Brian Rinaldi

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

Working with Static Sites

by Raymond Camden and Brian Rinaldi

Copyright © 2016 Raymond Camden and Brian Rinaldi. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc. , 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com .

Editor: Allyson MacDonald

Production Editor: FILL IN PRODUCTION EDITOR
TOR

Copyeditor: FILL IN COPYEDITOR

Proofreader: FILL IN PROOFREADER

Indexer: FILL IN INDEXER

Interior Designer: David Futato

Cover Designer: Karen Montgomery

Illustrator: Rebecca Demarest

January -4712: First Edition

Revision History for the First Edition

2016-11-17: First Early Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781491960875> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. Working with Static Sites, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

While the publisher and the author(s) have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author(s) disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-491-96087-5

[FILL IN]

Table of Contents

Preface.....	vii
1. Why Static Sites?.....	11
Benefits of Static Sites	12
Static Sites are Fast	12
Static Sites are Secure	13
Other Benefits	13
What Kinds of Sites can Go Static?	13
What are Static Site Generators?	14
What You Will Learn	16
2. Building a Basic Static Site.....	19
Welcome to Harp	20
Your First Harp Project	23
Working with Layouts and Partials	26
Working with Data	29
Generating a Site	35
Building Camden Grounds	35
Going Further with Harp	46
3. Building a Blog.....	47
Blogging with Jekyll	47
Your First Jekyll Project	49
Writing a Post	53
A quick introduction to Liquid	56
Working with Layouts and Includes	58
Adding Additional Files	60
Working with Data	62

Configuring your Jekyll Site	66
Generating a Site	67
Building a Blog	68
Going Further with Jekyll	74
4. Building a Documentation Site.....	77
Characteristics of a Documentation Site	77
Choosing a Generator for Your Documentation Site	78
Our Sample Documentation Site	79
Creating the Site	82
Installing Hugo	82
Generating the Initial Site Files	83
Configuring the Hugo Site	84
Adding Content	85
Creating the Layout	87
Going Further	93
5. Adding Dynamic Elements.....	95
Handling Forms	95
WuFoo Forms	96
Google Docs Forms	103
Formspree	109
Adding a Comment Form to Camden Grounds	111
Adding Comments	113
Working with Disqus	114
Adding Comments to the Cat Blog	117
Adding Search	118
Creating a Custom Search Engine	118
Adding a Custom Search Engine to a Real Site	122
Even More Options	123
6. Adding a CMS.....	125
CloudCannon	126
Creating a Site on CloudCannon	126
Editing a Site on CloudCannon	129
Where to Go From Here	134
Netlify CMS	134
Setting up the Netlify CMS	134
Where to Go From Here	142
Jekyll Admin	142
Setting Up Jekyll Admin	143
Editing a Site in Jekyll Admin	143

Where To Go From Here	147
More Options	147
Forestry.io	148
Lektor	149
Headless CMS	151
7. Deployment.....	153
Plain Ole' Web Servers	153
Cloud File Storage Providers	154
Hosting a Site on Amazon S3	154
Hosting a Site on Google Cloud Storage	159
Deploying with Surge	163
Deploying with Netlify	170
Summary	182
8. Migrating to a Static Site.....	183
Migrating from Wordpress to Jekyll	183
Other Migration Options	188
Hugo	189
Middleman	190
Hexo	191
Harp	191
Many More...	191
A. Appendix Title.....	193
Index.....	195

Preface

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This element signifies a tip or suggestion.



This element signifies a general note.



This element indicates a warning or caution.

Using Code Examples

Supplemental material (code examples, exercises, etc.) is available for download at https://github.com/oreillymedia/title_title.

This book is here to help you get your job done. In general, if example code is offered with this book, you may use it in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: “*Book Title* by Some Author (O'Reilly). Copyright 2012 Some Copyright Holder, 978-0-596-xxxx-x.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

Safari® Books Online



Safari Books Online is an on-demand digital library that delivers expert **content** in both book and video form from the world's leading authors in technology and business.

Technology professionals, software developers, web designers, and business and creative professionals use Safari Books Online as their primary resource for research, problem solving, learning, and certification training.

Safari Books Online offers a range of **plans and pricing** for **enterprise, government, education**, and individuals.

Members have access to thousands of books, training videos, and prepublication manuscripts in one fully searchable database from publishers like O'Reilly Media, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kauf-

mann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, Course Technology, and hundreds **more**. For more information about Safari Books Online, please visit us [online](#).

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at <http://www.oreilly.com/catalog/<catalog page>>.

To comment or ask technical questions about this book, send email to bookquestions@oreilly.com.

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

Acknowledgments

CHAPTER 1

Why Static Sites?

In the beginning, the web was only static sites. In fact, the **first web site** was (and still is) technically a static site.

World Wide Web

The WorldWideWeb (W3) is a wide-area [hypermedia](#) information retrieval initiative aiming to give universal access to a large universe of documents.

Everything there is online about W3 is linked directly or indirectly to this document, including an [executive summary](#) of the project, [Mailing lists](#), [Policy](#), November's [W3 news](#), [Frequently Asked Questions](#).

What's out there?

Pointers to the world's online information, [subjects](#), [W3 servers](#), etc.

Help

on the browser you are using

Software Products

A list of W3 project components and their current state. (e.g. [Line Mode](#), X11 [Viola](#), [NeXTStep](#), [Servers](#), [Tools](#), [Mail robot](#), [Library](#))

Technical

Details of protocols, formats, program internals etc

Bibliography

Paper documentation on W3 and references.

People

A list of some people involved in the project.

History

A summary of the history of the project.

How can I help?

If you would like to support the web..

Getting code

Getting the code by [anonymous FTP](#), etc.

Figure 1-1. The first web site from CERN

[first_web_site.png](#)

Of course, no one called them “static sites” back then as the entire web consisted of static HTML documents - there was no non-static alternative.

We've clearly come a long way since then, both in terms of the underlying technologies that build the web and web sites as well as how we expect a web site to look and behave. So, why would static sites be a worthwhile option for today's web?

First, let's explore some of the benefits of static sites before we dive into how the changing technology behind static sites (i.e. static site generators - the topic of this book!) are making them viable again.



A Deeper Background on Static Sites

If you are looking for a far more comprehensive look at what static sites are and how they fit into the larger development ecosystem, I previously wrote a short report on Static Site generators for O'Reilly. The report covers a good deal of detail behind the history of static sites, how they differ from dynamic sites built with content management systems or blog engines and some details about the available static site generator options. Rather than repeat the information entirely here, you can [download it for free from O'Reilly](#). Trust me, it's a quick and easy read!

Benefits of Static Sites

Of the many reasons that static sites are coming back into fashion, two stand out:

- Static sites are fast;
- Static sites are secure.

Static Sites are Fast

All developers seem to understand that web site performance is critical. For example, [recent studies](#) have shown that users tend to abandon sites that take longer than 3 seconds to load (with an under 2 second load time being considered optimal for mobile). Achieving that level of web site performance, however, can be difficult.

By their very nature, static sites load extremely fast. This is because every visitor is served the exact same HTML without the bottlenecks caused by a server-side language, database or any kind of dynamic rendering. Plus static files are extremely easy to cache and even serve via a CDN (content delivery network), making them even faster for the end user. Plus, once you eliminate dynamic rendering from a database, you've eliminated numerous points of failure that often cause sites to be unresponsive or completely fail.

Static Sites are Secure

Sadly, it is not uncommon nowadays for us to hear about a site being the target of a SQL injection or cross-site scripting (XSS) attack, two of the most common types of website security breaches. Oftentimes, hackers have gained access to a site via a vulnerability in the code, many times due to an unpatched CMS. However, with a static site, there is no database to breach and no server-side platform or CMS with unpatched vulnerabilities.

Speaking from personal experience, even a tightly patched and locked down CMS can be vulnerable. And finding then repairing the damage done by a breach can be extremely time consuming and difficult.

Obviously, static sites will not eliminate every vulnerability (what will?), but they do not only narrow the window of opportunities available to any hacker, but can also limit the amount of potential damage they can do should they somehow gain access.

Other Benefits

While those are the two key benefits of static sites, there are certainly others. Some of the ones I commonly cite when speaking on the topic are:

- *Flexibility* — There is no CMS framework to work within and, thus, no limitations on how you can build your static site.
- *Hosting* — Because there's no need for a database or server-side language support, hosting a static site can be anywhere from inexpensive to completely free, depending on your specific needs.
- *Versioning* — Since a static site is made up of static files, it is extremely easy to version everything using version control like Git (and GitHub).

With all these benefits, why would you not choose to use a static site? Well, in truth, only certain kinds of sites can realistically work as static only.

What Kinds of Sites can Go Static?

There are drawbacks to using static sites. For instance, while some amount of dynamic data is possible on a static site using external API calls or third-party services, a static site is simply not suitable for sites that require a large amount of dynamic data or content personalization. Also, from a development and content contribution standpoint, static site generators (i.e. the tools frequently used to build static sites — and what this book is about) can have a steep learning curve. Lastly, deployment (which we'll talk about in Chapter 7) can get complex, making static sites less than ideal for content that changes very frequently.

Keeping those things in mind, the sites that tend to work best as static sites are those that are content-focused, update infrequently (once or twice a day, at most, I'd say)

and do not require a high degree of user interaction or personalization. Here are some examples of the types of sites that work well as static sites:

- *Blogs* — This is the most common use-case (in fact, many static site generators default to a blog template). Blogs are content-focused by design and, in many cases, user interaction is limited to comments, where services like [Disqus](#) can fill the requirement.
- *Documentation* — In my own experience, this is the second most common use for static sites as documentation is purely about a fixed set of content that tends to update infrequently, but which the user really expects to get quickly (often on the go). Static sites fit these needs perfectly while providing the potential added benefit that they are easy to host on services like GitHub, for both versioning and community contribution, if it is desired.
- *Informational Sites and Brochureware* — Much of the web is actually made up of fairly simple web sites. I might create a site for an event I am running, a web “brochure” for by small business, an informational site for my community, etc. For any of these sorts of sites, a CMS is overkill but we want updating to be quick and painless when necessary. A static site (using a static site generator) can fit this need perfectly.

What are Static Site Generators?

So up to this point, we've mostly been talking about how static sites behave, not how they are built. But, if we were still in the days of building static sites using Dreamweaver or — heaven forbid — FrontPage (remember that?), the benefits of having a static site would be outweighed by the pain of building and maintaining one.

Static site generators solve the pain of building and maintaining a static site. The fundamentals of a static site generator are extremely simple: they take in dynamic content and layout and output static HTML, CSS and JavaScript files. There are literally [hundreds of static site generators](#), but essentially they all do exactly the same thing and, for the most part, function similarly.

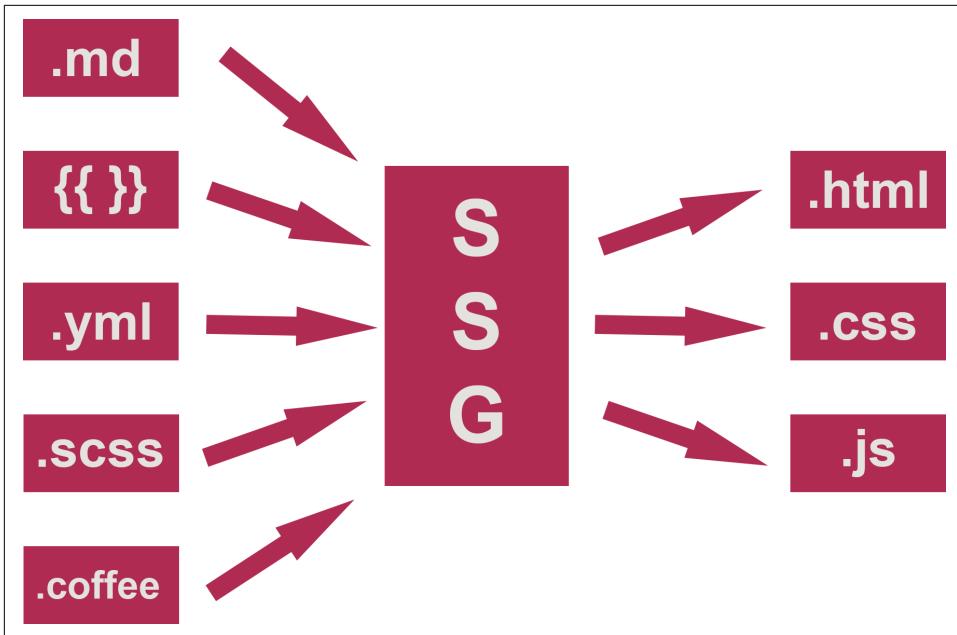


Figure 1-2. A simple illustration of what a static site generator does. It takes a lightweight markup language (ex. Markdown), a templating language (ex. Handlebars) and data/metadata (ex. Yaml) as well as sometimes a CSS preprocessor (ex. Sass) and a compile-to-JavaScript language (ex. CoffeeScript) and generates HTML, CSS and JavaScript.

[ssg.png](#)

Most static site generators...

- Use one or more templating languages (e.g. Liquid, Handlebars, Jade). This is a key part of any static site generator as it allows you to build a layout/theme for your site and plug in dynamic content during the build.
- Use one or more **lightweight markup languages** (e.g. Markdown, AsciiDoc, reStructuredText). While every static site generator that I've ever tried (and I've tried too many) supports straight HTML, using a lightweight markup language makes it quick and easy to write content using any text editor (provided you've learned the syntax of course).
- Are run via the command line (e.g. Terminal/Command Prompt). While it is becoming somewhat more common for static site generators to include some sort of GUI, nearly every static site generator is primarily designed as some sort of command line tool.
- Include a local development server. This allows you to quickly develop and test locally before building and deploying your changes. Typically changes are automatically watched and rebuilt during local development.

- Are extensible. In most cases, if your static site generator doesn't support a feature or language that you wish to use, there is a built-in plugin architecture that allows you to add that in (provided you can code in the language the tool is built upon, of course).
- Support file-based data formats (e.g. YAML, TOML, JSON). Lightweight markup languages and HTML are used for long form content, while file-based data allows you to structure any sort of arbitrary data independent of its display.

What You Will Learn

So, now that we've got a basic understanding of what a static site generator is? How can we use them?

The first issue to resolve is trying to figure out which option to use? This isn't an easy decision since, as of this writing, there are currently **445 different available options**. Even after filtering out projects that haven't been recently updated, we're still left with hundreds of potential tools.

Static Site Generators						
The definitive listing of Static Site Generators — all 445 of them!						
Stars	Name	License	Language	Created	Updated	
89	ABlog	MIT	Python	2 years ago	12 days ago	
35	Ace	MIT	Ruby	6 years ago	a month ago	
306	acrylamid	BSD-2-Clause	Python	6 years ago	5 days ago	
19	AkashaCMS	JavaScript	4 years ago	3 months ago		
34	Akashic	Common Lisp	4 years ago	16 days ago		
245	antwar	MIT	JavaScript	2 years ago	3 days ago	
	Appernetic	Web	9 months ago			

Figure 1-3. The [StaticSiteGenerators.net](#) site is a definitive list of nearly every static site generator project.

So, how do you choose?

Over the next few chapters we'll look at some of the more mature and popular options for developing static sites including Jekyll, Hugo and Harp. Not only does each have a different underlying language (Ruby for Jekyll, Go for Hugo and Java-Script for Harp), which may be an important consideration, but each also has its own pros and cons. We'll look at building some of the common use cases discussed above (a basic informational site, a blog and a documentation site) using these tools in ways that take advantage of their relative merits.

However, once we've built the basics of our static site, we need to add in some dynamic features, like comments on our blog posts or a site search. Or, if we have to support content contributors that aren't comfortable writing posts in Markdown via a text editor, we might want to add a CMS-like back end to our site to allow for easy editing. We'll take a look at multiple solutions that solve each of these problems.

After your static site is complete, it's time to deploy it. While this can be as easy as simply FTPing files onto a server, in most cases you'll want to automate the process or take advantage of services that can manage the build process for you. So, we'll explore at a variety of tools and services that can ease the deployment process.

Finally, you may be evaluating static sites as an option to replace an existing site that already uses a tool like Wordpress or some other CMS. For these cases we'll dive into tools that are available to ease the process of migrating to a static site generator from a CMS by bringing over your existing content.

In the end, we hope to provide you with both the broad overview of the existing static site generator ecosystem, while diving into the actual implementation details of how to accomplish your goals with these tools.

CHAPTER 2

Building a Basic Static Site

For our first static site, we'll start with something incredibly simple. For some, "brochure-ware" is a derogatory term referring to a website that looks like it was copied directly from a marketing brochure. While some sites certainly *are* simple copies of marketing material, that doesn't mean there's anything particularly *wrong* with them either.

There are times where a website will be nothing more than one or two pages of content. For example, maybe you're launching the Next Big Thing(tm) as part of your plan to become a dot-com millionaire. While your product is in development, you may simply need a "Coming Soon" site with one page.

Other examples of simple sites include:

- A restaurant site providing hours, address, and a menu
- A "landing" page for your mobile app that provides links to the various app stores
- A portfolio page acting as a resume for someone looking for a job

In all of these examples, the entire "site" may consist of nothing more than one or two unique pages. A static site generator may be overkill for such a thing but in my experience, small projects have a way of growing, not the other way around. All of the benefits that static site generators provide us will simply become *more* useful over time as the site (possibly) grows with new content.

For our first example site, we're going to create an online presence for a coffee shop called Camden Grounds. While not a terribly imaginative name, you've probably run across simple coffee shop/restaurant sites before, so you've got a basic idea of how they work. For Camden Grounds, the site will consist of:

- A home page which is mostly pretty images.

- A menu showing a list of various coffees, teas, cookies, and more.
- A list of locations because - wouldn't you know it - Camden Grounds is actually a chain.
- An “About Us” page talking about the history of the company. No one is ever going to read this, but it's pretty standard for such sites.

That's a grand total of four pages, and, as I said above, a static site generator may be a bit much for this, but the good news is that we will be “future-proofed” when we expand the site in the future. For such a simple site, we'll look at our first static site generator in the book, Harp.

Welcome to Harp

Harp (harpjs.com) is a light-weight static site generator. It is rather simple, and at times limited, but can be easy to pick up for people new to static sites.

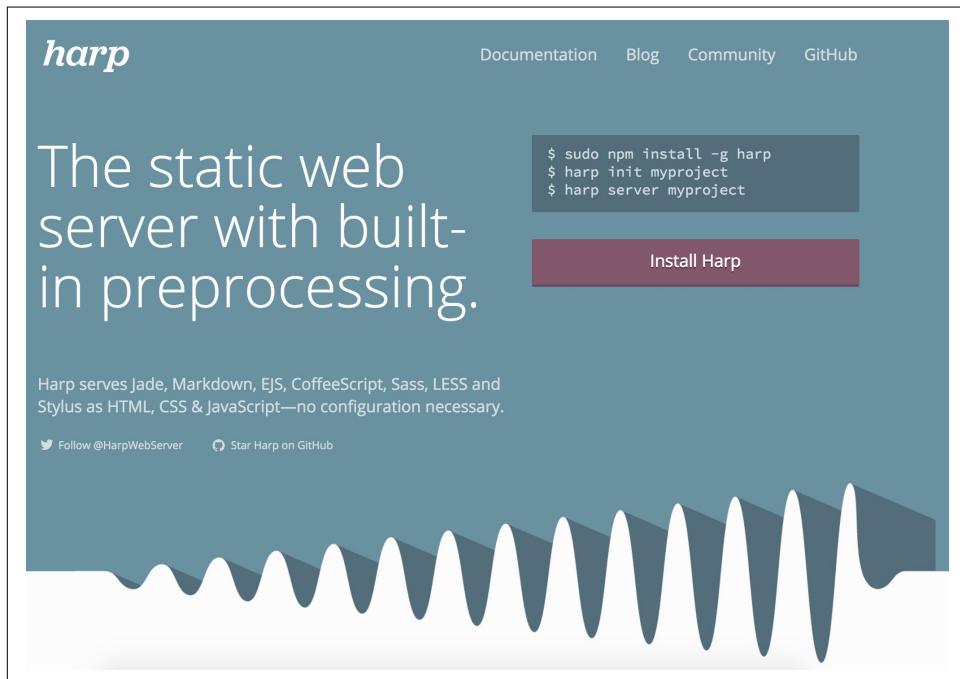
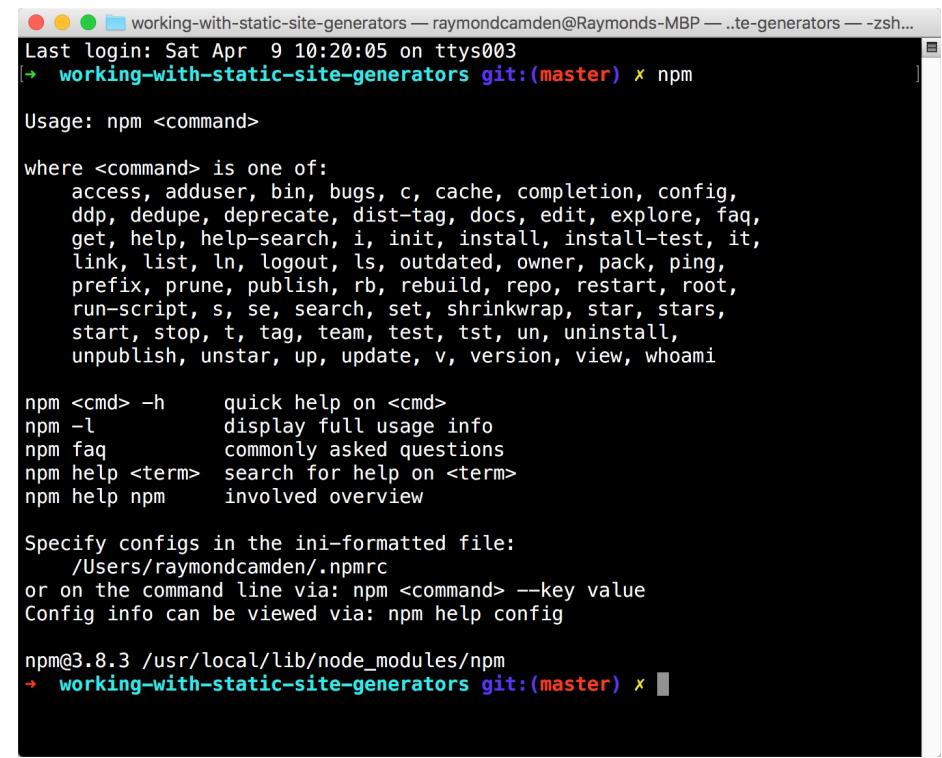


Figure 2-1. The Harp website

Harp requires minimal setup and focuses more on convention over configuration. That's a fancy way of saying that as long as you follow a few simple rules in terms of organization, Harp will just “work”. Let's start with installing Harp itself.

Before you begin, you'll need to ensure npm is installed on your system. npm stands for "Node Package Manager" and is used to handle common installation for programs. Imagine, for example, that a program requires some other program. Now imagine that program requires something else. The cool thing about npm is that an author can define what's required for it to run and the package manager handles the rest.

Because of how helpful this is, a great number of utilities make use of npm. You may actually have it installed already. The quickest way to see is to simply go to your terminal (or command prompt in Windows), and type npm.



```
working-with-static-site-generators — raymondcamden@Raymonds-MBP — ..te-generators — zsh...
Last login: Sat Apr 9 10:20:05 on ttys003
→ working-with-static-site-generators git:(master) x npm

Usage: npm <command>

where <command> is one of:
  access, adduser, bin, bugs, c, cache, completion, config,
  ddp, dedupe, deprecate, dist-tag, docs, edit, explore, faq,
  get, help, help-search, i, init, install, install-test, it,
  link, list, ln, logout, ls, outdated, owner, pack, ping,
  prefix, prune, publish, rb, rebuild, repo, restart, root,
  run-script, s, se, search, set, shrinkwrap, star, stars,
  start, stop, t, tag, team, test, tst, un, uninstall,
  unpublish, unstar, up, update, v, version, view, whoami

npm <cmd> -h      quick help on <cmd>
npm -l            display full usage info
npm faq           commonly asked questions
npm help <term>  search for help on <term>
npm help npm     involved overview

Specify configs in the ini-formatted file:
  /Users/raymondcamden/.npmrc
or on the command line via: npm <command> --key value
Config info can be viewed via: npm help config

npm@3.8.3 /usr/local/lib/node_modules/npm
→ working-with-static-site-generators git:(master) x
```

Figure 2-2. Checking for NPM

If this command fails to work, or if you know for sure you don't have npm installed, the easiest way to get it is to install Node.js. Don't worry - you don't have to actually *know* how to use Node nor will you ever need to learn it. (Although Node.js is pretty darn cool and recommended anyway.) Head over to nodejs.org and download the installer for your platform.

Assuming you've got Node installed and npm works for you, you can install Harp like so:

```
npm install -g harp
```

To confirm you installed Harp correctly, simply run `harp` in your terminal and you should see a nice usage summary.

```
→ working-with-static-site-generators git:(master) ✘ harp
Usage: harp [options] [command]

Commands:

  init [options] [path] Initialize a new Harp project in current directory
  server [options] [path] Start a Harp server in current directory
  multihost [options] [path] Start a Harp server to host a directory of Harp projects
  compile [options] [projectPath] [outputPath] Compile project to static assets (HTML, JS and CSS)

Options:

  -h, --help      output usage information
  -V, --version   output the version number

Use 'harp <command> --help' to get more information or visit http://harpjs.com/ to learn more.
```

Figure 2-3. Checking that Harp is installed.

Harp has a few different features, but it's simplest use is to enable a web server in a directory of files. Instead of just serving up files as they are, Harp supports different preprocessors that enable you to build dynamic resources.

To write HTML, you can use Markdown, Jade, or EJS (Embedded JavaScript). By simply using the proper extension (`.md`, `.jade`, and `.ejs`), Harp will automatically convert the syntax for each particular preprocessor into HTML. And of course, you can use regular HTML files too.

For CSS, you can use Less (`.less`), Sass (`.scss`), or Stylus (`.styl`). Regular CSS files work just fine too.

For JavaScript, the only option you have is CoffeeScript. Any file with a `.coffee` extension will be converted from CoffeeScript into JavaScript.

If it isn't obvious, you can mix and match any of the above preprocessors in any way you see fit. If you don't really care about CSS preprocessing and like JavaScript, then skip those preprocessors and just decide on which HTML preprocessor you want to use. Harp doesn't care and lets you decide which one works best for you and your particular project.

Your First Harp Project

The Harp CLI supports creating an initial “seed” application, but it may be easier to start simpler. Create an empty folder, it doesn’t matter what it’s called, and add a file called index.md. For our first test we’ll use Markdown, and again, this is completely arbitrary. Listing 1 shows the contents of the file. You can find this in the GitHub repo of code samples for this book in `code/harp/demo1/index.md`.

```
Hello World  
====  
  
This is a page. Woot.
```

At the command line, ensure you’re in the same directory containing the file, and run `harp server`. You should see Harp start up and describe where it is running.

```
→ demo1 git:(master) ✘ harp server  
-----  
Harp v0.20.3 - Chloi Inc. 2012–2015  
Your server is listening at http://localhost:9000/  
Press Ctrl+C to stop the server  
-----
```

Figure 2-4. Starting the Harp server

If you know open your web browser to the address and port shown above, you’ll see the web page rendered via Harp:

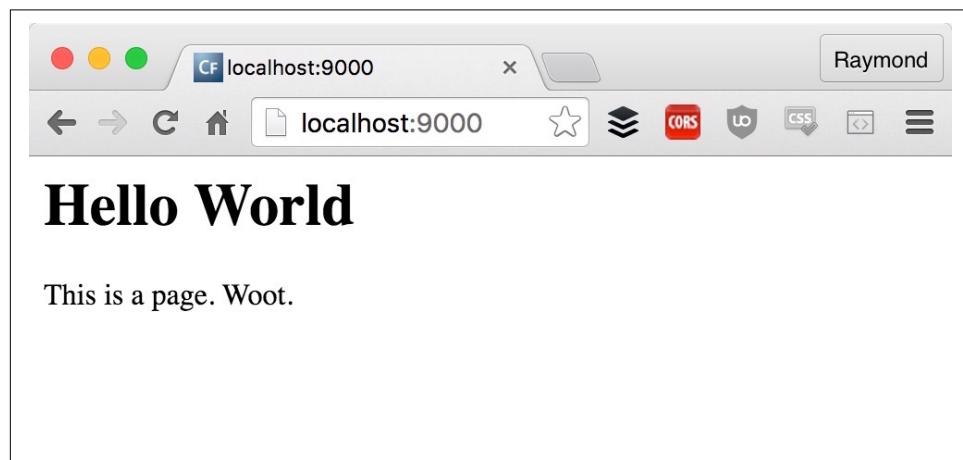


Figure 2-5. Harp rendering the Markdown page.

Notice how the Markdown was automatically converted into HTML. If you view source, or open up your browser developer tools, you can see this for yourself:

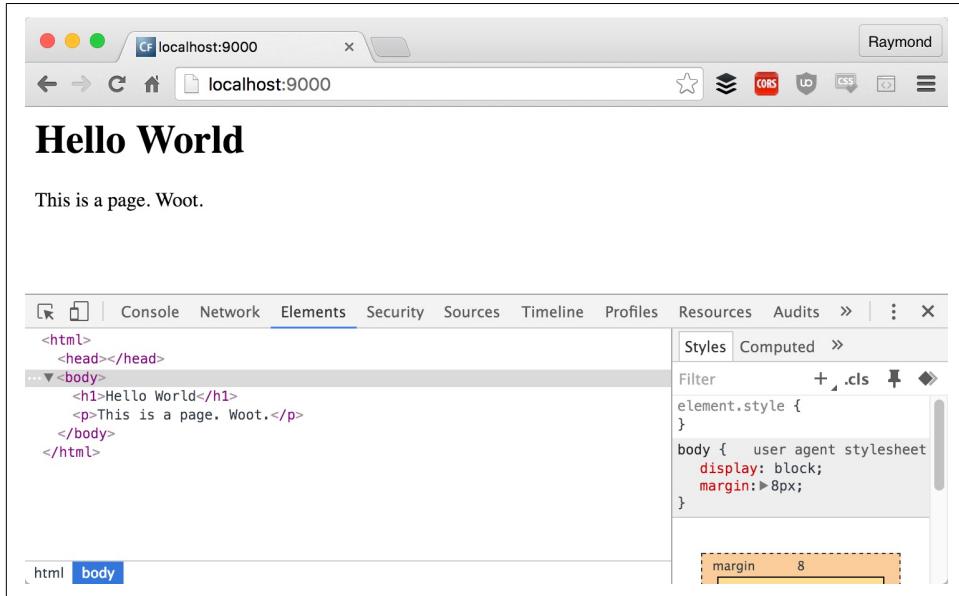


Figure 2-6. The Markdown code has been converted into HTML.

Along with converting the Markdown code into HTML, Harp actually makes your file available as an HTML file in the URL as well. If you actually try to access `index.html`, Harp will recognize that `index.md` represents this URL and will serve that file up. This is an important thing to remember. No matter what preprocessor you use, when you actually *link* to pages, or resources, you'll always use the “proper” type for what that preprocessor supports. So for example, if you want to add in a CSS file and you're using Less, you don't link to `styles.less`, but rather `styles.css` instead.



Learning More about Markdown

If you want to learn more about Markdown, check out its reference guide here: <http://daringfireball.net/projects/markdown/>.

Now let's make another file. `test.jade` will use the Jade template syntax. Listing 2 shows the contents.

```
h1 This is Jade  
  
a(href="index.html") Go Home
```

In the previous example, we've used Jade to create a header and a link back to the home page. Notice that we're linking to index.html. Again, you do *not* want to link to index.md as the final site, when static, will not include those extensions. Here is the result along with the rendered HTML shown in Chrome's Dev Tools.

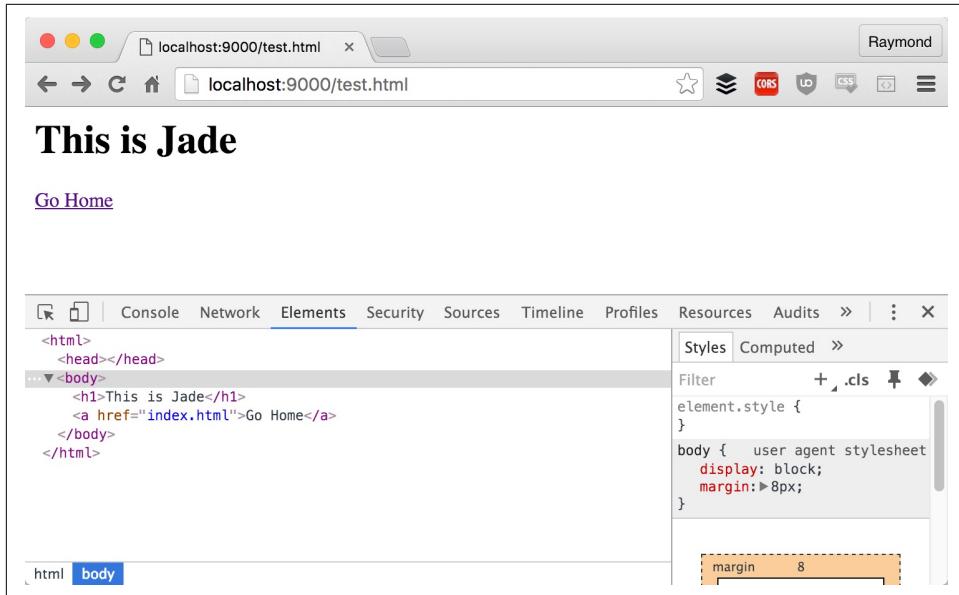


Figure 2-7. Jade code converted into HTML.



Learning More about Jade

You can learn more about Jade at its home page: <http://jade-lang.com>.

Finally, let's make another file to demonstrate EJS support. EJS is an older templating language and kind of resembles PHP or Classic ASP in some ways. Unlike Jade and Markdown which work from almost an abstraction layer over HTML, EJS requires you to write your HTML in - well - HTML. You only use EJS's template language when outputting something dynamic. Here's an example of that in action:

```
<h1>This is EJS</h1>
```

```
The time is: <%= new Date() %>
```

In the previous listing, the dynamic aspect begins with `<%=` and ends with `%>`. The inner portion is pure JavaScript and will be executed when the template is rendered. You can find this file from the book repo as `another-test.ejs`. As before, you can open

this in your browser by doing to another.html. You'll see the HTML rendered as is, but notice how the JavaScript was executed.

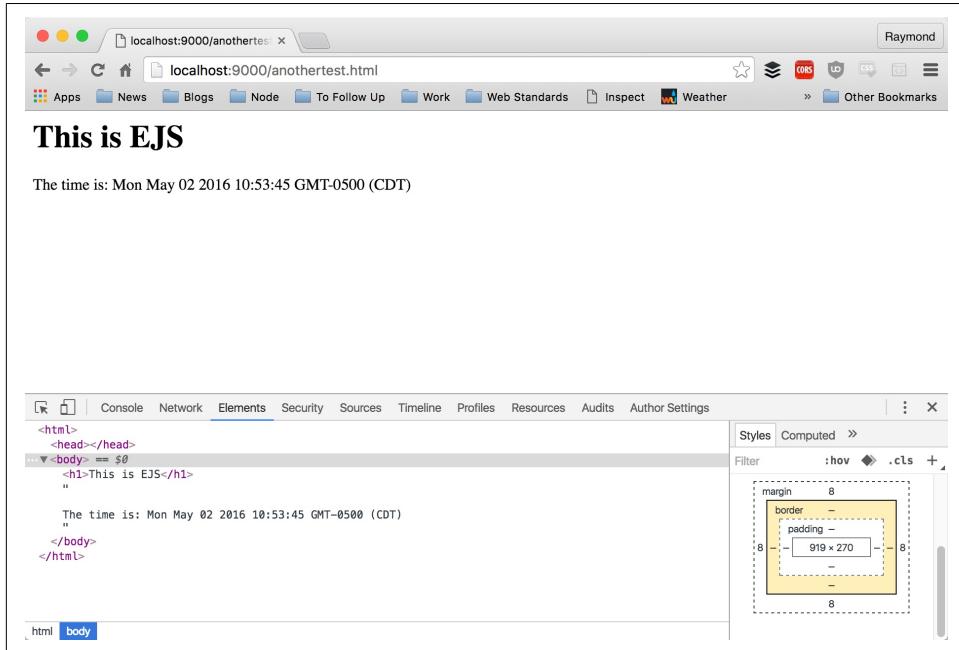


Figure 2-8. EJS code converted into HTML.



Learning More about EJS

You can learn more about EJS at its home page: <http://www.embeddedjs.com/>.

As you can see, you've got multiple different ways to write your HTML. We didn't even touch on the CSS preprocessors. So which do you choose? My suggestion is to play with both and get a feel as to which works best for you. The rest of the examples in this chapter will use Jade, but that is a **completely arbitrary** decision!

Working with Layouts and Partials

Now that you've seen how Harp preprocesses templates for you, it's time to kick it up a notch. One of the benefits that a CMS typically provides is the idea of a site wide layout. So for example, you may have a "design theme" for your site that sets a header, footer, basic colors, etc. The CMS then applies that layout to your content. Moving to static doesn't mean giving up that feature, of course. Harp provides support for layouts by looking for a file named `_layout.ext` where `ext` represents whatever prepro-

cesser you want to use. That means a request for `foo.ejs` will look for `_layout.ejs` or `_layout.jade`. Yes, you can mix up your layout preprocessor and regular content pre-processor but that's probably not a good idea.

Harp passes a variable called `yield` to the layout file that includes the contents of the file requested. Let's see this in action. If you are working with the code from the book repository, you can find the next example in the `demo2` folder. We've copied the files from `demo1`, but also added a new file, `_layout.jade`:

```
html
  head
    title Harp Site
  body
    != yield
    hr
    p Copyright #{new Date().getFullYear()}
```

This layout isn't terribly exciting, but notice in the middle where we've included the `yield` variable. This "sucks in" the page that was actually requested and puts it in the layout file. Beneath that a `hr` element is used along with a bit of code to include the current year in a copyright notice.

Now when you request any of the previous files you'll see them wrapped in the layout.

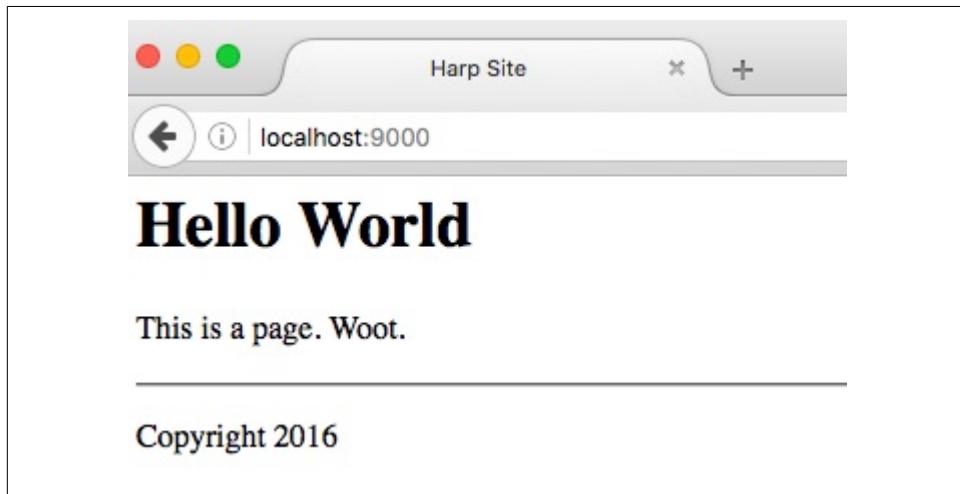


Figure 2-9. Layout being applied to pages.

By default Harp will look in the current folder before folders higher in the project directory. This means you can create a customized layout for a subdirectory if you want. Later on when we look at how to work with data you'll see another way to modify layout as well.

Partials are simply a way for one template to include another. Imagine you have a bit of boilerplate legal text, for example the typical license agreement text that no one actually reads. If you need that text in a few different pages, Harp provides a simple way to include that text in your templates. Let's look at an example.

```
h1 This is Jade  
  
!= partial("_legal")  
  
a(href="index.html") Go Home
```

In the Jade template above, the `partial` function is passed `_legal` as an argument. This tells Harp to look for a file named `_legal.ext` where the extension can be `.jade`, `.ejs`, `.md`. The extension is **not** included in the call. Here's a version in EJS.

```
<h1>This is EJS</h1>  
  
<%- partial("_legal") %>  
  
The time is: <%= new Date() %>
```

Same basic concept as the Jade version - you tell Harp the name of the file minus the extension. Unfortunately, you can't use the `partial` function within a Markdown file. You can include a Markdown file, but a Markdown file itself can't include items. As for what you put in the included file, it can be whatever you want. Here is a Jade template that simply outputs a tiny bit of legalese.

```
pre This is some boring legal text.
```

In the following screenshot, you can see both the Jade and EJS template including the same file.

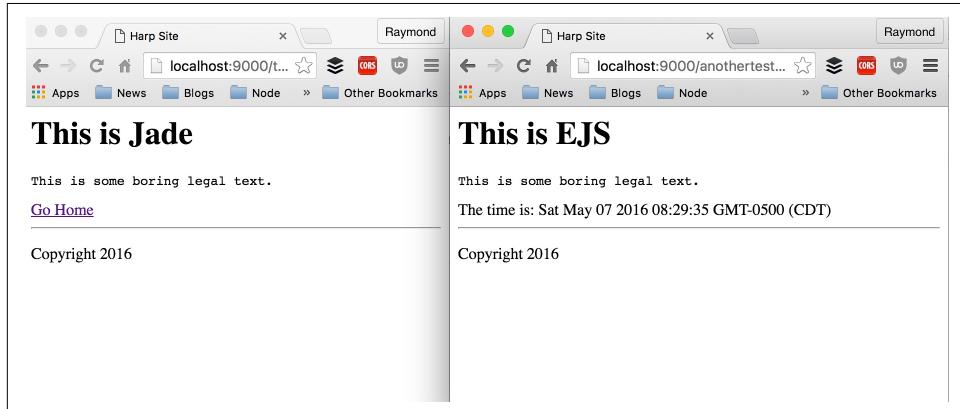


Figure 2-10. Examples of Partial.

Partials can be made even more powerful by passing variables to them. This covers the use-case where you want to include common code in templates but slightly tweak them every now and then. To pass a variable to an include loaded via `partial`, you simply include a second argument in your call. Here is an example.

```
h1 Testing Variables  
!= partial("_greeting", {name:"Raymond"})
```

In this Jade example (and of course you can do this in EJS as well), the second argument is a plain JavaScript object of name/value pairs. In this case, we have one value for `name`. You can have as many name/value pairs as you would like. On the partial side, you display the variable as you see fit. Here's `_greeting.jade`:

```
h3 Hello #{name}
```

Obviously this will display "Hello Raymond", and if you change the variable then the include will change as well.

All of the previous demos for the `partial` function may be found in the `demo3` folder. As a final note - you may wonder why the partials we used were named with an underscore in front. That is **not** a requirement. However, Harp has a feature where any filename beginning with an underscore will not be converted to a static file. For partials, it makes sense that we don't need them generated as is, so using an underscore in front of the name ensures we won't end up with stray static files we don't need.

Working with Data

So far we've seen how Harp can convert multiple types of templates into simple HTML pages (and don't forget Harp also supports this for CSS and JavaScript) as well as how to use layouts and partials. Now let's kick it up a notch and talk about how you can add data to a project to use within templates. Generally, data in Harp comes in two forms - global data and metadata. Global data is useful for things that apply to the site as a whole. For example, you may want to store a contact email address in data so you can easily change it one place. Conversely, metadata is more useful for describing specific parts of your site. So for example, given that you have a blog, you can use metadata to describe your blog entries (i.e., their titles and publication dates). Let's start with global data.

In your Harp project, you can add a file named `_harp.json`. The contents should be valid JSON and contain a top level "key" called `globals`. Inside of this should be a set of name/value pairs for whatever data you want to use in your site. Here is an example.

```
{  
  "globals":{
```

```

        "title": "My Site!",
        "owner": "Raymond Camden"
    }
}

```

In this example, we have two values - one for `title` and one for `owner`. To be clear, this is **completely arbitrary**. You can use whatever makes sense for you here. Your values also need not be simple strings. As long as it is valid JSON than you can use whatever you want, including arrays.

```

{
  "globals": {
    "title": "My Site!",
    "owner": "Raymond Camden",
    "subjects": ["Math", "Science", "Beer"]
  }
}

```

When defined, you can then use these variables in any of your templates, whether it be a regular template, layout file, or partial. How you use it depends on the language. For Jade, it would look like this:

```
p #{ owner }
```

In EJS it would look like so:

```
<p><%= owner %></p>
```

Let's look at an example of this in action. In the `demo4` folder you'll find a complete example that includes a home page, a layout, and `_harp.json` file. The contents of the `globals` match the previous code example so we won't share it again. Here's the home page.

```

h1 This is my site.

p Welcome to my site. Sorry this isn't more exciting.

p This site was made by #{owner}.

```

The `owner` value from `_harp.json` is used within a simple paragraph tag. Now let's look at the layout.

```

html
  head
    title #{title}
  body
    != yield
    hr
    p Copyright #{new Date().getFullYear()}

```

In this version, the layout is now dynamic based on the global variable. Note that we still have a dynamic copyright as well. You can mix and match global variables with

things defined locally on the template. We could replace that code with a global variable, but then we'd have to edit the year every New Year's Eve, and no one wants to do that.

Once run, you can see it in action. Simply change the values, reload the page, and you can see your change reflected immediately.

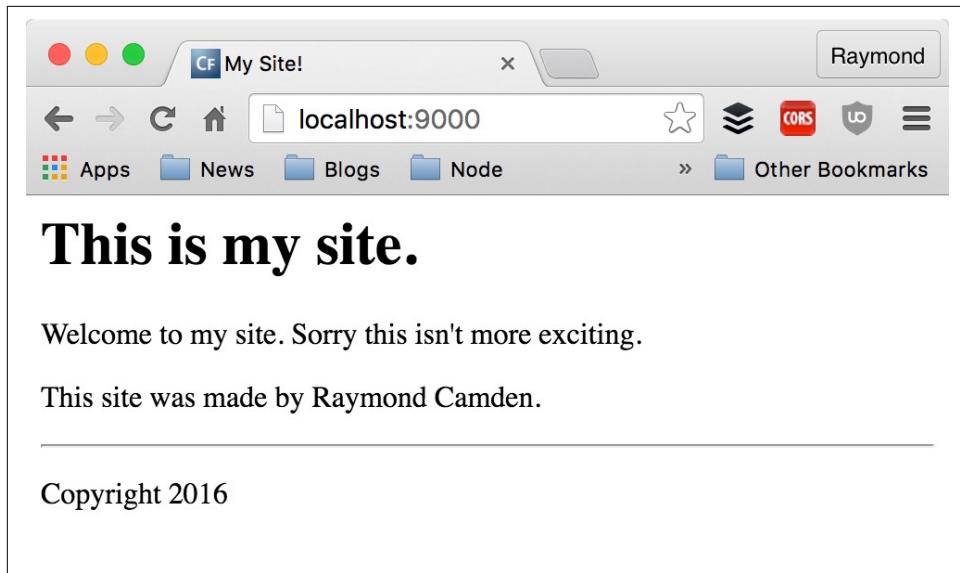


Figure 2-11. Examples of Global Variables.

Now let's consider metadata. This can be a slightly confusing topic so let's start off with an example. Imagine you're building a simple store. Your website will consist of a home page and a product page for each of your products. The folder, `demo5`, will be the basis for this site. You can see it has a home page, a layout, and then a folder of products.

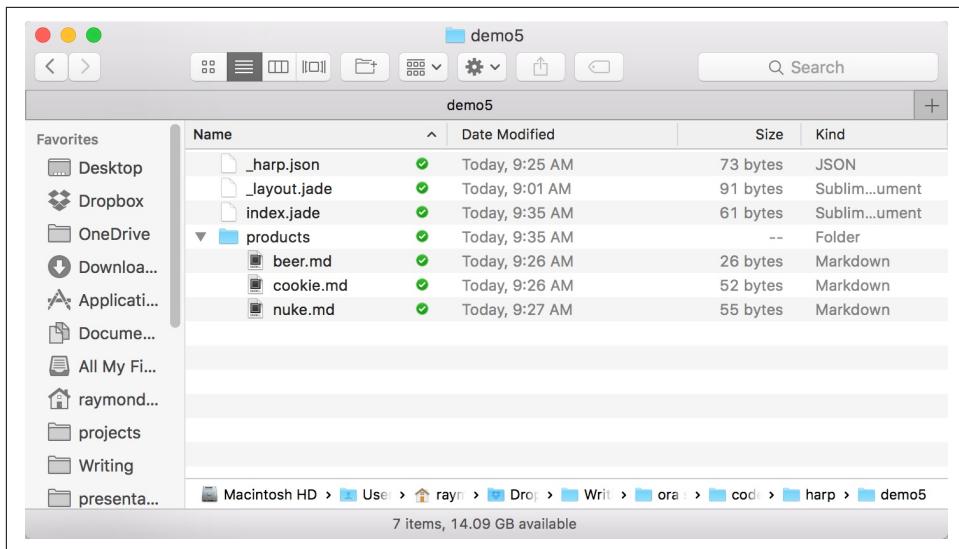


Figure 2-12. Version one of the product site.

If you run this with `harp server`, you can request the home page and then each individual product page. If you wanted to link to the products from the home page, how would you do it?

Since you only have three products, you could just hard code a simple list. That certainly works, but isn't really scalable. What if you had a hundred products? What if you wanted your home page to only list the most recently released products and keep a complete list of products on some other page? This is where metadata comes into play.

To begin, you create a new file in your `products` folder called `_data.json`. As with the `globals` file, it must be a valid JSON file. In this file, you describe your data. Let's consider the following example.

```
{  
  "nuke": {  
    "title": "A Nuke",  
    "price": 9.99  
  },  
  "cookie": {  
    "title": "Cookies",  
    "price": 2.99  
  },  
  "beer": {  
    "title": "Beer",  
    "price": 8.99  
  }  
}
```

In this JSON file we have three main parts - one for each product. The names, nuke, cookie, and beer match the file names from the demo. This will become important in a moment. Inside each block of data there's two variables - title and price. As before, this is arbitrary. But here's where things get interesting.

The first change is that Harp will now recognize this data and make it available to your templates. Any template can access this data via a new variable, public.products._data. The public variable is always available in Harp projects. The products key comes from the fact that we have a folder called products. Finally, _data maps to the _data.json file itself. What this means then is that we can create a dynamic list of products. Here's an updated home page that now creates a list of links.

```
h1 Welcome to our store

ul
  for product, link in public.products._data
    li
      a(href='/products/#{{link}}.html') #{product.title} ($#{product.price})

p This site was made by #{owner}.
```

Let's break this down. The for loop iterates over the data grabbing two values. product represents the individual product. Notice how we include the title and price. The link variable is the top level key in the JSON file, namely "nuke", "cookie", and "beer". Remember how we said we made it match the file names? That then lets us create links to the product files. Because we have access to the raw data, we could even do something complex, for example, resorting the values by price or by title. Here's a screen shot of the home page now including the links.

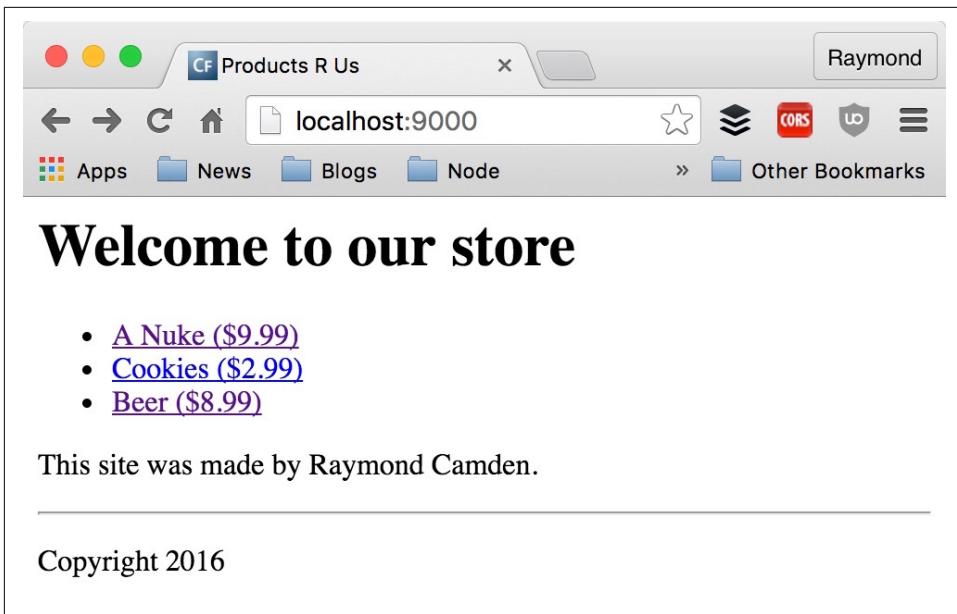


Figure 2-13. Product links.

Now you can click on the links and go directly to each product page. But here's where things get even more interesting. Consider the beer page.

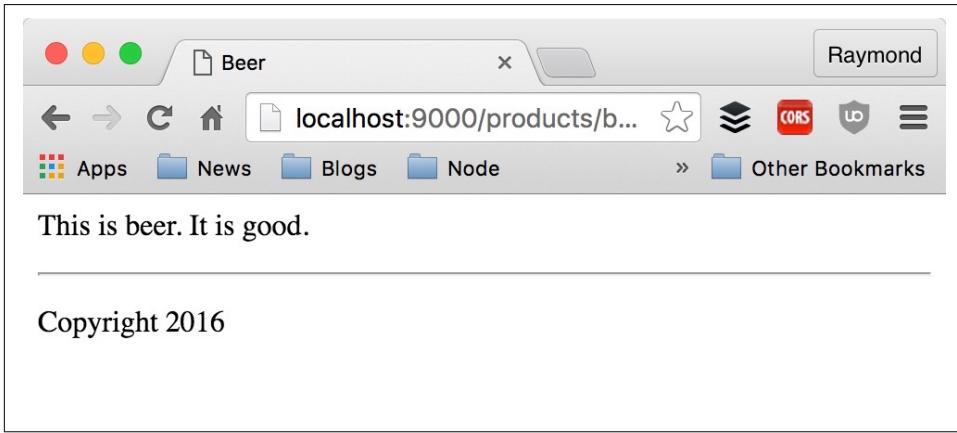


Figure 2-14. The Beer page.

Notice anything in particular? See how the title changed? What happened there? When Harp works with variables, it does something special with metadata. Harp recognized that `beer.md` was being loaded, and that “beer” matched the value in `_data.json`. Because of this match, when Harp checks variables, it will look in

`_data.json` before it checks globals defined in `_harp.json`. This allows you to dynamically update values on a page by page basis. It's perfect for cases where you may want a default, global title for a site but then have specific titles for each product.

Generating a Site

Now that you've seen the basics of how to use Harp and work with dynamic templates, how do you actually create static output? The basic command with the CLI is: `harp compile pathToCode pathToOutput`. A full example could look like this: `harp compile ./ ../output`. This tells Harp to compile the current directory and output the result in a folder above the current path named `output`. Given the input from the last demo (`demo6`), the output looks like this:

Name	Date Modified
index.html	Today, 11:10 AM
products	Today, 11:10 AM
beer.html	Today, 11:10 AM
cookie.html	Today, 11:10 AM
nuke.html	Today, 11:10 AM

Figure 2-15. Compiled output.

Notice how the `.jade` and `.md` files are now all regular HTML files. Also, all the "special" files (layout and data related) are gone.

Building Camden Grounds

We began this chapter by explaining what we would build as an example of a fairly simple site - a coffee shop named "Camden Grounds". We described this as a site with four pages - a home page, a menu, and list of locations, and a simple "About Us" that frankly no one will ever read. (But don't tell the client that!) To create this site, we first need a design. If you're like me (with little to no design skill), you'll probably want to either hire a designer or find a website template that you can use. Luckily, the website freewebsitetemplates.com actually has one called "Coffee Shop Web Template" (<https://freewebsitetemplates.com/forums/threads/coffee-shop-web-template.20350/>) that is both perfect and free.

As a word of caution, before deciding to use any template, it is a good idea to take a look at the source code behind it. It could be an incredibly good looking web page

with an absolute horrible mess of code behind it. In this case, the shop template was fairly simple and easy to work with. Here is the template in its original form:



Figure 2-16. Original template.

It matches perfectly with our requirements except for the “Blog” item. Luckily though we can just remove it. You can find the complete source code for this demo in the ch2/camdengrounds folder, but let’s go over the files one by one so you can see how it is put together.

First, let’s look at the template. This was created by looking at a few of the files from the template and determining which content remained the same. As you can imagine, this is the header and footer. Since a majority of the template is regular HTML, the listing below focuses on the dynamic, more interesting aspects of the template. See the original file (`_layout.ejs`) for the complete listing.

```

<div id="header">
  <a href="/index.html">
    
  </a>
  <ul>
    <li
      <% if(current.source == 'index') { %>
        class="current"
      <% } %>
    >

      <a href="/index.html">Home</a>
    </li>
    <li
      <% if(current.source == 'menu') { %>
        class="current"
      <% } %>
    >
      <a href="/menu.html">Menu</a>
    </li>
    <li
      <% if(current.source == 'locations') { %>
        class="current"
      <% } %>
    >
      <a href="/locations.html">Locations</a>
    </li>
    <li
      <% if(current.source == 'about') { %>
        class="current"
      <% } %>
    >
      <a href="/about.html">About Us</a>
    </li>
  </ul>
</div>
<div id="body">
  <%- yield %>
</div>

```

The first thing to make note of is in the header menu, note the use of a variable called **current**:

```

<li
  <% if(current.source == 'index') { %>
    class="current"
  <% } %>
>
  <a href="/index.html">Home</a>
</li>

```

The `current` object is a helper value that Harp provides to each template. It has two values: `source` and `path`. The `path` value represents the current “directory heirarchy” of the request. So given a request like so, `raymondcamden.com/products/weapons/foo.html`, the `path` value would be an array consisting of: `products`, `weapons`, `foo`. The `source` value is just the very end of the heirarchy, so given the same URL, you would get a value of `foo`. As you can imagine, this is a useful way of saying, “I’m on page `so` and so, do so and so.” The code snippet above basically handles adding a CSS class to each menu item when the user is on a particular page. Similar code is used in the bottom menu. Finally, don’t forget that your layout file has to include the `yield` variable to display the contents of the template.

Let’s now look at the pages. It will be easier to start with the simplest page, `about.html`, as it is just simple text. Remember that we don’t have to include the layout. Here is `about.ejs` with some of the boilerplate text removed to save space:

```
<div id="figure">
    
    <span>Lorem ipsum dolor sit amet.</span>
</div>
<div>
    <a href="about.html" class="about">About</a>
    <div>

        <h3>We Have Coffee for Everyone</h3>
        <p>
            Mauris sed libero ac neque lobortis aliquam. Vivamus vitae ultricies.
        </p>

        <h3>We Even Have Tea!</h3>
        <p>
            Sed a pretium risus, ut volutpat nunc. Donec blandit orci id sollicitin.
        </p>

        <h3>We Don't Have Beer</h3>
        <p>
            Proin dapibus, orci vitae bibendum laoreet, libero velit condimentum
        </p>

    </div>
</div>
```

There’s nothing dynamic here so the only real content is the text describing the company. Here it is as run under Harp - notice how the menu recognizes what page is being displayed.

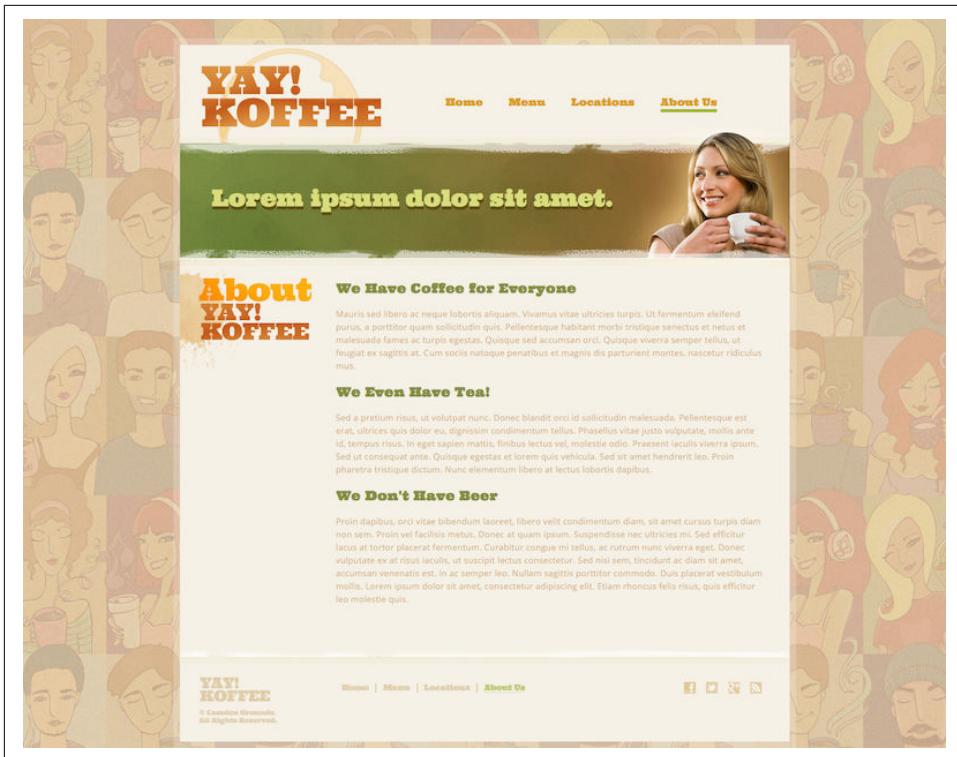


Figure 2-17. About Us Page.

Now let's move to the Locations page. This one will be a bit tricky. While there aren't many locations now, Camden Grounds hopes to grow into a mega-coffee-serving chain to rival that of a certain company out of Seattle. The locations page therefore should be dynamic to make it easy to add locations later on. Let's first look at a screen shot, and then we'll explain how it was built.

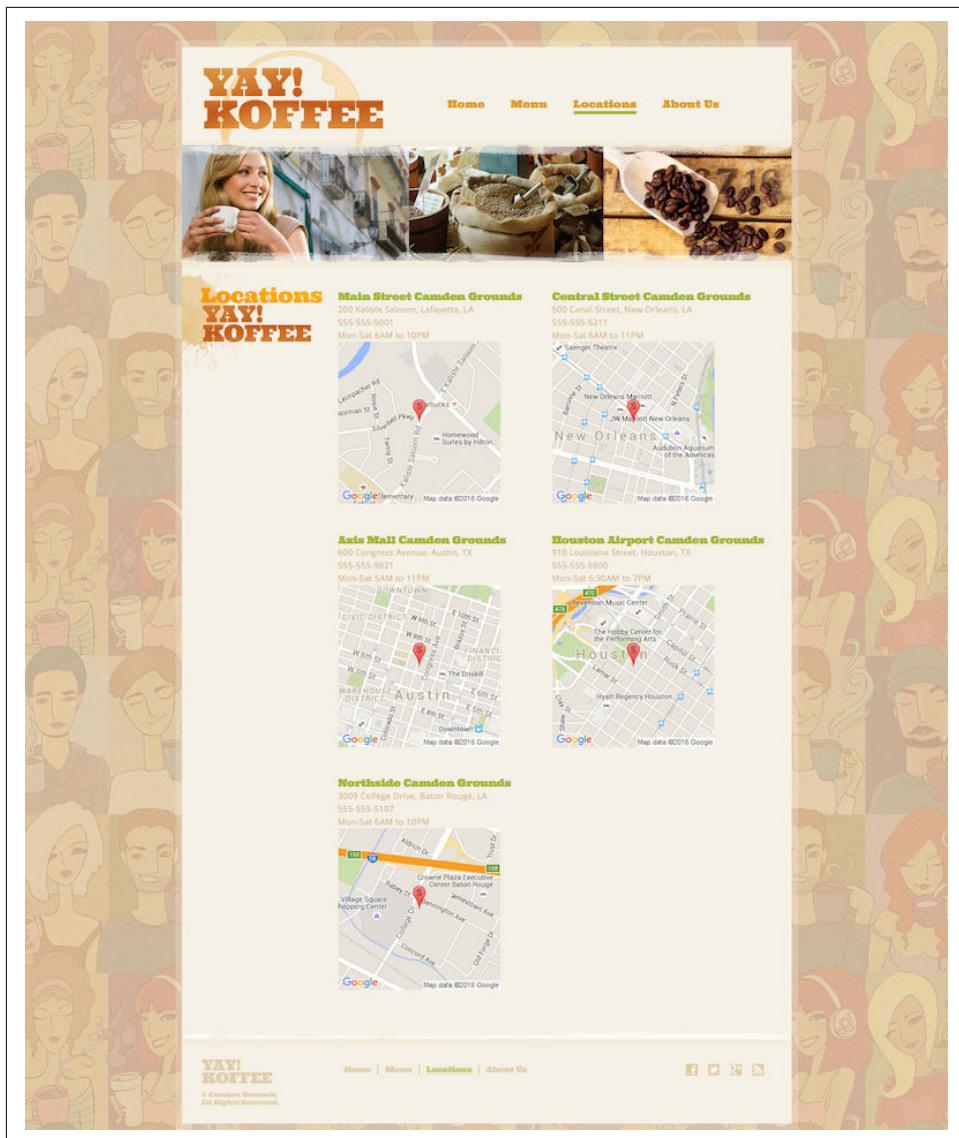


Figure 2-18. Locations Page.

Each location is rendered as a name, address, phone numbers, hours of operation, and a little map. How is this done? First, we added location data to the `_harp.json` file as a global. Here is that file, with a few locations removed for brevity.

```
{  
  "globals":{
```

```

    "locations": [
      {
        "name": "Main Street Camden Grounds",
        "address": "200 Kaliste Saloom, Lafayette, LA",
        "hours": "Mon-Sat 6AM to 10PM",
        "phone": "555-555-5001"
      },
      {
        "name": "Northside Camden Grounds",
        "address": "3009 College Drive, Baton Rouge, LA",
        "hours": "Mon-Sat 6AM to 10PM",
        "phone": "555-555-5107"
      }
    ]
  }
}

```

Now let's look at the template code.

```

<div id="figure">
  
</div>
<div>
  <a href="locations.html" class="locations">Locations</a>
  <div>
    <% for(idx in locations) { %>
    <dl>
      <dt><%- locations[idx].name %></dt>
      <dd><%- locations[idx].address %></dd>
      <dd><%- locations[idx].phone %></dd>
      <dd><%- locations[idx].hours %></dd>
      <dd>
        
      </dd>
    </dl>
    <% } %>
  </div>
</div>

```

The list of locations is iterated as a JavaScript array. Each item is output per the original template design but with some modifications to fit our data. Finally, the Google Static Image API (a very handy Google service that doesn't get a lot of attention) is used to display a small map of the location. Note that Google requires you to get a key to use this service. There is a free tier that is more than appropriate for simple static sites.

Fairly simple - but the big win here is that when Camden Grounds expands, all you need to do is edit the pure data to have the site updated. (And, of course, generate the static version and deploy it.)

Now it's time to kick it up a notch. To handle the menu for Camden Grounds, it will be a two step process. First, a unique page will be made for each menu item. We'll store that under a `coffees` directory. Each page will include text about the product, but for now we've kept it down to just one simple sentence. As before, let's look at a menu page and then we'll cover how it was built. Here is `coffee1.ejs`:

```
<%- partial("_coffee_header") %>

<p>
This is coffee 1.
</p>

<%- partial("_coffee_footer") %>
```

Obviously this would be a bit longer for a real product, but you can see where we're using partials to include header and footer layout for the item. (As an aside, you can do nested layouts in Harp, but it is somewhat complex.) The footer is just a few closing `div` tags, but the header is a bit dynamic:

```
<div id="figure">
    
    <span><%- name %></span>
</div>

<div>
    <a href="/menu.html" class="whatshot">What's Hot</a>
    <div>

        " style="float:right">
```

What's going on here? We've outputting a name and image, but where does this come from? Well remember that we can supply metadata for our content. Let's now take a look at that. (As with other code listings, we've trimmed it a bit.)

```
{
    "coffee1": {
        "name": "Coffee One",
        "price": 2.99,
        "image": "coffee1.jpg",
        "short": "Creamy"
    },
    "coffee2": {
        "name": "Coffee Two",
        "price": 9.99,
        "image": "coffee2.jpg",
        "short": "Rich"
    },
}
```

```

"coffee6": {
    "name": "Coffee Six",
    "price": 4.99,
    "image": "coffee6.jpg",
    "short": "Weird"
}
}

```

Each menu item has a key that matches with its file name, and then some basic information about the menu item. Notice we aren't using all of that data. You'll see where we actually do in a moment. Here is the page for the first coffee.



Figure 2-19. Menu item Page.

So what about the menu itself? Since we've created a few products and have metadata defined in `_data.json`, we can make use of it to render a dynamic menu.

```

<ul>
<% for(idx in public.coffees._data) { %>
<li>
    <a href="/coffees/<%- idx %>.html">
        "></a>
        <div>
            <a href="/coffees/<%- idx %>.html">
                <%- public.coffees._data[idx].name %></a>

```

```

<p>
<%- public.coffees._data[idx].short %>
&#36;<%- public.coffees._data[idx].price %>
</p>
</div>
</li>

<% } %>
</ul>

```

Remember that Harp places metadata in a `public.X` object where X represents the folder containing the data file. This lets us loop over each menu item and list out the name, short description, price, and image. Here is the menu.



Figure 2-20. Menu Page.

We're almost done. The final page is the home page. This is somewhat modified from the original template as we don't have a blog anymore. Instead, we're going to display 3 products. In theory we could sort these by the most expensive, or best selling, but for now we'll just grab the first three. Here's the code.

```

<div id="figure">
  
  <span id="home">Camden Grounds is the best darn coffee in the world!
    <a href="about.html">Find out why.</a></span>
</div>
<div id="featured">
  <span class="whatshot"><a href="menu.html">Find out more</a></span>
  <div>
    <%
      coffeeKeys = Object.keys(public.coffees._data);
      for(var x = 0; x<Math.min(coffeeKeys.length,3);x++) {
        coffee = public.coffees._data[coffeeKeys[x]];
    %>
    <a href="/coffees/<%- coffeeKeys[x] %>.html">
      "></a>
    <% } %>
  </div>
</div>

```

In general, this isn't too much different from the menu page, but instead we grab the keys from the data and loop until we hit either the total number of items or 3. The result is pretty much as you expect.



Figure 2-21. Home Page.

Obviously there's more that could be done to this site, but hopefully you can see how Harp really makes it easy to manage this simple little website.

Going Further with Harp

As we said in the beginning, we were only going to scratch the surface of the Harp static site generator. Here is a quick look at some of the features we did not cover in the chapter.

- While we documented it, we didn't show any CSS or JavaScript preprocessing. Don't forget that Harp allows for this as well.
- Harp provides access to a `_contents` variable that represents each folder in the project. This could be useful for generating a dynamic list of images for an art gallery.
- Harp sets a `environment` variable representing development versus production. This lets you toggle certain things based on where your code is running.
- Harp provides basic 404 and client-side routing support.
- Harp can be used inside another Node.js application.

For more information, see the Harp documentation: <http://harpjs.com/docs/>

CHAPTER 3

Building a Blog

Blogs are one of the most popular type of sites found on the Internet. In fact, WordPress, an open source blogging engine, is currently in use by over 60 million users.

At its heart, a blog is fairly simple. Like a diary, each entry in the blog is an individual story and typically presented to the user in reverse chronological order. Blogs will also usually have categories to organize entries. Visitors can then read more entries in a particular category to focus on things that may interest them more.

In this chapter, we're going to build a simple blog. The static site generator we'll be using in this chapter actually builds a blog out of the box, so not much time will be spent on building a *specific* blog per se, but we will demonstrate how to find a real blog template, implement it, and then create some temporary content so we can see the blog in action.



Last Minute Change

Jekyll released a major update right before the publication of this book. While basic operations are the same, some screen shots may look slightly different on your installation. Any serious issues will be reported in the books errata.

Blogging with Jekyll

Jekyll (<https://jekyllrb.com>) is a static site generator focused on creating blogs. While it certainly can be used to build non-blog sites, out of the box it's primary use is creating blogs. It also integrates natively with GitHub Pages which provides you a free hosting option for your site as well. (Since the site has to run on a GitHub project itself, naturally the blog then has to be *about* the project itself.)

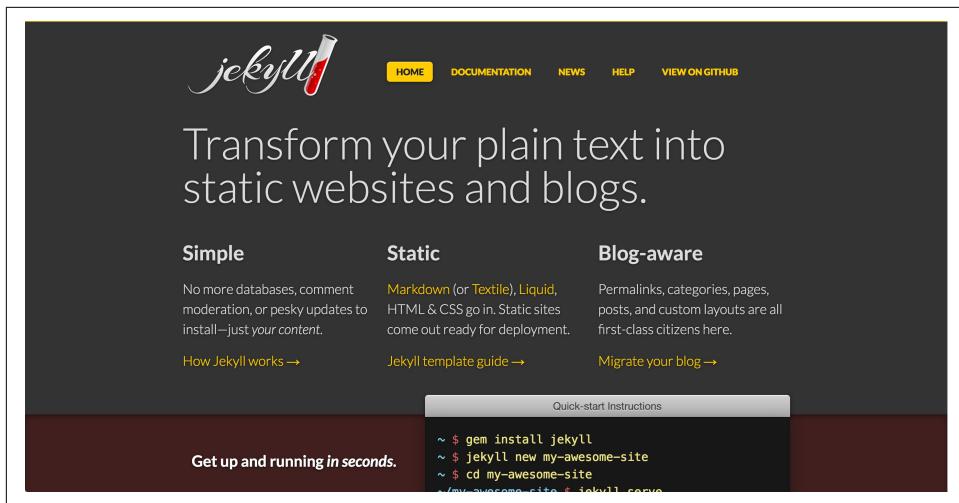


Figure 3-1. The Jekyll website

Installation of Jekyll is a bit more complex than other static site generators. Right away, be aware that Jekyll is **not** officially supported on Windows. You can find documentation for “making it work” on Windows, and that should be good enough to follow along with this chapter, but if your primary OS is Windows, you may to triple check how well it works before committing to Jekyll.

Before installing Jekyll, you need to install Ruby (<https://www.ruby-lang.org/en/downloads/>) and RubyGems (<https://rubygems.org/pages/download>). You won’t need to know anything about Ruby to use Jekyll (although it helps when building custom plugins), so don’t worry if you’ve never worked with the programming language before.

Once you have RubyGems installed, you can use the command line to install Jekyll:

```
gem install jekyll
```

Once installed, open up your terminal and type in `jekyll` to ensure it’s been installed correctly. The figure below shows what you should see.

```

➔ working-with-static-site-generators git:(master) ✘ jekyll
A subcommand is required.
jekyll 3.1.3 -- Jekyll is a blog-aware, static site generator in Ruby

Usage:

  jekyll <subcommand> [options]

Options:
  -s, --source [DIR]  Source directory (defaults to ./)
  -d, --destination [DIR] Destination directory (defaults to ./_site)
  --safe              Safe mode (defaults to false)
  -p, --plugins PLUGINS_DIR1[,PLUGINS_DIR2[,...]] Plugins directory (defaults to ./_plugins)
  --layouts DIR      Layouts directory (defaults to ./_layouts)
  --profile          Generate a Liquid rendering profile
  -h, --help          Show this message
  -v, --version       Print the name and version
  -t, --trace         Show the full backtrace when an error occurs

Subcommands:
  docs
  build, b           Build your site
  clean
  doctor, hyde       Search site and print specific deprecation warnings
  help
  new
  serve, server, s  Serve your site locally
  import             Import your old blog to Jekyll

```

Figure 3-2. Testing for Jekyll at the prompt

Jekyll supports creating content in either regular HTML or Markdown files. HTML files are served as is, but Markdown files will be converted to HTML first. On top of this, it supports a template language called Liquid (<https://github.com/Shopify/liquid/wiki>). Liquid is incredibly powerful and you'll see multiple examples in this chapter, but it integrates well with your existing HTML. This is where you'll create dynamic content locally that ends up being static when done. Adding another layer to the mix is "front matter". Front matter is metadata on top of the page that tells Jekyll how to parse the page. All of this will make a lot more sense once you see a few examples.

Your First Jekyll Project

Now that you have Jekyll installed, it is time to create your first site. Out of the box, Jekyll creates a "complete" but mostly empty blog. You'll have a layout, a home page showing your blog posts, and one written post. Creating a new Jekyll site is as easy as typing `jekyll new foo` where `foo` represents the path to the new site.

```

➔ ch3 git:(master) ✘ jekyll new demo1
New jekyll site installed in /Users/raymondcamden/Dropbox/Writing/ore static sites/ch3/demo1.
➔ ch3 git:(master) ✘

```

Figure 3-3. Creating a new Jekyll site

Jekyll will install various dependancies necessary for the site and default theme and then complete the installation. Now, change directories into the new directory the CLI created and start the server by running `jekyll serve`. (Jekyll provides a few alternatives to that command as well, like `jekyll s`.)

```
➔ demo1 git:(master) ✘ jekyll serve
Configuration file: /Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/demo1/_config.yml
  Source: /Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/demo1
  Destination: /Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/demo1/_site
  Incremental build: disabled. Enable with --incremental
    Generating...
      done in 0.48 seconds.
  Auto-regeneration: enabled for '/Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/demo1'
Configuration file: /Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/demo1/_config.yml
  Server address: http://127.0.0.1:4000/
  Server running... press ctrl-c to stop.
```

Figure 3-4. Starting up the Jekyll server

You can now open up your browser to the address reported by the command line:

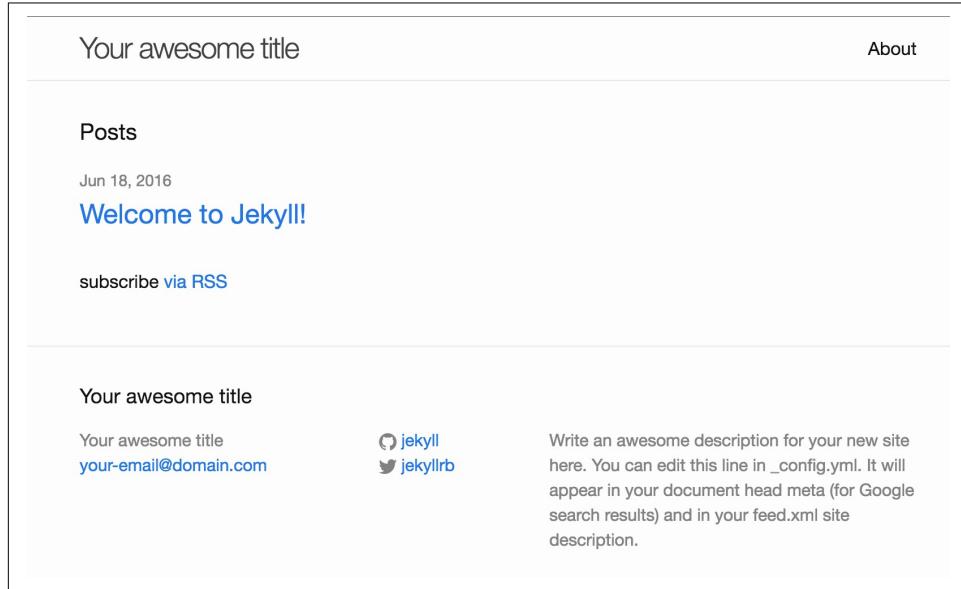


Figure 3-5. The default site created by Jekyll

There's a few things to note here. First off, the simplistic UI you see is just a default. It is absolutely **not** something you have to use. Later in the chapter you'll see an example that has its own custom UI. You also have two content examples. The [About](#) link on top shows a simple page while the [Welcome to Jekyll!](#) link shows you a sample blog post. You should definitely click that the link as the page provides some good explanatory text along with a sample of how code snippets are rendered.

Welcome to Jekyll!

Jun 18, 2016

You'll find this post in your `_posts` directory. Go ahead and edit it and re-build the site to see your changes. You can rebuild the site in many different ways, but the most common way is to run `jekyll serve`, which launches a web server and auto-regenerates your site when a file is updated.

To add new posts, simply add a file in the `_posts` directory that follows the convention `YYYY-MM-DD-name-of-post.ext` and includes the necessary front matter. Take a look at the source for this post to get an idea about how it works.

Jekyll also offers powerful support for code snippets:

```
def print_hi(name)
  puts "Hi, #{name}"
end
print_hi('Tom')
#=> prints 'Hi, Tom' to STDOUT.
```

Check out the [Jekyll docs](#) for more info on how to get the most out of Jekyll. File all bugs/feature requests at [Jekyll's GitHub repo](#). If you have questions, you can ask them on [Jekyll Talk](#).

Figure 3-6. A sample post from the Jekyll server

Now let's take a look at the directory structure of your Jekyll project.



Another Last Minute Change

Remember how we warned you about that last minute change? The very latest version of Jekyll uses a theme that runs from a special directory on your machine used by RubyGems. This means you will *not* see the theme files in your site. This is a cool change in that it keeps themes in one central location, but on the other hand, it makes it a bit difficult to “play” around and tweak a theme to your liking. In order to correct this, you can fix ask for the location of the theme by doing this:

```
bundle show minima
```

In the command above, `minima` is the name of the theme. This will return a path like this:

```
/Library/Ruby/Gems/2.0.0/gems/minima-2.0.0
```

Given that you're in the directory of the Jekyll site you just created, you can copy the files like so:

```
cp -r /Library/Ruby/Gems/2.0.0/gems/minima-2.0.0/ .
```

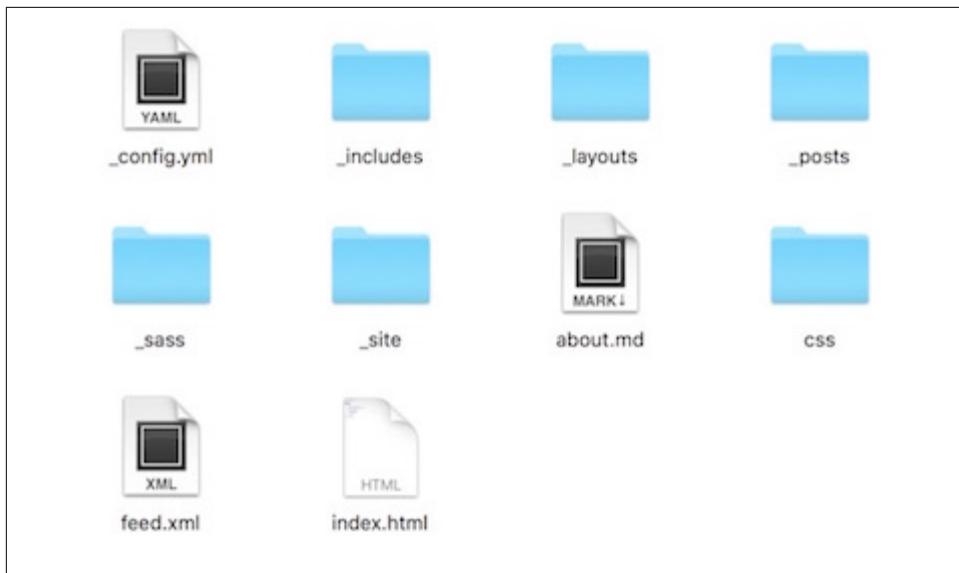


Figure 3-7. Folders and files created by Jekyll

Let's discuss these files and folders left to right, top to bottom:

- `_config.yml` is a configuration file that helps drive how Jekyll operates. You can provide configuration values here as well as in the command line, but it is easier to set things in the file so you don't have to repeat yourself at the command prompt.
- `_includes` is a folder specifically for items you will include in your templates. So for example, a common "boilerplate" legal agreement used in multiple locations could be written as an include. Jekyll refers to these files as "partials".
- `_layouts` is where you would put files that handle the layout of your site. You may only have one layout, you may have multiple. Jekyll supports creating as many as you would like.
- `_posts` is where you would create new files for blog content.
- `_sass` is for Sass CSS files. They will automatically be converted into CSS by Jekyll and copied to the `css` folder when creating your static site.
- `_site` is the static version of your site. This is created as you work. You can look at it now to see the static version of the blog you just created.
- `about.md` and `index.html` are files created by Jekyll by default. They represent a simple "About" page and a home page. You can remove any of these (although you should keep a home page) and add whatever makes sense. For example, if you wanted a "Contact" page, you could add it here. Note that using Markdown is optional and you can pick and choose which files you want to use Markdown with, as the default files demonstrate.

- Any other folder you create that Jekyll doesn't recognize, or work with, will be copied as is into your static site. So for example, if you added a `js` folder for your JavaScripts, these would be copied as is.



Additional Folders

There are more folders that Jekyll will automatically recognize and you'll see examples of them later in the book.

Writing a Post

Now that we've got a blog up and running, let's create a new post. When you write a new post, Jekyll expects you to follow a particular naming scheme for the file. The format is: `YEAR-MONTH-DAY-title.extension`. Create a new file called `2016-06-25-new-post.html`. By using `html` as the extension, we're specifically saying we don't want to use Markdown. Feel free to use `md` (or `markdown`) if you want. Also feel free to change the date values to the current date.

Every post you write must begin with "Front Matter", this is simply metadata about the post that lives on top of the page and is a format called "YAML". It may help to take a look at the YAML in the existing page created by Jekyll. (Note, your YAML may be slightly different based on the date you created the demo.)

```
---
layout: post
title: "Welcome to Jekyll!"
date: 2016-06-18 09:18:33 -0500
categories: jekyll update
---
```

YAML formats data in a simple key/value system. You can see four keys in the example above and four values. You can get more complex if you need to. The Wikipedia page (<https://en.wikipedia.org/wiki/YAML>) is a good resource for learning more about the YAML format, but you won't (typically) need to get too complex in your Jekyll files.



Dates in front matter and file system

You may have noticed that both the file name itself and the front matter include date information. The date in the front matter will take precedence over the file name. If you don't need to specify a time and you're happy with the date in the file name, simply remove the date from the front matter.

Go ahead and open the new file and create the front matter and content:

```
---
layout: post
title: "Hello World!"
date: 2016-06-25 09:51:33 -0500
categories: general
---

<p>
This is a test!
</p>
```

The first item, `layout`, tells Jekyll what layout file to use for rendering content. We'll look at layouts in the next section. The `title` value is - obviously - what will be the title for the post. The `date` value represents when you created the content. You can specify whatever value you want here. However - if you specify a time in the future, Jekyll will assume this is content for publication at a later time. (You can tell Jekyll to display posts in the future by using `--future` at the command line or by editing your configuration file.) The final value, `categories`, represents what categories this blog entry should go in. These categories are completely arbitrary. You can look at the values used in the initial post as an example, and going forward, pick ones that make sense for your site. So if you're creating a pet blog, categories could include "dog", "cat", and "dragon". Finally, the actual content of the post is just one paragraph. Save the file, and then reload your Jekyll blog.

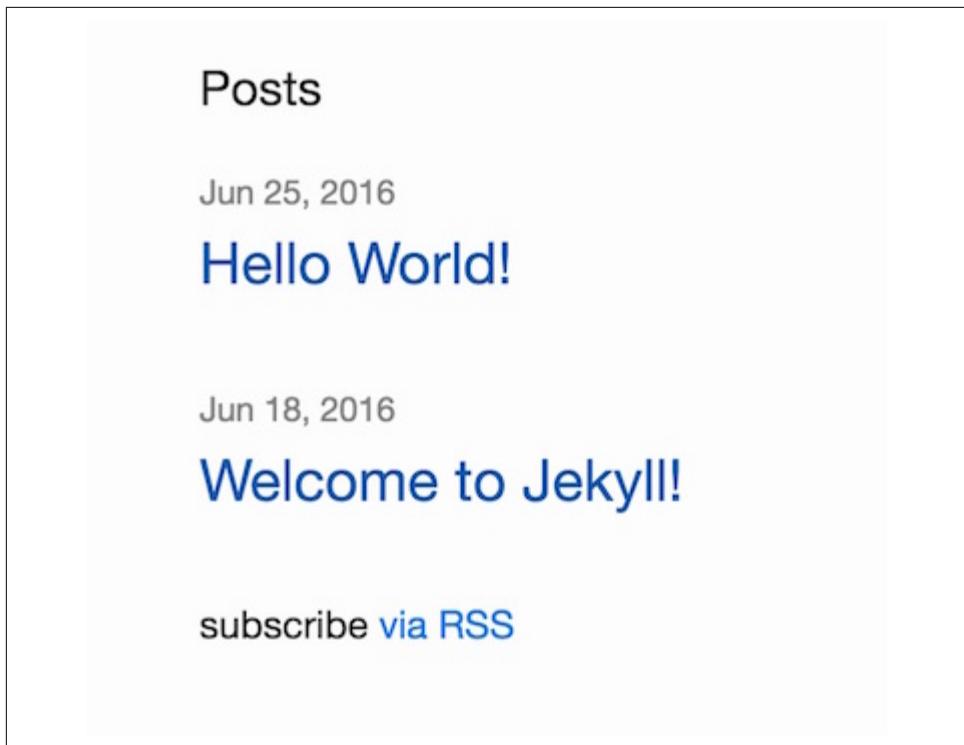


Figure 3-8. Your new post shows up automatically.

As you can see in figure 3-8, Jekyll automatically included the post. Notice how it shows above the previous post - that's because (hopefully) you used a date that is after the initial post. If you modify the date and set an earlier value for the year, it will automatically move down below the first post. Click the link and you'll end up on this URL (and again, the actual date values may be different):

`http://localhost:4000/general/2016/06/25/new-post.html`

There's a few things to note here. First, the category of the blog entry is included in the URL. Personally, I don't think that's good, and you'll see how to modify this when we get to configuration of Jekyll later. The next three items represent the date of the post. This is **not based** on the filename of the post, but what the date was in the YAML. Generally you want these to agree just to make things easier to understand, but they don't need to. The final part of the URL, new-post, comes from the filename. Generally this value is based on the title of the post. So for example, if my post was titled, "My new cat is named Sinatra!", you would use a file name of YEAR-MONTH-DAY-my-new-cat-is-named-sinatra. Essentially this is a file-safe version of the title. It's lowercased and missing the punctuation from the title.

Go ahead and click on the link to open the post. You'll see an example in figure 3-9:



Figure 3-9. Your new post.

The layout of the post was handled by Jekyll. If you modify the front matter to tweak the title and modify the HTML, you can then reload the page and see your changes. If you chose to use Markdown for your extension than you'll notice that Jekyll automatically converted it into HTML as well.

A quick introduction to Liquid

So how exactly did Jekyll create that list of blog entries on the home page? Jekyll includes a dynamic template language called Liquid. Liquid is a powerful templating language that includes the ability to output simple variables, loop over lists, conditionally show content, and include other templates. A full explanation of the language is beyond the scope of this book, but luckily you won't need to be an expert at it to use it with Jekyll. Liquid has its own home page for documentation on GitHub (<https://github.com/Shopify/liquid/wiki>) and you should bookmark it for easy reference while you work. You can also bookmark [Jekyll Cheat Sheet](http://jekyll.tips/jekyll-cheat-sheet/) (<http://jekyll.tips/jekyll-cheat-sheet/>) for Liquid tags specifically used in Jekyll. Let's take a look at the home page as an example of Liquid in work.

```
---  
layout: default  
---
```

```

<div class="home">

  <h1 class="page-heading">Posts</h1>

  <ul class="post-list">
    {% for post in site.posts %}
      <li>
        <span class="post-meta">
          {{ post.date | date: "%b %-d, %Y" }}</span>

        <h2>
          <a class="post-link" href="{{ post.url | prepend: site.baseurl }}">{{ post.title }}</a>
        </h2>
      </li>
    {% endfor %}
  </ul>

  <p class="rss-subscribe">subscribe
    <a href="{{ "/feed.xml" | prepend: site.baseurl }}>via
    RSS</a></p>

</div>

```

There's two things in particular you want to notice. First, the loop block that begins with `{% for ... %}`. Liquid uses `{% ... %}` as a marker for logic (looping and conditionals).

Next notice how the date is displayed: `{{ post.date ... }}`. Liquid uses `{{ ... }}` when doing simple variable replacements. You can see multiple examples of this in the template. Sometimes the values are just displayed as they are: `{{ post.title }}`. Sometimes they are modified: `{{ post.date | date: "%b %-d, %Y" }}`. In the previous example, the pipe character acts like a filter - in this case doing date formatting. Later examples show a `prepend` filter that - as you can guess - put content in front of a variable.

The variables, like `site.posts`, aren't a Liquid feature, but comes from Jekyll giving your template access to the content of your site. If you open up `feed.xml`, you'll see a similar loop used to generate the RSS feed for the site. There are a whole set of variables Jekyll provides to your templates and you can find the complete list at the Jekyll Variables documentation page (<https://jekyllrb.com/docs/variables/>). Let's modify the home page to include one of the values - the post excerpt.

Within `index.html`, add the following line before the closing `` tag:

```

{{ post.excerpt }}

```

Reload the home page and you'll notice that a portion of the blog entry now shows up on the home page.

The screenshot shows a web page with a light gray background and a white content area. At the top left of the content area, the word "Posts" is written in a small, dark font. Below it, the date "Jun 25, 2016" is shown in a smaller font. To its right, the title "Hello World!" is displayed in a larger, bold blue font. Underneath the title, the first paragraph of the post content, "This is a test!", is visible in a smaller black font. Further down, another date "Jun 18, 2016" appears, followed by the title "Welcome to Jekyll!" in blue. Below the titles, a descriptive text block explains how Jekyll handles post excerpts, mentioning the `_posts` directory, re-building the site, and running the `jekyll serve` command.

Jun 25, 2016

Hello World!

This is a test!

Jun 18, 2016

Welcome to Jekyll!!

You'll find this post in your `_posts` directory. Go ahead and edit it and re-build the site to see your changes. You can rebuild the site in many different ways, but the most common way is to run `jekyll serve`, which launches a web server and auto-regenerates your site when a file is updated.

Figure 3-10. Post excerpts.

So how did Jekyll create this excerpt? By default Jekyll grabs the first paragraph of content from your post. You have multiple ways to change this, though. First, in your front matter, you can simply write your own excerpt. Or, you can define a “marker” in your post for where an excerpt should be created. So for example, many blogging engines will let you write `<!--more-->` to define where the excerpt should end.

Working with Layouts and Includes

Now that you've seen how to write a post, and had a quick introduction to the Liquid templating language, let's continue by looking at how Jekyll renders the content around your posts - the layout.

As you've already noticed, the default Jekyll site has a rather simple layout. The expectation is that you'll replace it with your own custom theme. But while the layout is simple, the actual support for working with layouts in Jekyll is really well done. Let's go over the basics.

With Jekyll, layouts are defined in the front matter of every page within your site. If you open up `index.html`, `about.md`, or any of your posts, you'll see a `layout` value defined within the YAML on top of the page. If you do not specify a layout value, then no layout is used.

The value specified in the front matter should *not* include an extension. So while your layout file may be called `profile.html`, you would only use the value “`profile`” when specifying it as your layout.

A layout file includes the contents of the file that uses it by outputting a variable called `content`. Here is an incredibly simple and short example of a layout:

```
<html>
<head></head>
<body>
{{ content }}
</body>
</html>
```

Any file using this layout would be “wrapped” with the HTML above {{ content }} as well as the HTML below it.

Layout files themselves can also call other layout files. This means you can define a core, site wide layout and then a specific layout for your blog posts. You can see an example of this yourself in the default site produced by the Jekyll command line. Open up post.html in the _layouts folder and you’ll see that it specifies default as the layout for itself. That means a post specifying post for its layout would first be wrapped by post.html and then default.html.

Finally - all of your layout files should be included in the _layouts folder. You get this folder by default when you create a new Jekyll site.

Of course, sometimes you may want to re-use content across a site in a non “wrapped” format. You may simply have some text or boilerplate HTML that you need to use in multiple locations. To support this feature, Jekyll includes a include command that will simply insert the contents of another file into the one currently being rendered. You can see an example of this in the default layout:

```
<!DOCTYPE html>
<html>

{% include head.html %}

<body>

{% include header.html %}

<div class="page-content">
<div class="wrapper">
{{ content }}
</div>
</div>

{% include footer.html %}

</body>

</html>
```

In the layout file above, you can see two includes being used - one for header.html and one footer.html. Notice that you *do* need to specify a file extension but you do *not* need to specify a folder. Jekyll will automatically look for these within the _includes

folder. These includes can be dynamic as well and make use of Liquid. This comes in handy especially for things like your page title.

If you view the About page, notice that the title in your browser tab is “About”. If you view source, you’ll see it in the code as well: `<title>About</title>`. Where did this come from? Two places actually. First, notice the front matter for about.md includes a title value: `title: About`. This gets picked up by Jekyll and becomes a variable that can be used by Liquid. If you open up `_includes/head.html`, you’ll see this in action:

```
<title>{%
  if page.title
    {{ page.title | escape }}%
  else %
    {{ site.title | escape }}%
  endif %}</title>
```

The `site.title` variable is also something that can be tweaked. We’ll discuss it later in the Configuration section.

Finally, you can also call includes and pass arguments to them. As an example:

```
{% include youtube-video.html id="4nx7g60Ldig" %}
```

In your include file, the variable `id` is referenced via the `include` scope, as demonstrated below:

```
<a href="https://www.youtube.com/watch?v={{ include.id }}">Watch video</a>
```

Hopefully the benefits of layouts and includes are already known to you, but in case they aren’t, it’s helpful to remind ourselves why this is such a powerful feature. In most sites, the layout consists of a particular set of tags that drives the look and feel for what you see. Everything from the structure (menu on the top, or menu on the right) to the design (dark green header with blue text) is consistent across the site giving every page a uniform look and feel. Moving to static web sites without an engine like Jekyll driving it means having to repeat this code in every file, greatly increasing the amount of work required to update the design or make changes across the site. By using layouts, and includes, you’re saving time (and money) implementing your site.

Adding Additional Files

While Jekyll is generally focused on creating blogs and blog posts, your site will most likely have additional pages that aren’t posts. When you created your initial Jekyll site you saw an example of this, the About page. There is nothing special about this file outside of the fact that it isn’t within the `_posts` folder. By simply including any HTML file, or Markdown file, in the root of the Jekyll folder, they will automatically be parsed for their Liquid tags, wrapped in layouts, and available in the final static version of your site.

While you have an example of this already, go ahead and create a new file in the root of the `demo1` folder called `cats.html`. The following listing shows the contents of this file.

```

---
layout: default
title: Cats
---

<div class="home">
  <h1 class="page-heading">Cats</h1>

  <p>
    If my blog were a collection of cats instead of posts, I'd have {{ site.posts.size }} cats.
  </p>
</div>

```

While this is a rather trivial example, it shows how you can create content in a Jekyll site totally unrelated to the blog itself. This isn't a post, just supporting content. And as you can see, it can contain Liquid tags. In this case, we've asked for the total number of posts. Since the file is in the root of the Jekyll folder, it will be found in the root of the domain, in this case, <http://localhost:4000/cat.html>. You're free to create any subdirectories you want of course, but obviously want to avoid folders Jekyll uses for special purposes, like `_layouts`.

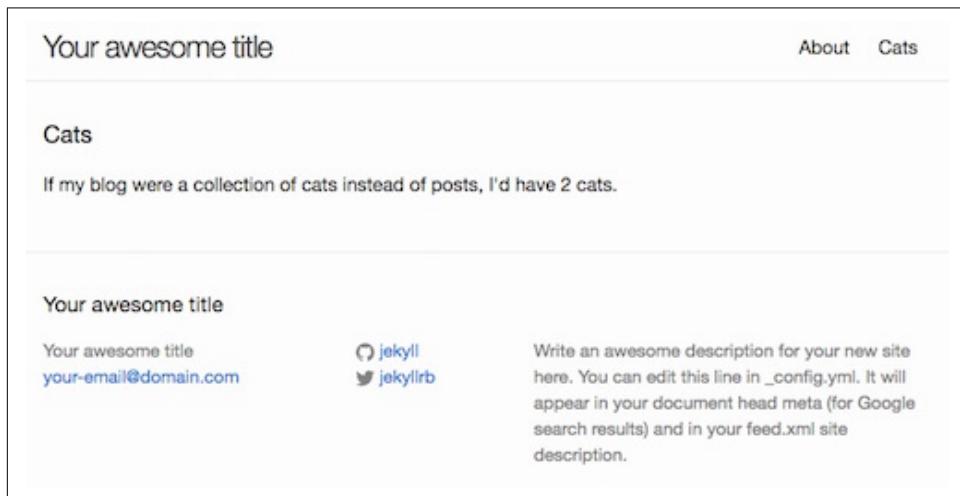


Figure 3-11. The new cat page.

Notice how the header automatically added a link to the new page? You can actually find the code behind that logic in the include `header.html`. It is included via this logic:

```

{% for my_page in site.pages %}
  {% if my_page.title %}
    <a class="page-link"

```

```
    href="{{ my_page.url | prepend: site.baseurl }}">>{{ my_page.title }}</a>
  {% endif %}
  {% endfor %}
```

This just shows again an example of the type of data Jekyll provides to your templates. In this case, you can actually ask Jekyll for a list of known pages. Since you just created a new page, Jekyll knows about it and your code can work with it.

Finally, if you don't care for the URL used by Jekyll for the new page and you don't want to rename the file itself, you can specify the link to use in your front matter by adding a `permalink` value. This will override the URL Jekyll uses that's based on the filesystem.

Working with Data

Now you know the basics of creating a Jekyll site, creating content and pages, and how layouts work. It's time to kick it up a notch and work with some more powerful features. The first one we'll tackle is data files. Data files are a way to provide data to a Jekyll site in an abstract fashion. So what exactly does that mean?

Imagine for a moment you have a simple page representing your board of directors. Now imagine that board of directors are all cats. Because - why not? While you could build this out manually, you realize that you may want to get this list of directors and use it elsewhere. Or you may want to provide multiple versions of this. Perhaps a list of only left-handed members versus those who use the right hand. Jekyll provides a way for you to define a generic "data set". This data can consist of anything you can possibly imagine. Once defined, and you'll see how that's done in a minute, you can then access this data in your blog posts, pages, layouts, and includes and use it as you see fit.

To begin, your data files should live in a folder called `_data`. This folder is *not* created by Jekyll by default, so simply create it yourself. (If you've downloaded the GitHub version of the code samples, everything in this section may be found in the `demo2` folder.)

Inside that folder, you'll then create a file that represents your data. Jekyll will use the filename as a way to reference the data, so you want to ensure you use a sensible name. If your data is a list of cats, then your filename should be `cats.something`, where the extension defines how the data is represented. Jekyll's data feature supports YAML, JSON, or CSV files.

For this first demo, create a file called `cats.json`. This will represent a data set of cats. You can either use the code downloaded from the GitHub repo for the book or copy the following. Feel free to edit the names and other properties to your liking. (In fact, we've trimmed the data set a bit in the book listing!)

```
[  
  {  
    "name": "Elvis",  
    "age": 4,  
    "gender": "male",  
    "picture": "http://placekitten.com/200/300"  
  },  
  {  
    "name": "Sinatra",  
    "age": 8,  
    "gender": "male",  
    "picture": "http://placekitten.com/400/330"  
  }  
]
```

In the code sample above you can see two cats with four properties. Just to be clear - this is 100% arbitrary. Jekyll doesn't care what kind of data you store and how you store it. It simply makes it available to your site. The nature of the data (and what defines a cat for example) is completely up to you.

Now that you have this data defined, you can start using it in your templates. One simple example would be to list all the cats:

```
---  
layout: page  
title: The Board of Cats  
---  
  
<ul>  
  {% for cat in site.data.cats %}  
    <li>{{cat.name}}</li>  
  {% endfor %}  
</ul>
```

At top you see the normal YAML front matter for a Jekyll page. Under this is an unordered list. To iterate over the cats, we “address” it using `site.data.cats`. Again - Jekyll uses the file name to create the handle you use in your code. If the file name was `furrycreaturesofevil.json`, then the reference in our pages would be `site.data.furrycreaturesofevil`. For each cat, we simply output the name. Much like how Jekyll doesn't care what data you create, it doesn't care how you use it either. You're welcome to include data you may not need, or may use in the future. The result is pretty much as you expect - a list of cats.

The Board of Cats

- Elvis
- Sinatra
- Lilith
- Pig
- Robin

Your awesome title

Your awesome title
your-email@domain.com



Write an awesome description for your new site here. You can edit this line in _config.yml. It will appear in your document head meta (for Google search results) and in your feed.xml site description.

Figure 3-12. The list of cats - a cat list.

Of course, you can mix it up a bit as well. This example, male_cats.html, shows basic filtering of the data set.

```
---  
layout: page  
title: Male Cats  
---  
  
<ul>  
  {% for cat in site.data.cats %}  
    {% if cat.gender == 'male' %}  
      <li>{{cat.name}}</li>  
    {% endif %}  
  {% endfor %}  
</ul>
```

And here is a slightly more complex (cats_with_pics.html) example showing the entire data set being used:

```
---  
layout: page  
title: Cats with Pictures  
---
```

```
{% for cat in site.data.cats %}  
  <p>  
    <br/>  
    {{cat.name}} is a {{cat.gender}} cat and is {{cat.age}} years old.  
  </p>  
{% endfor %}
```

The result is absolutely lovely. In case you're curious, those pictures come from [Place-Kitten.com](#), a placeholder image service comprised completely of cat pictures.

Your awesome title ≡

Cats with Pictures



Elvis is a male cat and is 4 years old.



Sinatra is a male cat and is 8 years old.

Figure 3-13. This is an even better list of cats.

Configuring your Jekyll Site

One aspect we haven't dealt with yet is configuring your Jekyll site. Jekyll can be configured at the command line via flags and by editing the `_config.yml` file. Which you use is obviously up to you and what your particular site needs. Most likely if you simply want to test a particular change it would make sense to use a command line flag. If you are sure you want Jekyll to act a particular way going forward, then use the config file.

There are quite a few different options for configuration so developers should check the documentation (<https://jekyllrb.com/docs/configuration/>), but let's consider a few simple examples of things you may want to tweak right away.



Updates to Configuration

Note that if you decide to play with a local Jekyll site while reading, and you're highly encouraged to do so, Jekyll will not notice changes to your configuration automatically. You will need to stop and restart the server to see them reflected.)

`title`: Your awesome `title` is used not by Jekyll but by the default layout. If you open the `include/header.html`, you'll notice this: `{{ site.title }}` The variable, `site.title`, comes directly from the configuration file. You can add your own variable to the configuration file and then access it via the `site` object. So this is an example of a configuration value that doesn't necessarily drive how Jekyll operates, but is used within your layouts as a sort of global variable. You can see the same thing with the `email`, `twitter_username`, and `github_username` values. You can add, or remove these, as you see fit.

`port` is used locally when testing your Jekyll server. You won't see this in your default `_config.yml`. If you don't like the default port (4000), or perhaps you want to run multiple Jekyll servers at once, you can tweak the port. You can either specify it in your config file or pass it via the command line: `jekyll server --port=4001`

`permalink` is used by Jekyll to determine how URLs are created for posts. While you can configure this on a per-post basis, most likely you'll want to set this up one time for the entire site. By default, Jekyll uses a permalink that includes the category of the post itself. Looking at the site used in the `demo1` folder, the URLs for a post look like this:

```
http://localhost:4000/general/2016/06/25/new-post.html  
http://localhost:4000/jekyll/update/2016/06/18/welcome-to-jekyll.html
```

In the first URL, only one category was specified (general) and in the second post two were used (jekyll and update).

While this may be fine, most sites use a URL that includes the dates only and not the categories. Jekyll lets you change this by using the `permalink` setting. Jekyll supports defining permalinks by using various tokens. The default looks like this:

```
/:categories/:year/:month/:day/:title.html
```

Each token is prefixed with a colon and represents data from the information about the post. Notice the last token, `:title`, will be automatically formatted so it is safe for a URL. So to remove the category values from the links, you can simply edit your `_config.yml` file to include the new permalink value:

```
permalink: /:year/:month/:day/:title
```

You can see the complete list of tokens at Jekyll's Permalink documentation (<https://jekyllrb.com/docs/permalinks/>). While this is easy to tweak, you'll probably want to decide which style you prefer before you launch your site.

Generating a Site

You may have noticed that Jekyll constantly rebuilt your site as you worked on it. It stored the static version in the subdirectory `_site`. If you want to generate a static version manually, simply run `jekyll build`. By default it will use the current directory as the source and `_site` as the destination, but both can be changed via command line arguments. Here's an example: `jekyll --source ../blog --destination ../output`

Here's an example of the output from the `demo3` site. Note that the posts each end up in unique folder based on the date. This is because of the changes made to the `_config.yml` file in the permalink value. If you compare this to the output from `demo1`, the folder structure is different.

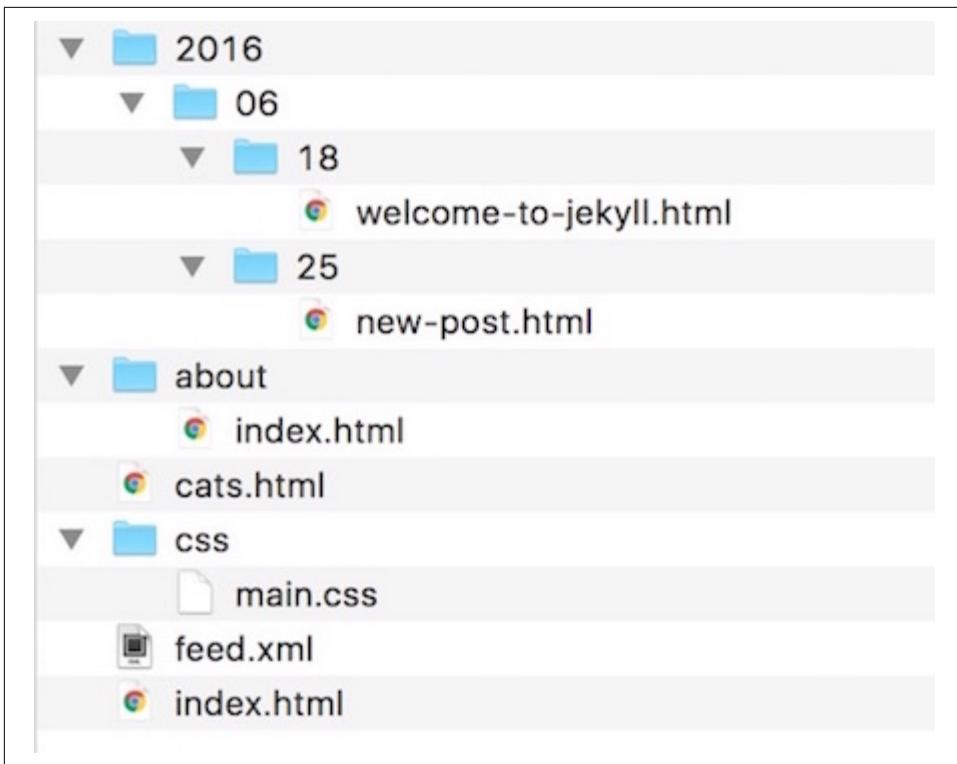


Figure 3-14. Static output from Jekyll.

Building a Blog

Since Jekyll caters heavily to blogs, just creating a new Jekyll site gives you a blog out of the box. In theory, you're done and we can wrap the chapter up right here. Obviously though you probably won't go live with Jekyll's default look and feel for your site. For our blog, we're going to discuss how you can find a theme, customize it, and make it your own.

There's quite a few open source, and commerical, blog templates out there. Turns out - many of them come with Jekyll versions as well. One such theme, Clean Blog (<http://startbootstrap.com/template-overviews/clean-blog/>), looks great and includes both a 'regular' version and a Jekyll template. Go ahead and download the Jekyll version (from the GitHub repository at <https://github.com/BlackrockDigital/startbootstrap-clean-blog-jekyll>), unzip, and try running it from the command line. Remember to stop any other Jekyll server you have up.

The first time you try to run the Jekyll site, you might encounter an error:

```
+ catblog git:(master) ✘ jekyll s
Configuration file: /Users/raymondcamden/Dropbox/Writing/ora static sites/ch3/catblog/_config.yml
  Dependency Error: Yikes! It looks like you don't have jekyll-feed or one of its dependencies installed. In order to use Jekyll as currently configured, you'll need to install this gem. The full error message from Ruby is: 'cannot load such file -- jekyll-feed' If you run into trouble, you can find helpful resources at http://jekyllrb.com/help/!
jekyll 3.1.3 | Error:  jekyll-feed
+ catblog git:(master) ✘
```

Figure 3-15. An error from the template.

What you're seeing here is the fact that the Jekyll template makes use of a "gem", a Ruby package, as part of it's plugin functionality. This is described in more detail in the docs (<https://jekyllrb.com/docs/plugins/>), but for now all you need to know is how to install the code necessary for the template. Luckily this is pretty simple: `bundle install`.

```
+ catblog git:(master) ✘ bundle install
Fetching gem metadata from https://rubygems.org/......
Fetching version metadata from https://rubygems.org/..
Fetching dependency metadata from https://rubygems.org/.
Resolving dependencies...
Rubygems 2.0.14.1 is not threadsafe, so your gems will be installed one at a time. Upgrade to Rubygems 2.1.0 or higher to enable parallel gem installation.
Using public_suffix 2.0.3
Installing colorator 1.1.0
Installing ffi 1.9.14 with native extensions
Using forwardable-extended 2.6.0
Using sass 3.4.22
Installing rb-fsevent 0.9.8
Installing kramdown 1.12.0
Using liquid 3.0.6
Using mercenary 0.3.6
Installing rouge 1.11.1
Using safe_yaml 1.0.4
Using jekyll-paginate 1.1.0
Using bundler 1.13.6
Installing addressable 2.5.0
Using rb-inotify 0.9.7
Using pathutil 0.14.0
Using jekyll-sass-converter 1.4.0
Installing listen 3.0.8
Using jekyll-watch 1.5.0
Using jekyll 3.3.0
Using jekyll-feed 0.8.0
Bundle complete! 3 Gemfile dependencies, 21 gems now installed.
Use bundle show [gemname] to see where a bundled gem is installed.
+ catblog git:(master) ✘
```

Figure 3-16. Adding the Gem required for the template.

If you get an error that `bundle` isn't a valid command, then first install it: `gem install bundler`

Once done, the blog should start up fine. But there's another hitch you may not even notice. If you try to view the blog at `http://localhost:4000`, where you viewed all the previous examples, you'll notice that you get an error instead. If you go back to your command line prompt, you'll see why:

```

➔ startbootstrap-clean-blog-jekyll-master jekyll s
Configuration file: /Users/raymondcamden/Downloads/startbootstrap-clean-blog-jekyll-master/_config.yml
  Source: /Users/raymondcamden/Downloads/startbootstrap-clean-blog-jekyll-master
  Destination: /Users/raymondcamden/Downloads/startbootstrap-clean-blog-jekyll-master/_site
  Incremental build: disabled. Enable with --incremental
    Generating...
      done in 0.311 seconds.
  Auto-regeneration: enabled for '/Users/raymondcamden/Downloads/startbootstrap-clean-blog-jekyll-master'
Configuration file: /Users/raymondcamden/Downloads/startbootstrap-clean-blog-jekyll-master/_config.yml
  Server address: http://127.0.0.1:4000/startbootstrap-clean-blog-jekyll/ ↗
  Server running... press ctrl-c to stop.

```

Figure 3-17. See the non-standard path?

This template uses a feature of Jekyll that lets you define a particular subdirectory for a blog. This would be handy when a blog is a subset of your entire site. While handy, this is something we'll tweak. When you enter the correct URL, you'll see your new blog in all its glory.



Figure 3-18. The blog template in action

So at this point we have one main job. We need to modify this blog so it matches our design requirements. For example, we can replace that header image with something more appropriate. The title, too, should be changed. If you scroll to the bottom, you'll also notice various social media icons that have to change as well.

What you change, of course, depends on what you feel is necessary for your blog. You may find a template that is, out of the box, perfect. For now, let's focus on changing a few things to make it more appropriate for our subject matter - cats.

Let's begin by looking at `_config.yml`.

```

# Site settings
title: Clean Blog
header-img: img/home-bg.jpg
email: your-email@yourdomain.com
copyright_name: Your/Project/Corporate Name
description: "Write your site description here. It will be used as your sites meta description as
baseurl: "/startbootstrap-clean-blog-jekyll"
url: "http://yourdomain.com"
twitter_username: SBootstrap
github_username: davidtmiller
facebook_username: IronSummitMedia
email_username: your-email@yourdomain.com

# Google Analytics
# To enable google analytics, uncomment below line with a valid Google Tracking ID
# google_tracking_id:

# Build settings
markdown: kramdown
highlighter: rouge
permalink: pretty
paginate: 5
exclude: ["less","node_modules","Gruntfile.js","package.json","README.md"]

gems: [jekyll-paginate, jekyll-feed]

```

From top to bottom, let's make these changes:

- Change the title to “The Cat Blog”
- Leave header-img alone. We *will* be changing it, but in this case the configuration is specifying a file name. We can replace it later.
- Change email to your own email.
- Change description to “A blog about cats!”
- Completely remove baseurl.
- Change url if you want, but it need not be a real site.
- Change twitter_username to your own, or keep it, but we'll keep this social link.
- But then remove github_username and facebook_username. Again, we're kind of making some arbitrary decisions here. Feel free to *not* remove them and change them to new values if you want.
- Finally, set email_username to the same value you used for email

You can find the final version of `_config.yml` in the code repository for the book, but go ahead and kill your Jekyll server, restart it, and see your changes.

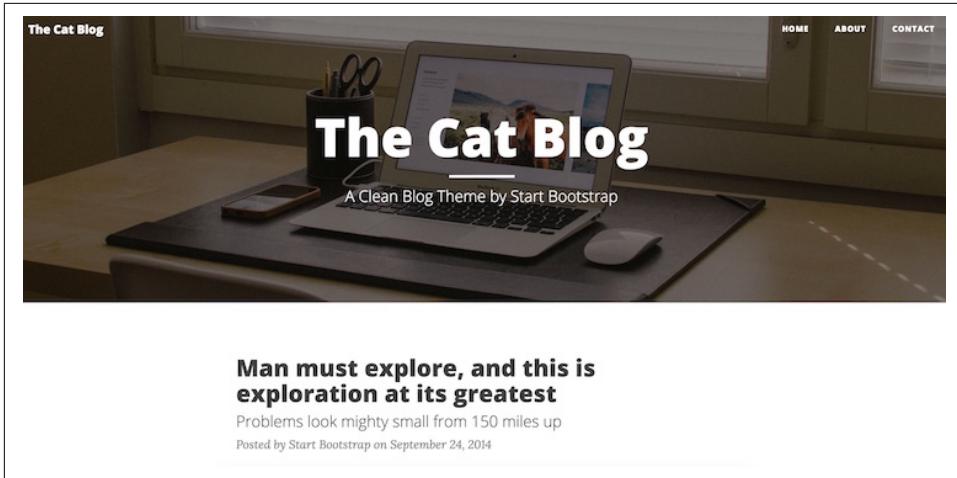


Figure 3-19. Our cat blog is getting there!

Now let's fix that header image. While you can find any image you want, we'll make use [Placekitten.com](#). The config file specified the header file as `img/home-bg.jpg`. If you open that up and check the size, you'll see it is 1900 pixels wide and 872 pixels high. Using Placekitten, you can generate an image of the same size by going to this URL: <https://placekitten.com/1900/872>. On that page, simply download the graphic and save it over the existing file name. Now when you reload, you'll see your new header.

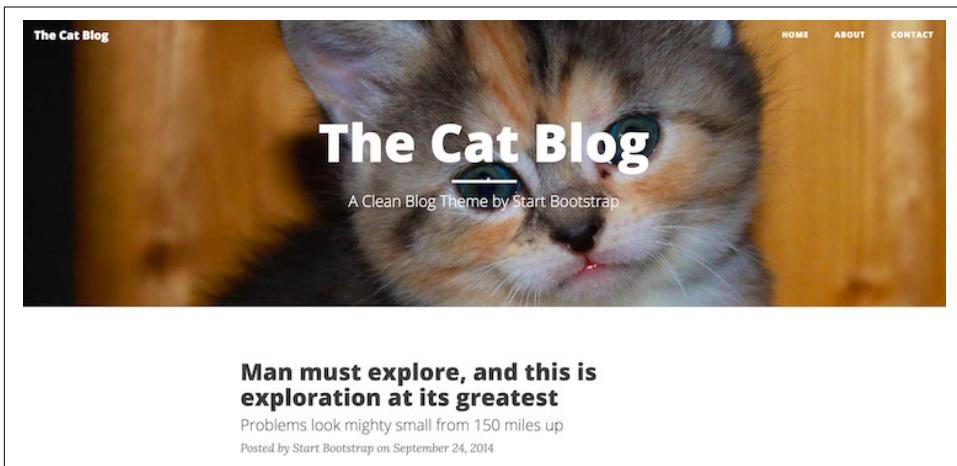


Figure 3-20. That's a beautiful cat header!

Now let's change the text on the home page. You can see that the title loaded from configuration, but what about the tagline beneath it, "A Clean Blog Theme by Start

Bootstrap”? If you open up `index.html`, you’ll see that this is actually defined in the front matter:

```
---
```

```
layout: page
```

```
description: "A Clean Blog Theme by Start Bootstrap"
```

```
--
```

Go ahead and change that description to: “A Blog about Cats” - reload - and you’ll see the change. Now let’s do a few more tweaks. You may have noticed two links in the top of the site - “About” and “Contact”. We can’t support “Contact” forms yet (you’ll see how later in the book) so just delete the file (`contact.html`) from the folder. Reload, and it’s gone from the header. Like you saw earlier, Jekyll is actually driving this based on the files in the folder. Just removing the file is enough to remove it from the navigation.

Now open up `about.html` and modify the “Lorem ipsum” text. It doesn’t matter what you type, but just go ahead and change it so you can see it in action. Like the home page, this file has a `description` field you can change as well. Notice that this page also points to a header file (`img/about-bg.jpg`). If you want, you can change this as well. The image I used was from Placekitten again (<https://placekitten.com/1200/600>). Here’s the final version of the About page from the book repository.

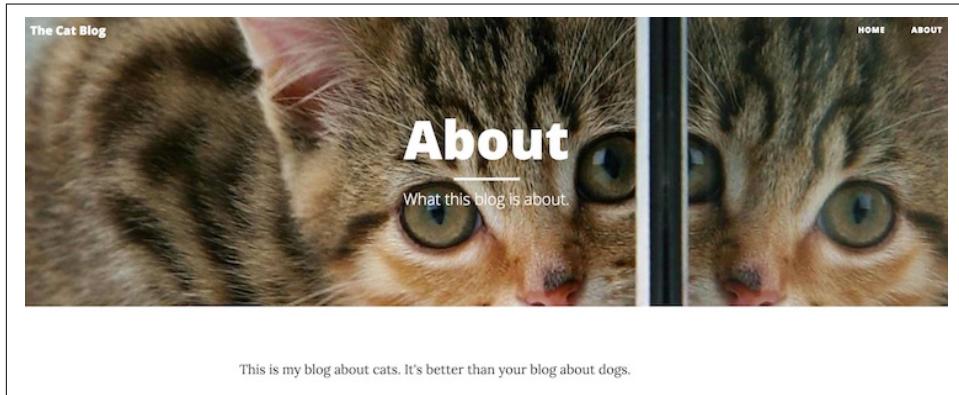


Figure 3-21. The Cat-flavored About page.

And with that - we’re done with the site layout and related pages. Once again, “done” is a relative term here. You can continue to tweak the layout to your heart’s content. Now let’s handle the blog content. Right now there are six files in the `_posts` directory. Go ahead and delete all but one file. Take that last file, and rename it to match the current date and a name that’s more appropriate. For example: `2016-07-23-welcome-to-my-blog.markdown`. You’ll notice the template used Markdown files for its posts. If you want to change the post to HTML, simply rename the extension as well. Open up the file, and let’s update the front matter and text.

```
---
```

```
layout:      post
title:       "Welcome to my blog!"
subtitle:    "Cats are worth blogging about!"
date:        2016-07-23 12:00:00
author:      "Raymond Camden"
header-img:  "img/post-bg-06.jpg"
---
```

Everything there should be pretty standard, but note that the time isn't being used by the template so it doesn't matter what value you set there. For the text, just type what you want. I used text from Cat Ipsum (<http://www.catipsum.com>), a "Lorem Ipsum" generator with a cat flavor. Lastly, I generated another custom cat picture from Placekitten with this URL: <https://placekitten.com/888/400>

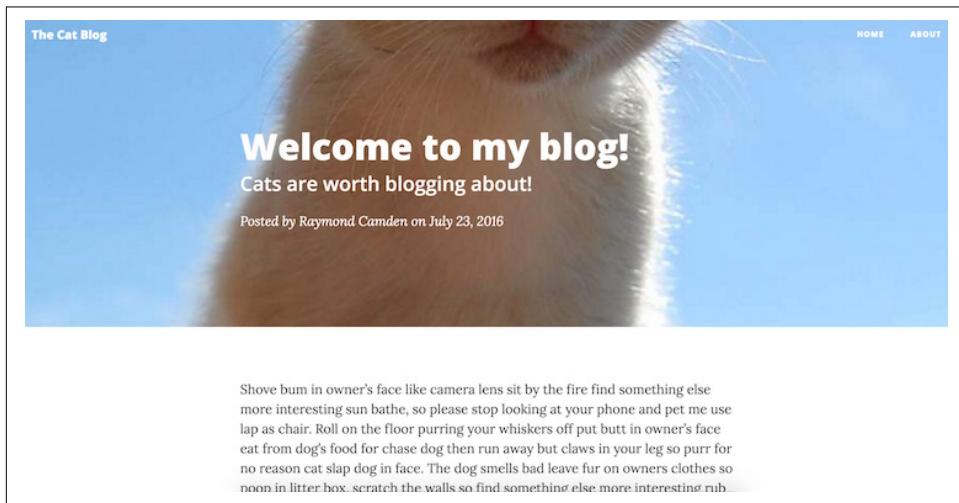


Figure 3-22. Our first blog post.

And that's it! Again, you could easily customize this even more, but now you're ready to launch your own cat blog. (And to be clear, this is highly recommended.)

Going Further with Jekyll

Liked what you saw in Jekyll? Here's a quick tour of some of the things we didn't cover but may excite you even more.

- Probably one of the most powerful features is automatic pagination. Jekyll can handle correctly "sizing" the page as well as creating all of the additional files to handle showing each page of content. You also get an easy way to support adding the "Previous" and "Next" links in whatever fashion you want. You can see this in

action if you add 5 more posts to the Cat blog project created earlier. It's set to show 5 posts per page.

- Jekyll also has a powerful plugin system (<https://jekyllrb.com/docs/plugins/>) for adding additional functionality to your site. This can include generators, which make new files, new “tokens” you can use in your templates, and even extensions to the CLI. It's very diverse and powerful. Speaking from my own experience, I had to build a plugin once and even though I didn't know Ruby, it wasn't difficult to add.
- Jekyll has a concept called “Collections” (<https://jekyllrb.com/docs/collections/>), which lets you define your own set of content that can act like posts and pages. This is a power feature you may not need for a simple blog, but it's great that Jekyll includes this functionality.
- Finally - while we're going to talk about deploying your static sites later in the book, Jekyll is used by GitHub and their GitHub Pages (<https://pages.github.com/>) feature. If you want to host static content for your GitHub project, you can easily use Jekyll along with it. If your Jekyll site uses plugins, be sure to check GitHub's docs as they only support a subset of plugins.

For more information, see the Jekyll documentation: <https://jekyllrb.com/docs/home/>. You can also find support at the official Jekyll forums: <https://talk.jekyllrb.com>.

CHAPTER 4

Building a Documentation Site

While there are many different types of documentation sites — from hardware or software documentation to project or policy documentation — most of these sites share some common characteristics.

Characteristics of a Documentation Site

First, they tend to have multiple, and frequently numerous, contributors. In the case of project or policy documentation, for instance, these may be exclusively internal contributors within a given company (though, depending on the size of the company, they may be external to the team that builds and maintains the documentation site). However, in a software world increasingly dominated by open source, many documentation sites have to deal with a large number of external contributors.

Second, a typical documentation site is fairly simple and straightforward in terms of features and design. Most of the time the layout and design of documentation is intentionally simple and geared towards readability over style. Outside of things like comments or maybe runnable example, a documentation site rarely includes any complex, dynamic functionality.

Third, most documentation sites change relatively infrequently. Oftentimes a documentation site receives periodic significant updates with occasional minor updates in between.

While none of these characteristics is a requirement for choosing a static site, they do enable documentation sites to be in a position to take the most advantage of the benefits of using a static site generator.

All that being said, there are some drawbacks to choosing a static site generator for building a documentation site. It generally requires that a developer or development

team be involved in the creation of the site and often in the regular build and deployment of updates. Also, there is no built-in authoring interface that would be comfortable for a non-technical author/contributor. This can be a challenge if your documentation writers are used to authoring in WYSIWYG interfaces rather than markup languages like Markdown.

Overcoming the Drawbacks

It's worth noting that there are a number of ways to overcome the challenges listed here for documentation sites. For instance, the build and deployment process could be handled via services that make this a simple one step process. Also, the editor could be improved by leveraging one of the available static site content management systems (CMS) available. We'll discuss these options later [chapter 6](#).

Choosing a Generator for Your Documentation Site

Pretty much any static site generator will work for creating a documentation site. In fact, many documentation sites already use Jekyll, including well-known projects like [PhoneGap](#) and [Kendo UI](#). Many smaller software projects choose Jekyll primarily because of the built-in integration into [GitHub Pages](#). In this chapter, however, we will use [Hugo](#), a Go-based static site generator that has been rapidly growing in popularity recently.

Why choose Hugo?

The primary reason I would recommend using Hugo for documentation sites is that the build process is extremely fast. While this can be insignificant for smaller sites, many documentation sites can grow very large and unwieldy, making the build process a bottleneck in updating and deploying the site.

The secondary reason is that Hugo does not come preconfigured for running a blog - a format that is not suitable for most documentation sites. This makes it easier to set up Hugo for whatever site structure that you need for your documentation site without needing to restructure the default setup. In fact, Hugo even offers a number of [community themes](#) that are designed specifically for building a documentation site.

For the remainder of this chapter, we'll look at how to build a basic documentation site using Hugo.

Building a Doc Site with Hexo

Hexo is a JavaScript-based static site generator available on npm. As an alternative to building a documentation site in Hugo, Bruno Mota has [written a tutorial on Site-Point](#) showing how to use Hexo for this purpose.

Our Sample Documentation Site

We're going to build a documentation site for an esoteric programming language called [LOLCode](#).



Figure 4-1. The LOLCode home page

[lolcode_homepage.png](#)

While LOLCode is designed to be a humorous language, with an intentionally difficult syntax to comprehend, it does have a full [language specification](#). Right now, that language spec is a simple and boring plain text on GitHub, shown below.

LOLCODE Specification 1.2

FINAL DRAFT — 12 July 2007

The goal of this specification is to act as a baseline for all following LOLCODE specifications. As such, some traditionally expected language features may appear "incomplete." This is most likely deliberate, as it will be easier to add to the language than to change and introduce further incompatibilities.

Formatting

Whitespace

- Spaces are used to demarcate tokens in the language, although some keyword constructs may include spaces.
- Multiple spaces and tabs are treated as single spaces and are otherwise irrelevant.
- Indentation is irrelevant.
- A command starts at the beginning of a line and a newline indicates the end of a command, except in special cases.
- A newline will be Carriage Return (/13), a Line Feed (/10) or both (/13/10) depending on the implementing system. This is only in regards to LOLCODE code itself, and does not indicate how these should be treated in strings or files during execution.
- Multiple commands can be put on a single line if they are separated by a comma (,). In this case, the comma acts as a virtual newline or a soft-command-break.
- Multiple lines can be combined into a single command by including three periods (...) or the unicode ellipsis character (u2026) at the end of the line. This causes the contents of the next line to be evaluated as if it were on the same line.
- Lines with line continuation can be strung together, many in a row, to allow a single command to stretch over more than

Figure 4-2. The existing LOLCode language specification

lolcode_lang_spec.png

We're going to borrow that language spec and spice it up a little bit. Using a real spec will allow us to build an example that meets a "real world" use case, even if it is still relatively simple enough to remain useful as an example.

Now, those of you that know me, know that I have no design skill whatsoever (though that may still be more than Ray). To overcome this obstacle, we'll use a freely available design for a basic documentation site called [DocWeb](#). Our completed site will look like the image below.

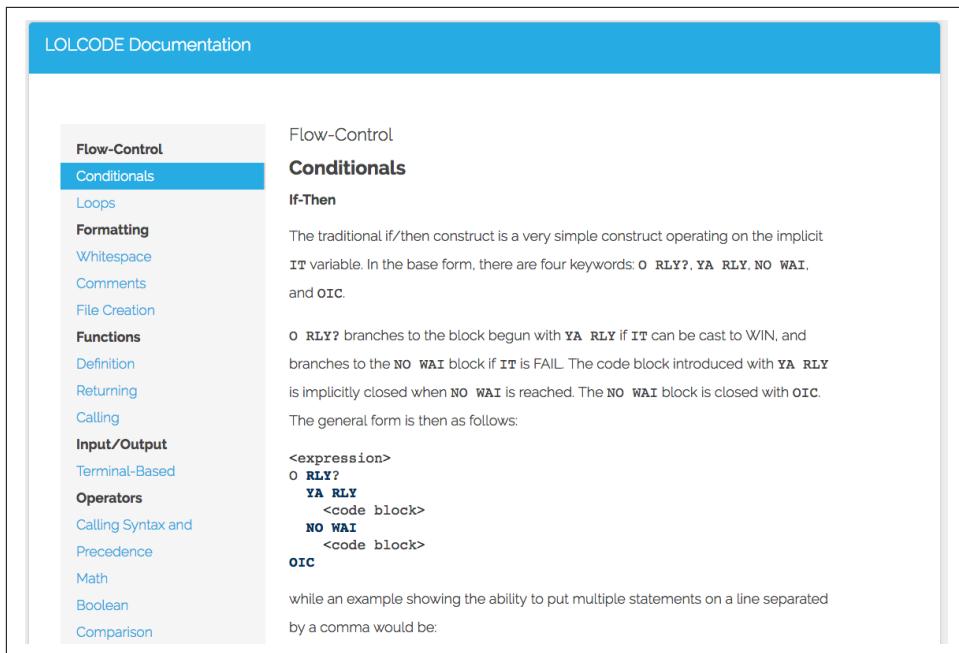


Figure 4-3. The LOLCode language spec site built with Hugo using the DocWeb template

lolcode_hugo_site.png

The site itself is actually a single page with the navigation on the left being anchor links that will scroll to the appropriate section. The navigation will remain visible as the user scrolls through the long page, and the highlighted item will automatically change depending on which section the user has scrolled to.

This form of single page documentation is common for things like language API docs or even getting started guides. However, even though this is a single page, it is comprised of many separate pieces of content, as we'll see shortly. In fact, if you had a larger, more complex project to document, where a single page may be unsuitable, it'd be relatively easy to convert this example to use a page per content item rather than a single page for the entire project.

Pre-built Hugo Themes

Although we'll be customizing our own design using the DocWeb template, Hugo has a large number of pre-built themes that you can use, all of which can be found at themes.gohugo.io. There are even a few specifically designed for documentation sites such as [Bootsie Docs](#) and [Material Docs](#).

Creating the Site

Now that we've reviewed what we are building, let's start building it!

Installing Hugo

If you are on a Mac, the simplest way to install Hugo is via [Homebrew](#). If you have Homebrew installed, enter the following into a Terminal window.

```
brew update && brew install hugo
```

On Windows, you'll need to download the [latest release](#) as a zip file. (Note that the latest release as of this writing was v0.15). Hugo is simply an executable file contained within the downloaded zip. However, to use it properly, you'll need to add it to your PATH variable.

Start by creating a folder at C:\hugo\bin. Copy the executable into this folder (you'll want to rename it to simply hugo.exe).

Assuming that you are on Windows 10...

1. Start by double-clicking the start button and choosing "system".
2. Next click on the "Advanced system settings" option.
3. Press the "Environment Variables..." button at the bottom of the window.
4. Choose the "Path" variable and click the "Edit" button.
5. Press the "Nw" button and enter "C:\hugo\bin".
6. Click "OK", then "OK" again and finally "OK" one last time to close all the windows.



Help with Installation

If you aren't on Windows 10 or prefer not to use Homebrew, the Hugo documentation has you covered. For additional Windows options, see the "[Installing on Windows](#)" page. Likewise, for additional Mac options, see the "[Installing on Mac](#)" page.

To test that your installation is configured correctly, open the Command Prompt/Terminal and enter `hugo help`. You should receive a long list of commands and options similar to what is seen in the image below.

```
[brinaldimac:~ Rinaldi$ hugo help
hugo is the main command, used to build your Hugo site.

Hugo is a Fast and Flexible Static Site Generator
built with love by spf13 and friends in Go.

Complete documentation is available at http://gohugo.io/.

Usage:
  hugo [flags]
  hugo [command]

Available Commands:
  server      A high performance webserver
  version     Print the version number of Hugo
  config      Print the site configuration
  check       Check content in the source directory
  benchmark   Benchmark hugo by building a site a number of times.
  convert     Convert your content to different formats
  new         Create new content for your site
  list        Listing out various types of content
  undraft    Undraft changes the content's draft status from 'True' to 'False'
  import     Import your site from others.
  gen         A collection of several useful generators.
```

Figure 4-4. You can verify your install by using the help command.

[hugo_help.png](#)

Generating the Initial Site Files

Now that we've confirmed that we have Hugo installed, we can generate the initial site files using the command line utility (CLI). Start by opening a Command Prompt/Terminal in the folder where you would like your project to reside.

To create the folder structure and files needed to start our new Hugo site enter the following command (where "docsite" is the name of the folder that will be created):

```
hugo new site docsite
```

Hugo will generate the base files and directories necessary for a new site. Inside our "docsite" directory, we should see the following files and folders.

▶	archetypes	Today, 11:28 AM	--	Folder
▶	config.toml	Today, 11:28 AM	107 bytes	TOML file
▶	content	Today, 11:28 AM	--	Folder
▶	data	Today, 11:28 AM	--	Folder
▶	layouts	Today, 11:28 AM	--	Folder
▶	static	Today, 11:28 AM	--	Folder

Figure 4-5. The default files and directories generated by Hugo.

At the moment, all of the directories are empty. Unlike many other static site generators, Hugo does not generate a default theme/layout or any dummy content. Let's start with configuring the site.

Configuring the Hugo Site

By default, Hugo uses [TOML](#) for configuration, though you can choose to use [YAML](#) if you prefer. The site configuration is located in the config.toml file within the root directory of the site. It contains only a few default values.

```
baseurl = "http://replace-this-with-your-hugo-site.com/"
languageCode = "en-us"
title = "My New Hugo Site"
```

These values should be fairly self-explanatory. Because we don't have a real URL, we'll just set the baseurl value to the root and port of the web server that Hugo will use for local testing.

```
baseurl = "http://localhost:1313"
languageCode = "en-us"
title = "LOLCODE Documentation"
```

One important aspect of a documentation site for a language like LOLCode is code highlighting. There are multiple options for code highlighting in a Hugo site, but the standard solution is to do server-side code highlighting (i.e. the highlighting is added during the build phase) with [Pygments](#).

To get Pygments working, we need to have Python installed. If you don't already have it, you can [download the proper version here](#) (if you are on Windows, I suggest that you choose the option to add Python to your PATH during installation).

Once Python is installed, open a new Command Prompt/Terminal window and install Pygments.

```
pip install Pygments
```

If you'd like to test that Pygments installed correctly, enter `pygmentize -h` via the command line. This should print out the Pygments command-line help documentation.

Now that Pygments is installed, let's configure Hugo to use it. We'll add the following to our config.toml file.

```
pygmentsStyle = "colorful"
pygmentsUseClasses = false
pygmentsCodeFences = true
```

The first option chooses one of Pygments default styles for the code coloring. The second tells Pygments to add the code colors directly rather than depend on the pygments.css file. The third will allow us to use [GitHub style code fences](#) in our Mark-

down to set code blocks and the language to use for highlighting. A GitHub style code block would look like this:

```
```javascript
// my JavaScript code here
```
```

Hugo offers more options for code syntax highlighting than we're discussing here. For details on those, visit the [Hugo syntax highlighting documentation](#).



Client Side Code Coloring

One thing to consider, depending on the size of your site, is that the server-side code coloring shown here via Pygments can significantly impact build times. If this might be an issue for your site, you might consider using one of the [client-side code coloring options](#) like Highlight.js or Prism.

We're done with what we need in the configuration file. Of course, there are a ton of other options available to us when building a site with Hugo. You can view the full list of configuration options in the [Hugo configuration documentation](#). For now, however, let's move on to adding content and creating a layout.

Adding Content

One of the benefits of Hugo, in my opinion, is that it is not very prescriptive about how you organize or name your content. Content is typically in the /content folder, but beyond that, it's up to you how you handle things from there. In our case, our final content will actually all be rendered on to a single page, so the organization isn't terribly important.

As with most static site generators, Hugo natively supports content written in Markdown. However, it also supports two other markup languages: [AsciiDoc](#) (via [AsciiDoctor](#)) and [reStructuredText](#). For our documentation site, we'll be using Markdown.

Rather than have us write or convert all the LOLCode documentation here, I've provided a zip file containing all of the Markdown files. Simply [download it](#) and unzip it within the /content folder of your Hugo site. All of the Markdown files should end up within a directory of /content/lolcode.

Converting to Markdown

Rather than rewrite all of the LOLCode documentation to the Markdown format, I chose to use a converter. Thankfully there are a number of converters available. To convert HTML to Markdown, I often use [to-markdown](#). To convert rich text to Mark-

down, I'll use [Mark It Down](#). Finally, to convert Word documents to Markdown, I'll use [Word to Markdown Converter](#).

It's important to note, however, that in the case of this last converter, you'll need to do some manual cleanup both before (I usually remove images as they get Base64 encoded) and after (some headers and other formatting don't translate perfectly). Nonetheless, if you need to convert any document to Markdown, any of these tools will save you significant time.

As is typical for static site generators, Hugo uses front matter at the beginning of each content document. Essentially, front matter is metadata about that particular piece of content, including everything from the title and the publish date, to the URL and categories. By default, front matter is written in TOML, but YAML and JSON are also allowed if you prefer.

The easiest way to create a new post and ensure that it has the necessary front matter is to use the command line. Assuming that you are in the root of your site, you can simply enter the following command to create a new post named "welcome to my blog."

```
hugo new welcome-to-my-blog.md
```

The generated post will have front matter that looks like the following:

```
+++
date = "2016-06-07T19:30:29-04:00"
draft = true
title = "welcome to my blog"

+++
```

The `date` is a timestamp that will be used to determine the publish date of the post. The `draft` key indicates that this post is not yet published; you can set this to false or remove it to take the post out of draft status. Finally, the `title` is, obviously, the title of the blog post or content.

Since we've already added our content, we don't need to create a new content item (so delete the post if you created it). But let's take a look at the metadata in our documentation content, as it includes some items not in the default front matter.

```
+++
date = "2016-05-05T08:41:21-04:00"
draft = true
title = "Arrays"
categories = ["Types"]
categories_weight = 6
+++
```

We've already discussed the `date`, `draft` and `title` values (as you can see, this is the documentation covering arrays in LOLCode). However, we've added some taxonomy via the `categories` value. This value can contain an array of categories, although, in this case, we've only added one for "Types". As you'll see when we create the layout, we're going to use this `categories` value to group our documentation (in this case, an Array is a data type, thus the category "Types").

The `categories_weight` value also helps us to organize the way we will display the content on the page. In this case, we want certain content to come before others within the "Types" category. The value here indicates that we want "Arrays" to be the sixth item in the order of the content within the "Types" category.

There are a lot of additional front matter metadata values that you can add to your Hugo content. Refer to the [documentation](#) for full details.



Required Front Matter in Hugo

There's some discrepancy between the [Hugo documentation on front matter](#) and the reality. The documentation indicates that title, date, description and taxonomies are all required. However, as we discussed above, Hugo itself doesn't include either description or taxonomies when it generates content for you and I've never found Hugo to error when either of these keys are missing.

Creating the Layout

So our site is configured and it has content, but if we were to ask Hugo to serve it at this point and open it in a browser, we'd still get nothing. Why? Well, there is no layout to tell Hugo how to display the content.

As we discussed earlier, we aren't going to use a pre-built theme in this case, but rather will create our own layout using [DocWeb](#) as a template. To start, we need to download that template and place the zip in a temporary location.

Once you've unzipped the DocWeb layout, let's start moving some of the assets over to our LOLCode documentation site. First, let's copy the `css`, `img` and `js` folders into `/static/assets` within our Hugo site. Placing files in the `static` directory will copy them over to the generated site without processing them - essentially leaving them as is.

Next, copy `index.html` from the DocWeb files and place it under the `/layouts` directory within our Hugo site. Technically, at this point, we have a home page layout (and that will be the only page for this site), but it's all in a single HTML file and the content within Hugo isn't populating. Let's fix that.



Go Templates

By default, Hugo uses the Go Template language for templates, though it also natively supports [Amber](#) and [Ace](#) templates. We'll cover some of the basics of creating Go Templates here, but if you want a full overview, the Hugo documentation provides a [great introduction](#) or you can review the official [Go Template documentation](#).

We'll start by creating some "partials," which are essentially includes. These allow us to break the template into reusable parts, making them easier to maintain.

Let's create a new folder under `/layouts` called "partials" and create a file named `head.html`. Then take the `<head>` portion of `index.html` and copy it into the new `head.html` document. Back in `index.html`, place the following code where the `<head>` used to be.

```
 {{ partial "head.html" . }}
```

Within `layouts/partials/head.html`, let's add some dynamic site data. First, let's clean the head up a bit by removing the `<link>` for RSS (of course, you can make an RSS file for your site using Hugo, we just won't for this demo as it is a fixed set of documentation rather than a site with regular updates) and the stylesheet for `prettify.css` (`Prettify` is a code highlighting tool and we've already configured Hugo to handle our code highlighting).

Next, let's replace the existing hardcoded site title with the site title we defined earlier in the site configuration file.

```
<title>{{ .Site.Title }}</title>
```

Finally, let's modify the link to the stylesheet to reference the base URL of the site to ensure that the reference URL is always correct.

```
<link rel="stylesheet" href="{{ .Site.BaseURL }}assets/css/style.css">
```

Both `Title` and `BaseURL` are site variables that Hugo makes available to use within our templates. You can find a full list of them in the [Hugo documentation](#).

Your completed `head.html` partial should look something like this:

```
<head>
  <title>{{ .Site.Title }}</title>
  <meta charset="utf-8">
  <meta name="description" content="">
  <meta name="HandheldFriendly" content="True">
  <meta name="MobileOptimized" content="320">
  <meta name="viewport" content="initial-scale=1.0, minimum-scale=1.0, maximum-scale=1.0, user-scalable=no">
  <link href="http://fonts.googleapis.com/css?family=Raleway:700,300" rel="stylesheet" type="text/css">
  <link rel="stylesheet" href="{{ .Site.BaseURL }}assets/css/style.css">
</head>
```

It's a good idea to split out portions of your site into partials such as this one, which can be easily reused within additional templates across your site. For example, we could also cut the footer portion of the template and place it into a layout/partials/footer.html file.

```
<section class="vibrant centered">
  <div class="">
    <h4> This documentation template is provided by <a href="http://www.frittt.com" target="_blank">
  </div>
</section>
<footer>
  <div class="">
    <p> © Copyright LOLCODE. All Rights Reserved.</p>
  </div>
</footer>
```

In place of the footer above, we'll add `{{ partial "footer.html" }}` within layouts/index.html.

Before we move on, let's fix one last thing in our layouts/index.html file. At the very bottom of the page are some script tags. Let's add the `.Site.BaseURL` to the ones we need (i.e. jquery.min.js and layout.js) and remove the two tags for prettyify.js (again, we're not utilizing this for code coloring). Our scripts should look like this now.

```
<script src="{{ .Site.BaseURL }}assets/js/jquery.min.js"></script>
<script src="{{ .Site.BaseURL }}assets/js/layout.js"></script>
```

Once you've done this, you'll also want to open static/js/layout.js and remove the call to `prettyPrint()`; at the bottom of the file to avoid generating JavaScript errors.

Let's create one last partial for our template to encompass the site navigation. This one will be a little more complex as we'll want Hugo to dynamically populate the navigation based upon our content.

First, copy out the portion of HTML code in layouts/index.html, starting with the opening `<ul class="docs-nav">` and ending with the closing `` tag, and place it into a new file named layouts/partials/nav.html. Make sure to add a `{{ partial "nav.html" . }}` where the `` block used to be in layouts/index.html.

As you can see by examining the code for the existing, hard-coded navigation items, there are two levels of navigation elements. There are section headers that are bolded, but not clickable. For example:

```
<li>Getting Started</li>
```

And navigation elements which are clickable. For example:

```
<li><a href="#welcome" class="cc-active">Welcome</a></li>
```

If you recall, each of the content items we added had a category defined under categories in the front matter, as well as a `categories_weight`. We'll use the categories as

the section headers. Then we'll place each content item within that category as a clickable navigation item, the order being determined by the `categories_weight`.

Start by deleting all but the first section header (i.e. "Getting Started") and first navigation item (i.e. "Welcome"). Now, let's loop through all of the available categories defined by our site content. Hugo has a site variable called `.Site.Taxonomies.categories` that contains a structure (technically call a `map` in the Go language) of all of the site's categories. To loop through this structure, we're going to use the `range` keyword. Here's what the loop looks like.

```
<ul class="docs-nav">
  {{ range $name, $taxonomy := .Site.Taxonomies.categories }}
    <li><strong>{{ $name | title }}</strong></li>
      <li><a href="#welcome" class="cc-active">Welcome</a></li>
    {{ end }}
  </ul>
```

In the above code, `$name` is the key value from the map (`.Site.Taxonomies.categories`) and `$taxonomy` is the value, which, in this case, will be another map containing all the content within that category. Within the section header list element, we are outputting the value of the key.

The use of the pipe is something special within Go templates. In this case, we are saying to use a

[Hugo template function](#) to title case the value of the `$name` variable. Pipes can be utilized in more complex fashions, which you can learn more about in the [Hugo documentation](#).

Now let's loop through and create the links to the pages within each category.

```
<ul class="docs-nav">
  {{ range $name, $taxonomy := .Site.Taxonomies.categories }}
    <li><strong>{{ $name | title }}</strong></li>
    {{ range $taxonomy.Pages }}
      {{ if ne .Title "" }}
        <li><a href="#{{ .LinkTitle | urlize }}" class="cc-active">{{ .Title }}</a></li>
      {{ end }}
    {{ end }}
    {{ end }}
  </ul>
```

The `$taxonomy` variable makes a number of useful information available to output on our page. In this case, we are interested in iterating over `.Pages`, which has all of the pages within each given category. Within that loop, we are making sure the value of any given page's title is not empty (i.e. `if ne .Title ""` where `ne` indicates "not equal"). This is because there is one piece of content that we want to output on the page (`about_types.md`) that we do not want included in the navigation, so we've left the title off.

Since our final template will be a single page, each navigation item links to an anchor rather than a separate page. Thus, we've used another Hugo template function, `urlize`, to convert any spaces in the page variable `.LinkTitle` to dashes. Finally, within the link, we output the title of the content via `.Title`.

Now, let's try to run our site to see what we've done so far. Make sure that you've saved all files, including `nav.html`. Open a Terminal/Command Prompt window within the folder of the documentation site and enter:

```
hugo server --buildDrafts
```

The `--buildDrafts` flag is necessary as all of our posts are still in draft mode (i.e. `draft = true` in the front matter). It tells Hugo to build all our posts, including any in draft mode.



Changing the Draft Status

You can manually take a post out of draft simply by changing the draft flag in the front matter. However, you can also do so via the command line. For example, the following command will take the arrays post out of draft status:

```
hugo undraft content/lolcode/arrays.md
```

Opening `http://localhost:1313` (the default server URL for Hugo) in your browser, we should now see a site that looks like this:

The screenshot shows a Hugo documentation site with a sidebar and a main content area. The sidebar on the left contains a tree view of content categories: Flow-Control, Conditionals, Loops, Formatting, Whitespace, Comments, File Creation, Functions, Definition, Returning, Calling, Input/Output, Terminal-Based, Operators, and Calling Syntax and. The main content area has a header "Free Documentation HTML Template | Responsive". Below the header, there is a "Getting Started" section with a "Welcome" heading. The "Welcome" section contains the text: "Are you listening to your customers? As they say: You cannot improve what you cannot measure; but the paradox is you cannot measure everything – happiness, hatred, anger.. but you can measure customer satisfaction. Yes, you can measure customer satisfaction by analyzing likes and dislikes of your customers. You can gauge popularity of your website or products. You can also:

- See how many visitors like the new design of your website or logo
- Analyze what your readers want to see on your blog
- Understand how helpful the content on your support forum or website is
- Know the latest trends and user's opinion before launching a new product or service

".

Figure 4-6. Our documentation site with dynamic navigation added

[docsite_with_nav.PNG](#)

Our navigation is now outputting correctly, but the links don't lead anywhere because we still only have hard-coded content. Let's fix that (but, go ahead and leave the site open in your browser, Hugo will automatically refresh it as you save changes).

Open layouts/index.html. Within the `<div class="docs-content">` block, erase all but the first few lines. You should be left with the following:

```
<div class="docs-content">
  <h2> Getting Started</h2>
  <h3 if="welcome"> Welcome</h3>
  <p> Are you listening to your customers?</p>
</div>
```

We'll do the same loop here as we did for the navigation, but, in this case, we'll actually output the full content (i.e. `.Content`) of each post. The only other difference here is that we did not include the `if` statement that skipped content without a title, as we now want that content outputted.

```
<div class="docs-content">
{{ range $name, $taxonomy := .Site.Taxonomies.categories }}
  <h2>{{ $name | title }}</h2>
  {{ range $taxonomy.Pages }}
    <h3 id="{{ .LinkTitle | urlize }}">{{ .Title }}</h3>
    {{ .Content }}
    <hr>
  {{ end }}
{{ end }}
</div>
```

Once you save this page, Hugo will automatically regenerate the site files and refresh the open browser window with your finished site.

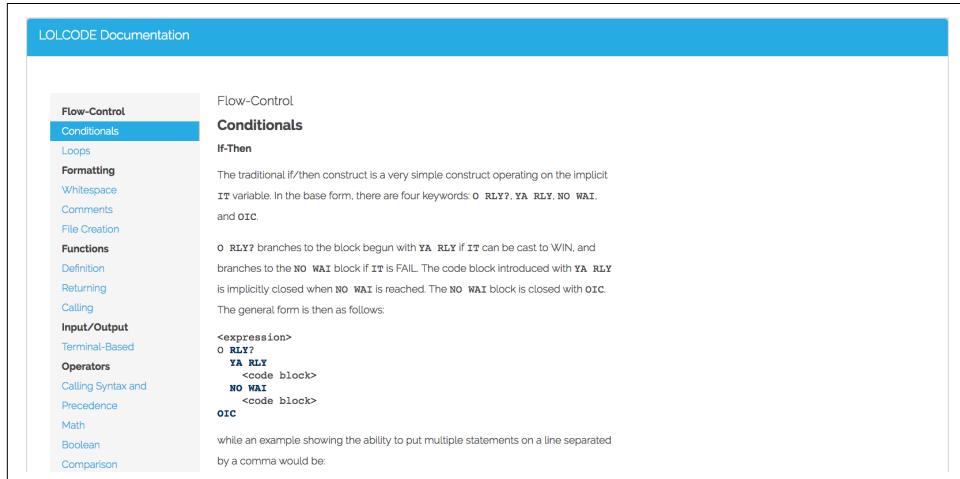


Figure 4-7. The completed LOLCode documentation site.

lolcode_finished.png

Clicking on any of the links in the navigation should now correctly scroll to the appropriate section of the page. You should also notice that code blocks are appropriately colored.

Congrats! You're done with your first documentation site built with Hugo. The completed code for this sample can be found on the [GitHub repository for this book](#).



The Scrolling Navigation

If you've tested the page we built, you may notice that some of the navigation items fall off the page and aren't visible until the entire page is scrolled to the bottom. This has to do with some JavaScript code in `static/js/layout.js` that is designed to have the navigation remain statically positioned as the user scrolls the page. Obviously, this is a bit of a flaw in the design of the DocWeb template when utilized for longer single-page documentation sites. The completed demo of this mitigates this issue through some [tweaks to the JavaScript](#), though it can still be improved, in particular to account for the speed of the user's scroll.

To generate the completed site files, use the command line (note that we're still using the `buildDrafts` flag since our content is still in draft mode):

```
hugo --buildDrafts
```

Going Further

Obviously, there are many different formats you can use for a documentation site. In this particular sample project, the documentation was made up of primarily short content, which made it a very good candidate for a single page with anchors. However, if your project has long form documentation, it would make more sense to use individual content pages. To do this, you'd need to add a template for the individual posts as `layouts/_default/single.html`. Of course, you could reuse the partials for the head, navigation and footer, making it easier to build this template quickly.

One thing worth noting here for anyone looking to convert an existing documentation project to Hugo is that there are a [number of converters](#) available. These include converters from popular static site generators like Jekyll but also from dynamic site engines and content management systems like Wordpress, Ghost and Drupal.

Honestly, the toughest challenge you are likely to run into when choosing to use Hugo, or any static site generator for your documentation site, is writing content in Markdown. While developers may be comfortable writing in Markdown, most documentation writers will be unfamiliar with it. Fortunately, the support for Markdown

in existing or standalone tools keeps getting more robust, even if it is typically not the WYSIWYG experience many writers will be most comfortable in. Also, it's worth noting that there are a number of third-party **front ends** for Hugo that aim to create an experience more similar to writing in Wordpress or other content management systems.

Finally, there's also a lot more that you can do with Hugo than was covered in this chapter. Thankfully, Hugo has extensive and **detailed documentation** (among the best of any static site generator, in my own opinion).

Adding Dynamic Elements

Now that you've made the decision to use a static site generator (wooh!), you may discover that you've forgotten one little detail that seems to *require* a dynamic application server of some sort. Perhaps you have a contact form. Maybe you want to add discussions to your blog. Does this mean you've made a horrible mistake and have to revert everything you worked on? Of course not!

Multiple different services exist that enable you to add dynamic features back to your site. In some cases, it is near impossible for the end user to even tell that anything "special" is being done at all. To them, it simply appears as if your site is doing whatever it needs to on the back end, when in reality you're still using simple static files.

In this chapter, we'll discuss how to add dynamic elements back to your static site. We'll focus on:

- Forms
- Comments
- Search

We'll also discuss some other options for different types of data, like events and generic sets of data.

Handling Forms

Since the dawn of the web itself, forms have been an integral part of how users interact with your site. Moving to a static site doesn't make forms themselves magically stop working. They still render. You can still interact with them using JavaScript. But if you want the form to do something on a server, or to communicate with someone else, then you'll going to be out of a luck when using a static site. Luckily there's a

number of options you have available to add basic form processing back to your site. Let's look at a few options.

WuFoo Forms

Let's get one thing out of the way first. WuFoo (<http://www.wufoo.com>) is the most expensive of all the options we'll be covering in this chapter. However, it has an *incredible* powerful form builder which can be very easy for non-professionals to use. Because of how it works, your clients can even edit the forms themselves without having to go through you or deploy a new version of the static site.



Figure 5-1. WuFoo's web site

At the time this book was written, WuFoo had a free tier that covered 3 forms and 100 responses. That's probably enough for a pretty small site, but most site will probably want to use the lowest commercial option, "Ad Hoc", which covers 10 forms and 500 entries. You can find higher tiers and current pricing here: <http://www.wufoo.com/pricing/1/>.

Go ahead and sign up for a free account, and in the form manager, click to create a new form.

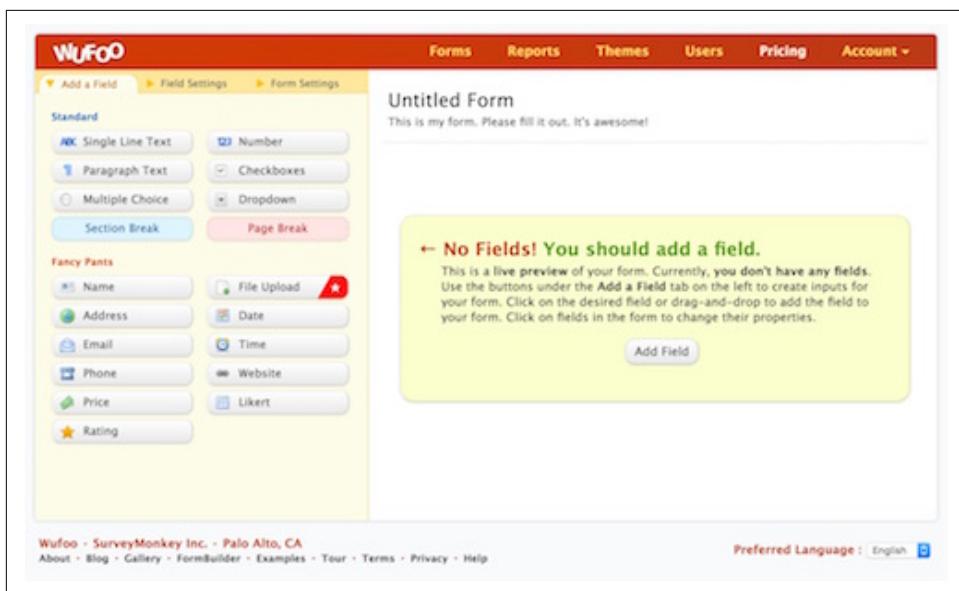


Figure 5-2. Beginning to create a new form

Next, simply drag and drop fields to create your form. Along with basic fields (text, multiple choice, etc), Wufoo has shortcuts for common things like name's (including first and last name fields) and email addresses. Here's a simple contact form created by dragging the Name, Email, and Paragraph Text fields.

Untitled Form

This is my form. Please fill it out. It's awesome!

Name

First Last

Email

Untitled

Figure 5-3. The first version of the contact form

Notice though that the form is kind of generic. Each part of the form can be edited to provide customized text. Simply click a field or block of text, and a new editor will appear:

Figure 5-4. Editing the top portion of the form

Here's the form after customizing the top portion and the paragraph text field:

Contact Form

Please fill out my contact form.

Name

First Last

Email

Your comments

Figure 5-5. The customized form

Once done, save the form, and Wufoo will offer to provide you a method to share it. You can test the form via a custom URL or embed it (either via JavaScript, iframe, or even WordPress). For now, take the JavaScript embed and click the “Copy Code” option.

Embed with JavaScript Embed with iFrame Embed on WordPress

Use this embed code to place your form into other pages.
This is the **recommended** way to embed your form.

```
<div id="wufoo-m1e6s3i81d2ymck">
Fill out my <a href="https://raymondcamden.wufoo.com/forms/m1e6s3i81d2ymck">online form</a>.
</div>
<div id="wuf-adv" style="font-family:inherit;font-size:small;color:#a7a7a7;text-align:center;display:block;">HTML Forms powered by <a href="https://raymondcamden.wufoo.com">Wufoo</a>
</div>
```

Copy Code

Figure 5-6. The embed code for your new form

Now that you’ve got the embed code, you can simply copy and paste it into any simple web page. In the source code for this book, you can find an example of the form

just created in the ch5/forms folder, named wufoo1.html. You are welcome to try running the file as is, but since this particular example if using the free tier, you won't be able to see any of the responses. You're encouraged to replace the embed with your own code from the Wufoo form editor. In order to make it easier, the file uses a comment to make it clear:

```
<!-- WuFoo embed here -->  
...  
<!-- End WuFoo embed -->
```

The screenshot shows a live contact form titled "Testing WuFoo". The form is titled "Contact Form" and has a sub-instruction "Please fill out my contact form.". It contains fields for "Name" (split into "First" and "Last" inputs), "Email" (a single input field), and "Your comments" (a large text area). A "Submit" button is at the bottom. Below the form, a credit line reads "HTML Forms powered by [Wufoo](#)".

Figure 5-7. The live form.

Unless you specifically said so while editing the form, nothing will be required. Also, when you submit the form, you'll end up on the Wufoo web site. You can change this back in your form settings.

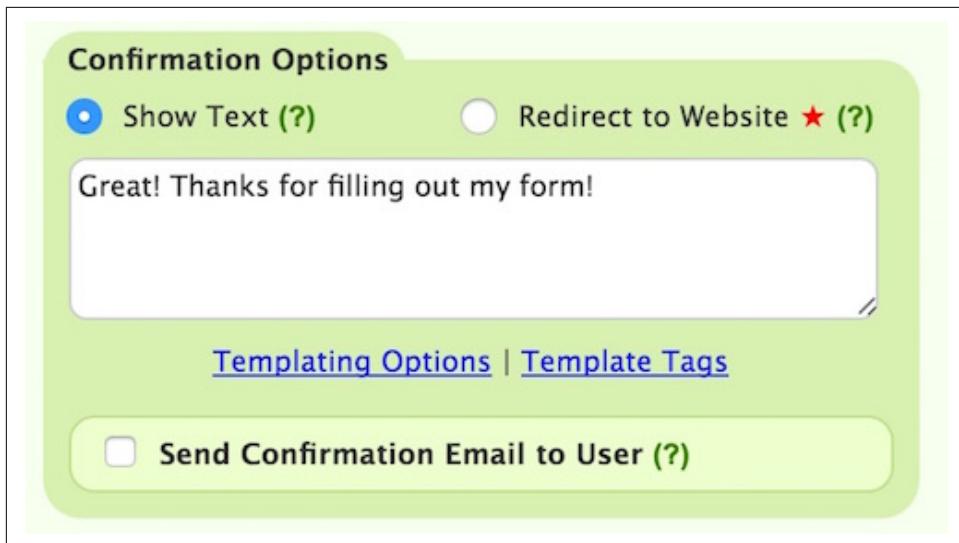


Figure 5-8. Form submission settings

The most likely option you will use here is the “Redirect” one to send users back to your site.

When submitted, results will be available in the form manager. You can select your form, click on the entries, and get a tabular view of the results. Selecting an entry will give you both the values from the submission as well as some other metadata about the form.

The screenshot shows the Wufoo interface for viewing a form entry. At the top, there's a navigation bar with links for Forms, Reports, Themes, Users, Pricing, and Account. Below the navigation is a contact form with fields for Name (Raymond Camden), Email (raymondcamden@gmail.com), and Your comments (Testing). To the right of the form are several action buttons: New (green), Edit (yellow), Email (blue), Print (grey), and Delete (red). A message box indicates "Viewing Entry #2". On the left, a sidebar displays the creation date (10 Sep 2016, 10:44:57 AM) and IP address (76.72.12.163). On the right, there's a comment section with a text input field, a dropdown for "Your Name" (set to raymondcamden), and a "Save Comment" button. At the bottom, a table lists recent entries, showing two entries: Raymond (Last: Camden, Email: raymondcamden@gmail.com) and another entry (Last: , Email:).

Figure 5-9. Details of the form submission

You'll probably also want to get the form submission sent to you via email. Back on the form manager, click the Edit link to add notifications. On this page you can set up both email notifications and even a SMS message!

The screenshot shows the Wufoo notification settings page. It features three main sections: "to My Inbox" (with fields for Email Address and Set Reply To), "to My Mobile Device" (with fields for Cell Phone Number and Carrier), and "to Another Application" (represented by a large green circle with a white plus sign). Each section has a "Save" button at the bottom. Above the sections, there's a yellow bar with the text "Send Notifications from Contact Form ...". A small link "Subscribe to RSS Feeds" is located in the top right corner.

Figure 5-10. Notification settings for WuFoo

There's a lot more to Wufoo than we'll cover here, and as we said in the beginning, this is not the cheapest option, but its power and ease of use are easily worth the price (for most cases).

Google Docs Forms

Google Docs has become a ubiquitous way for people to work on documents (as well as spreadsheets and slide decks) online and in a multiuser environment. Over the past few years it has become a serious competitor to office suites like Microsoft Office and Apple's iWork. One feature many people may not know about is the ability to design a form. Even better - it's free service you can use for your static site. Assuming you've got a Google account (and don't we all), simply open your browser to <https://docs.google.com/forms/>.

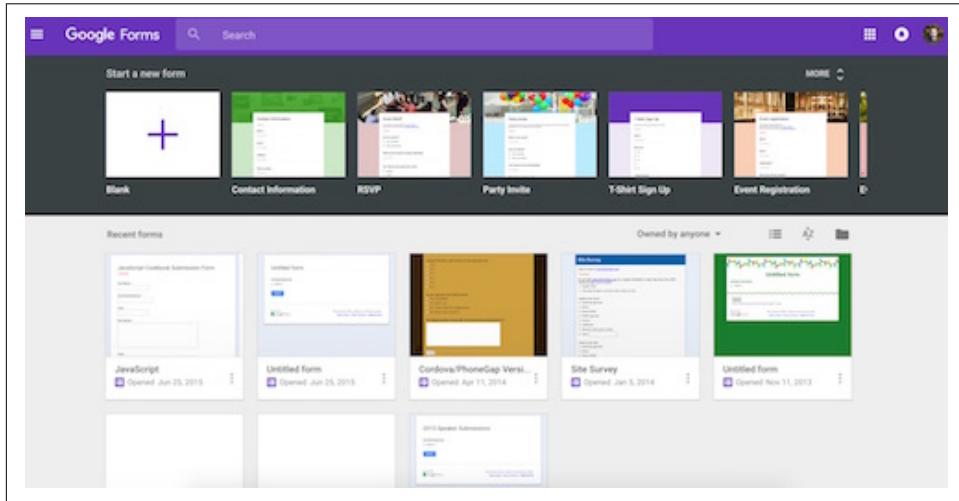


Figure 5-11. The Google Forms web site

You'll see a number of templates, including one for a contact form, but go ahead and create a blank one so you can see the process of creating your own form.

The screenshot shows a blank Google Form interface. At the top, there are tabs for 'QUESTIONS' and 'RESPONSES'. The main area is titled 'Untitled form'. Below the title is a section for a question, with the placeholder text 'Untitled Question'. To the right of the question text is a dropdown menu showing 'Multiple choice'. Underneath the question text, there are two radio button options: 'Option 1' and 'Add option or ADD "OTHER"'. At the bottom of the question section are three small icons: a square, a trash can, and a 'Required' toggle switch which is turned on. To the right of the question section is a vertical toolbar with several icons: a plus sign, a 'Tr' icon, a camera icon, a video camera icon, and a grid icon.

Figure 5-12. A blank Google Form

By default you are automatically editing the first question. You can change the text, question type, and other properties, by using the various options on the page itself. Here's the form after the first field has been changed to a short answer, required field with a new label.

This screenshot shows the same Google Form as Figure 5-12, but with changes made to the first question. The question text 'Untitled Question' has been replaced by 'Name'. The question type has been changed to 'Short answer', as indicated by the dropdown menu. The 'Required' toggle switch is still turned on. The rest of the interface remains the same, including the sidebar with its various icons.

Figure 5-13. Beginning to create the contact form

You can add more fields by using the menu to the right of the form. Note that validation options are a bit hidden. You need to click the three dots to the right of “Required” and select data validation. Validating an email field requires selecting “Text” and then “Email address”.

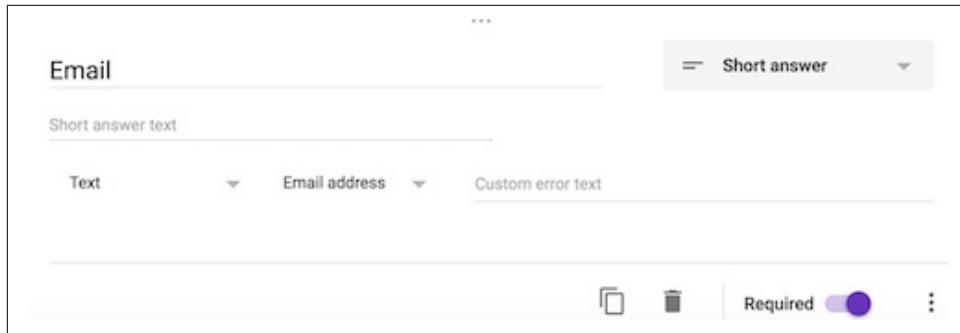


Figure 5-14. Adding a specific validation rule to a field

When you've gotten the form to your liking, you can preview it in your browser. Note that you do not have to add a submit button. Google does that for you.

A screenshot of a contact form titled 'Contact Form'. It contains three required fields: 'Name', 'Email', and 'Comments', each with a red asterisk. The 'Name' field has a placeholder 'Your answer'. The 'Email' field has a placeholder 'Your answer'. The 'Comments' field has a placeholder 'Your answer'. Below these fields is a blue 'SUBMIT' button. At the bottom of the form, there is a note: 'Never submit passwords through Google Forms.' The entire form is set against a light purple background.

Figure 5-15. The contact form in action

Once you've gotten your form to your liking, actually using it is a bit weird. You need to click the "SEND" button on top. This isn't terribly obvious, but this is where you'll find the embed options for the form.

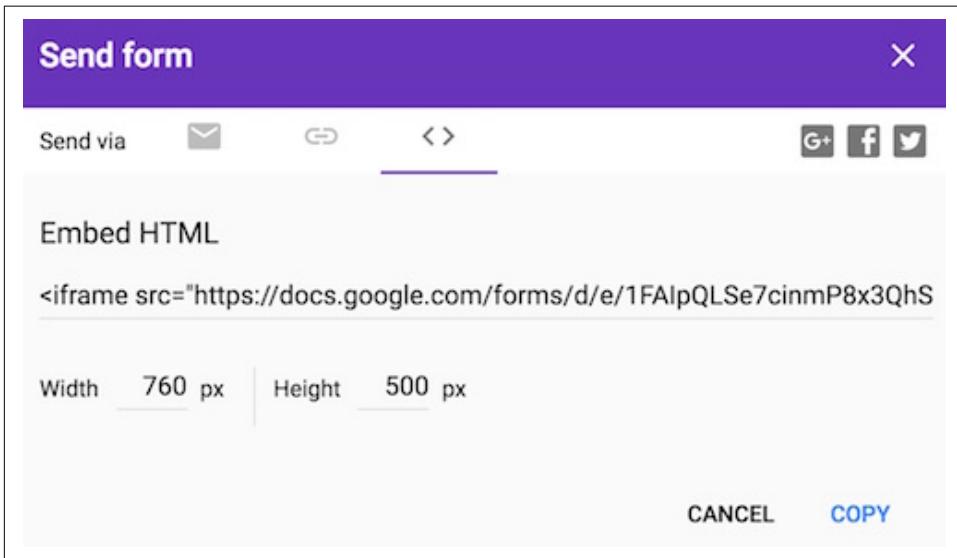


Figure 5-16. Reaching the embed options for Google Forms

As before, copy the code into your HTML file in your static site, and you're good to go. As before, you can find an example of this in the source code the book: ch5/forms/googl1.html. Here the form is in a simple HTML page.

Testing Google Docs

The image shows a screenshot of a Google Doc page titled "Testing Google Docs". Inside the document, there is an embedded "Contact Form". The form has a purple header with the title "Contact Form" and a red asterisk indicating it is required. Below the header, there is a light blue input field labeled "Name *". Underneath the input field, the placeholder text "Your answer" is visible. At the bottom of the form, there is another input field labeled "Email *". Both the "Name" and "Email" fields have red asterisks next to them, indicating they are required fields.

Figure 5-17. The embedded form

Notice how the iframe isn't necessarily sized correctly by default. You can customize this in the embed options and - unfortunately - you'll need to guess a few times to get it just right.

Once submitted, the user is presented with a simple message (that can be customized) and results will be available back in the Google Form editor.

The screenshot shows a Google Form interface with a purple header bar. The header includes tabs for 'QUESTIONS' and 'RESPONSES' (which shows '1'). Below the header, it says '1 response'. There are two buttons: 'SUMMARY' (highlighted in purple) and 'INDIVIDUAL'. To the right, there's a green plus icon, a vertical ellipsis, and a switch labeled 'Accepting responses' which is turned on. The main content area contains three entries: 'Name (1 response)' with the value 'Ray'; 'Email (1 response)' with the value 'raymondcamden@gmail.com'; and 'Comments (1 response)' with the value 'moo'.

Figure 5-18. Google Form responses

Email notifications are possible, but a bit weird. You'll use the “Add Ons” link on top to add notifications. It isn't just an option, but an optional extra. Once you've done that, you can tell Google to notify you. Once cool thing is that it can batch responses. This could be useful for a busy form.

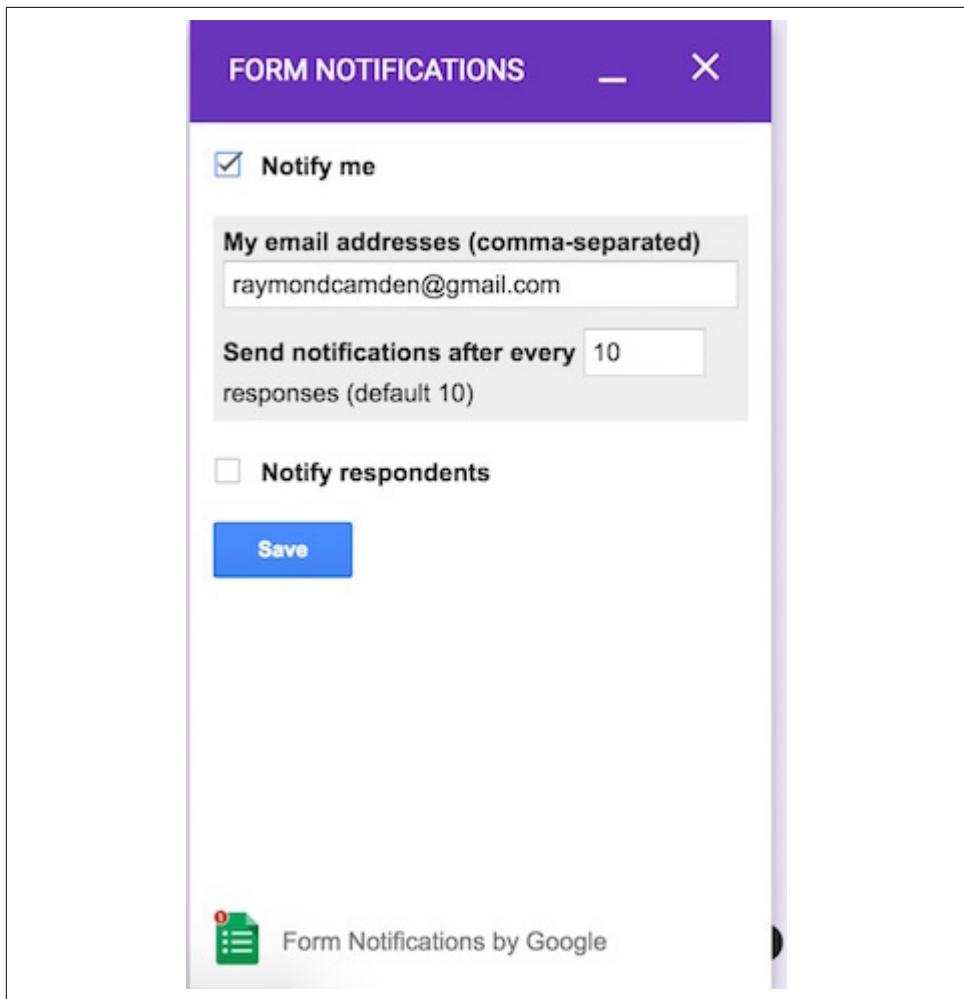


Figure 5-19. Setting up notifications

Formspree

Our next option is a service actually specifically built for static sites, Formspree (<https://formspree.io/>). How easy is Formspree? Let's look at an example. Begin with a form:

```
<form method="post">  
  <p>  
    <label for="name">Name: </label>  
    <input type="text" name="name" id="name">  
  </p>
```

```

<p>
<label for="email">Email: </label>
<input type="email" name="email" id="email">
</p>

<p>
<label for="comments">Comments: </label>
<textarea name="comments" id="comments"></textarea>
</p>

<input type="submit" value="Send Comments">

</form>

```

This is a simple form with three fields. The action is blank. To enable Formspree support, simply change the action to `https://formspree.io/someemail@some.domain`. So here is a real example (this can be found in ch5/forms/formspree1.html):

```
<form action="https://formspree.io/raymondcamden@gmail.com" method="post">
```

And literally - that's it. Open the form (after editing the action value to use your own email address), fill it out, and hit submit:

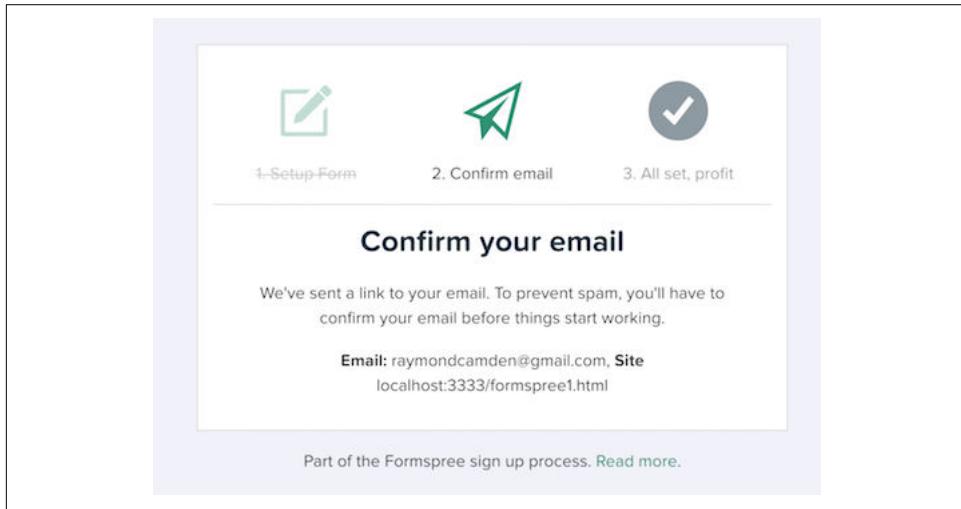


Figure 5-20. Formspree's confirmation dialog

What you're seeing is a one time security feature where Formspree wants you to confirm you want to get email from the service. They will do this for every unique referer URL and email address but again - it is a one time process. Click the confirmation link in your email and submit the form again. Now you'll see this:

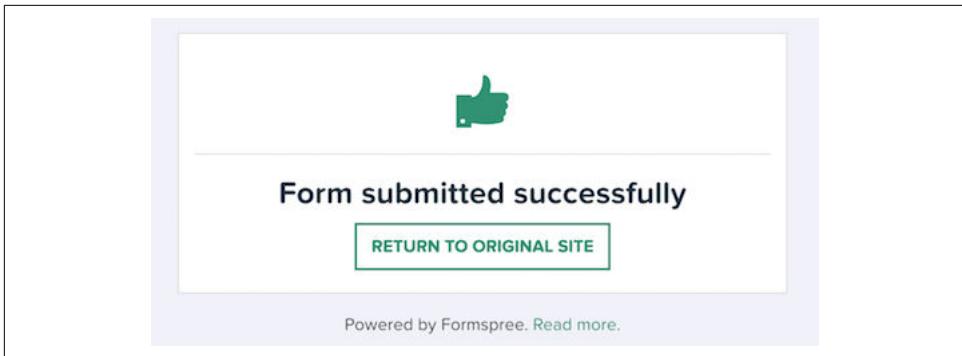


Figure 5-21. Formspree's success message

There's no signup. No cost (for the first one thousand submissions). It's as simple as that. Of course, you'll probably want to use a custom response. You can do that by simply adding a hidden form field (`ch5/forms/formspree2.html`):

```
<input type="hidden" name="_next" value="http://localhost:3333/thanks.html" />
```

Now when you submit the form, Formspree will send the user back to the URL provided in the value. (Note that Formspree will require you to validate the email again as this is a new file, and new URL, being used to send the form.)

You can even customize the subject by - yes - another hidden form field:

```
<input type="hidden" name="_subject" value="My Site's Comment Form" />
```

Formspree is truly an incredibly simple, and powerful, forms service for static sites. As stated above, the free tier covers one thousand emails a month which is incredibly generous. You can read about additional features at the website: <https://formspree.io/>

Adding a Comment Form to Camden Grounds

Now that we've seen a few examples of handling forms in a static site, let's update the earlier HarpJS based demo site, Camden Grounds, to support a contact form. If you've downloaded the code for this book from the GitHub repository, you can find an updated version of the demo site in `ch5/forms/camdengrounds`. The modifications to support a form were as follows:

- First, the menu (both top and bottom) was modified to add a link to `contact.html`. This is done in the `_layout.ejs` file. If you skipped over chapter 2, you should quickly give it a read to become familiar with the basics of Harp.
- Then a contact form was added, named `contact.ejs`.
- When Formspree is done processing the form, we want it to return to another new file, `thanks.ejs`.

- Finally, some minor CSS tweaks were done to the site. Note that the styling you'll see in the demo is not spectacular. Blame the developer (me!), not Harp or Formspree!

Let's begin by looking at the contact form.

```

<div id="figure">
  
  <span>Tell us your thoughts!</span>
</div>
<div>
  <div class="contactForm">

    <form action="https://formspree.io/raymondcamden@gmail.com" method="post">

      <% if(environment == "production") { %>
      <!--
          NOTE THIS IS NOT A REAL URL!
          Obviously someone may buy the domain and put something naughty.
      -->
      <input type="hidden" name="_next"
      value="http://www.camdengrounds.com/thanks.html" />
      <% } else { %>
      <input type="hidden" name="_next"
      value="http://localhost:9000/thanks.html" />
      <% } %>

      <p>
        <label for="name">Name: </label>
        <input type="text" name="name" id="name">
      </p>

      <p>
        <label for="email">Email: </label>
        <input type="email" name="email" id="email">
      </p>

      <p>
        <label for="comments">Comments: </label><br/>
        <textarea name="comments" id="comments"></textarea>
      </p>

      <input type="submit" value="Send Comments">

    </form>
  </div>
</div>
```

For the most part, this is just a vanilla form. We made things a bit fancy by making use of an advanced Harp feature, the `environment` variable. This lets us detect if we're running locally via the Harp server or using the compiled, static version. This way we

can submit the form both on our development machine as well on the live site. (And as the source code says, we're using a fake domain that's fake as of this moment, but may not be fake at the time of publication. Be careful!)

Here's the new contact form in the site itself. (And again, the CSS could be a bit nicer!)

A screenshot of a website for "YAY! KOFFEE". The header features the logo "YAY! KOFFEE" with a stylized coffee cup icon. The navigation menu includes links for Home, Menu, Locations, Contact (which is underlined to indicate it's the active page), and About Us. Below the menu, a large green banner with the text "Tell us your thoughts!" in yellow is overlaid on a background image of a smiling woman holding a coffee cup. The main content area contains three input fields: "Name:", "Email:", and "Comments:", followed by a "Send Comments" button. At the bottom of the page, there's a footer with the "YAY! KOFFEE" logo and copyright information, along with links for Home, Menu, Locations, Contact, and About Us, and social media icons for Facebook, Twitter, and YouTube.

Figure 5-22. The new Camden Grounds contact form

And that's it! The nice thing about Formspree is how simple it is to integrate it into your static site. Don't forget when you publish your site, you want to immediately fill out the form so Formspree can confirm your address the first time.

Adding Comments

One of the best ways to improve engagement with your site is by adding comments. Of course, whether or not comments is a good idea is totally based on the nature of the visitors to your site, but if you have a (generally!) good audience, then allowing them to comment can be an incredible way to increase traffic.

Multiple different services exist now to support adding comments to a site, whether static or dynamic. The most common options are:

- Disqus (<https://disqus.com>)
- Livefyre (<http://web.livefyre.com>)
- Facebook Comments (<https://developers.facebook.com/docs/plugins/comments/>)

The most popular, and easiest to test, is Disqus, so let's see what's involved in using it.

Working with Disqus

The first thing you'll need to do is create an account (<https://disqus.com/profile/signup/>). After you've signed up, you're prompted to select what you'll be wanting to do - either just writing comments or adding comments to your site.

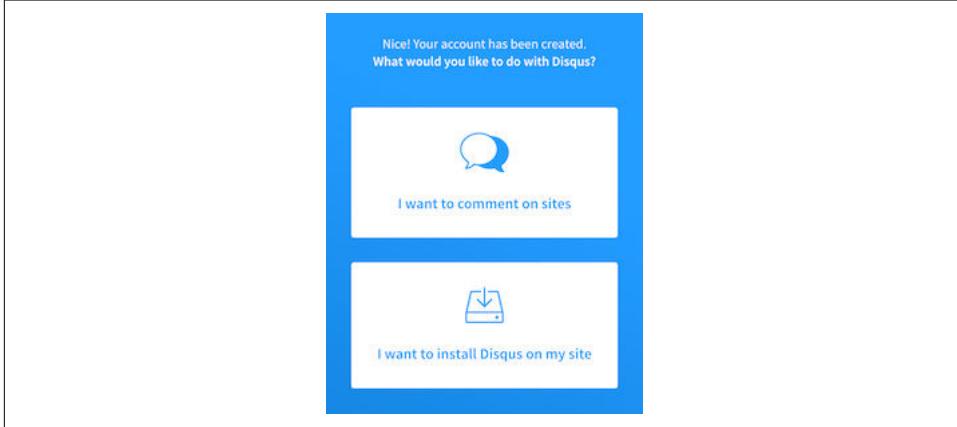


Figure 5-23. The Disqus signup process

Obviously you'll want to select the option for adding comments to your site. Next you'll be prompted to provide a name and category for your site. Enter whatever makes sense here.

A screenshot of a "Create a new site" form. It starts with a note that "All fields are required." The "Site Owner" field shows a profile picture of Raymond Camden and the name "Raymond Camden". Below it is a link to "login with a different account". The "Organization" section shows "Your Sites" and a note that "The organization is the group of sites you own." with a link to "Set an organization name". The "Website Name" field contains "Static Sites Book Test" with a note about the unique URL and a "Customize Your URL" link. The "Category" dropdown is set to "Tech". At the bottom is a large blue "Create Site" button.

Figure 5-24. Setting up your Disqus site

After a quick notice from Disqus, you'll be asked what platform you're using Disqus with. If you scroll to the bottom, you'll see an option for "Universal Code", that's the one you want.

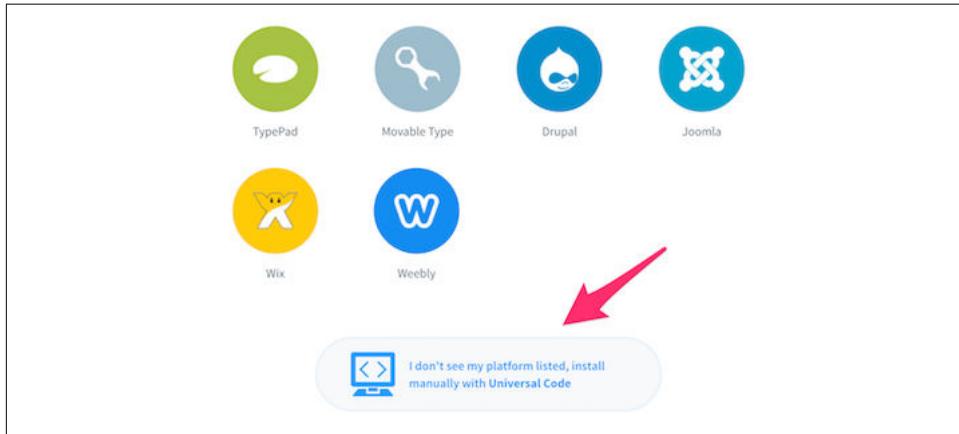


Figure 5-25. This is the option you are looking for...

On the next screen, you'll be given a script that you can copy and paste. There's other things you can tweak too (like configuration variables and CSS), but for now, simply copy that code.

A screenshot of a web page titled "Universal Code install instructions". It shows a code editor with a yellow numbered circle containing the number 1. The code is a JavaScript snippet for Disqus configuration. The code includes comments for defining configuration variables like PAGE_URL and PAGE_IDENTIFIER. A red "copy" button is visible at the top right of the code area.

Figure 5-26. The code to add Disqus to your site

Create any HTML file you want, or use the file from the GitHub repo ([cf5/comments/disqus1.html](#)) and copy in the Disqus code. Here's how it looks out of the box running on a local web server.

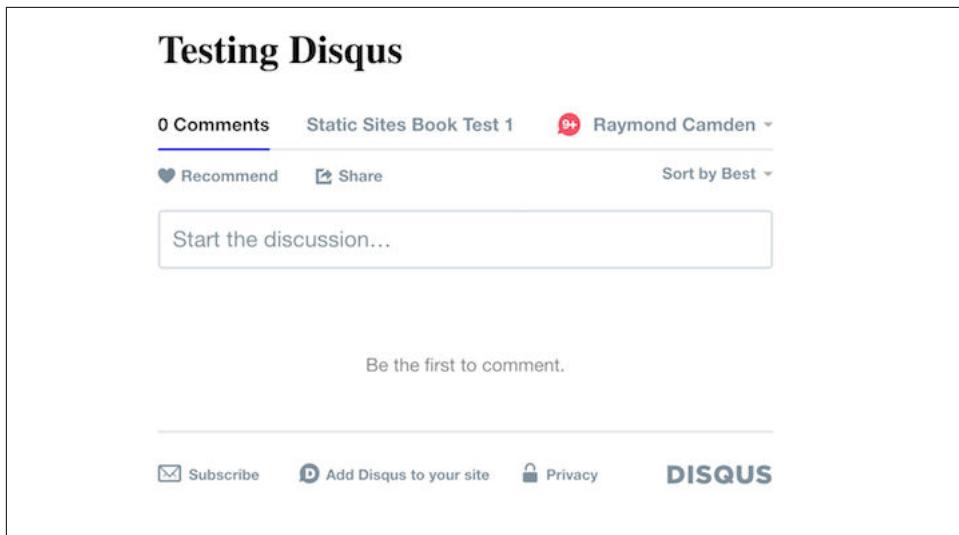


Figure 5-27. Comments for the test page

You can now start adding comments. If you use the code from the GitHub repo and do *not* change the code, you'll see the comments I added while testing, and comments from anyone else reading the book.

By default, Disqus uses the URL of the file to create a “thread”, or a unique set of comments. You can use the exact same code in another file and comments there will be different than the first file. You can see an example of this in `ch5/comments/disqus2.html`. The code is the exact same (except for the header), but the comments will be unique.

If you return to the Disqus site and click on the “Community” tab in the header, you can see an example of the tools Disqus provides you. You've got options for moderation, spam protection, even the ability to let other people become moderators of comments on your site.

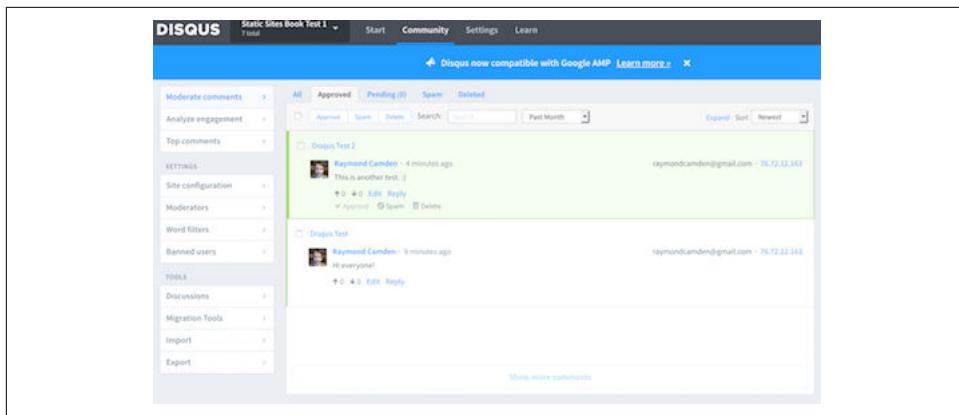


Figure 5-28. The Disqus Admin

You also get basic reporting options, but don't expect to see much immediately. I'm using Disqus for my blog (<https://www.raymondcamden.com>) so here is a screen shot from my site's dashboard:

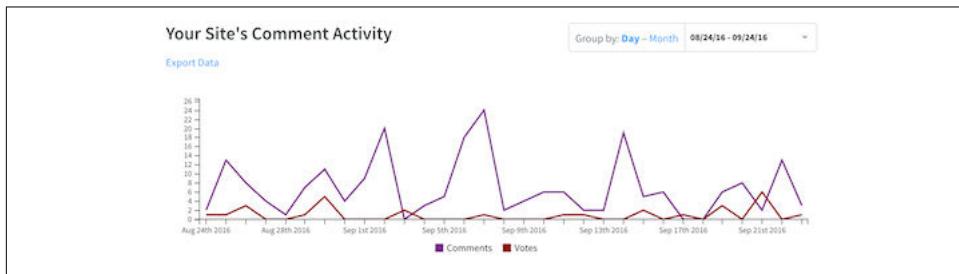


Figure 5-29. A chart showing comment activity on my blog

Adding Comments to the Cat Blog

In chapter 3, you learned how to use Jekyll to build a simple blog. Now let's update the blog to make use of Disqus comments. In the GitHub repository for this book, you'll find a new version of the blog in `ch5/comments/catblog`. Jekyll's template files are found in the `_layouts` folder. If you open `post.html`, you'll see that the Disqus code used earlier has been copied and pasted into the layout. The exact position was kind of arbitrary. I placed it over the pagination HTML but it could really go anywhere. Also, you don't need to restrict comments to blog posts only. If the blog also had an "About" page or other random pages, you can add commenting there as well.

And again... that's it. Once added to the posts layout, every blog post will now have comments!



Figure 5-30. The Cat Blog - now with comments!

Adding Search

For a small static site, like Camden Grounds shown earlier, a user can browse the entire site in minutes. For larger sites, like a blog, it would be near impossible to do so. This is where having a search engine of some form can be incredibly helpful. While a few options exist for adding search to a static site, the best option (at least in this author's opinion) is to use the undisputed king of search - Google.

Google provides a search called the Custom Search Engine (CSE). This service lets you create an embed that will use the power of Google's search index against a specific domain (or domains if you want). This means you can use a CSE to provide a Google-powered search for *only* your site's content, and nothing more. Let's take a look at how this is done.

Creating a Custom Search Engine

To begin, open your browser to <https://cse.google.com/cse> and login with your Google account. (Of course Google requires that.) You can then click the big obvious blue button "Create a custom search engine" to begin.

In the simplest form, a CSE consists of a site to search and a name:

Figure 5-31. Creating a new custom search engine

In case it isn't obvious, Google wants a *live* URL for an existing site. You can certainly use a domain that isn't live yet, but Google won't be able to index it obviously. This also means that if you are just now launching your static site, the search engine won't actually be able to return anything yet. That may make testing a bit difficult at first. But once your site is deployed, Google will index it rather quickly. You can enter any value you want here, but in order to get some content showing up immediately, you can enter the blog for my domain, raymondcamden.com. As the help text describes, you want to follow the domain with /* to signify you want to search everything under the domain. Finally, for the name of the search engine, enter anything appropriate. For now we'll use "Static Site Test". Hit the "Create" button and you're good to go.

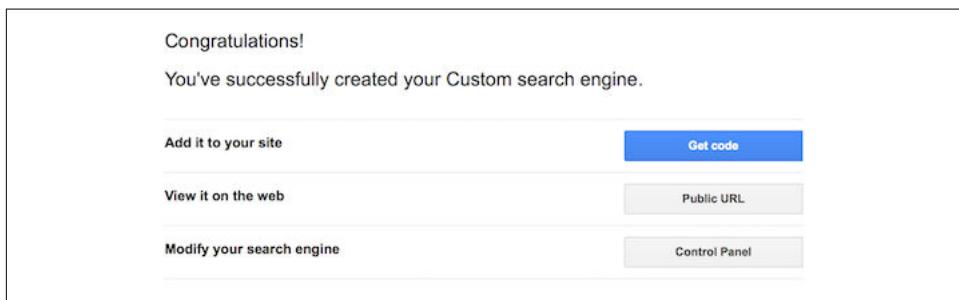


Figure 5-32. You've built the custom search engine

There's quite a few options you can provide to customize how the CSE works, but go ahead and click the button to get the code. Take that code and drop it in any simple HTML page. If you've downloaded the code for this book, you'll find an example in `ch5/search/test.html`. By default, the embed provides both the search field and submit button. It takes the full width of the enclosing content which means if you *don't* use CSS to restrict it, it will take up the entire width of the window.



Figure 5-33. The embed in action!

You can go ahead and type something to test the search. If you're using the version from the GitHub repo or my URL for the search engine, try searching for "star wars".

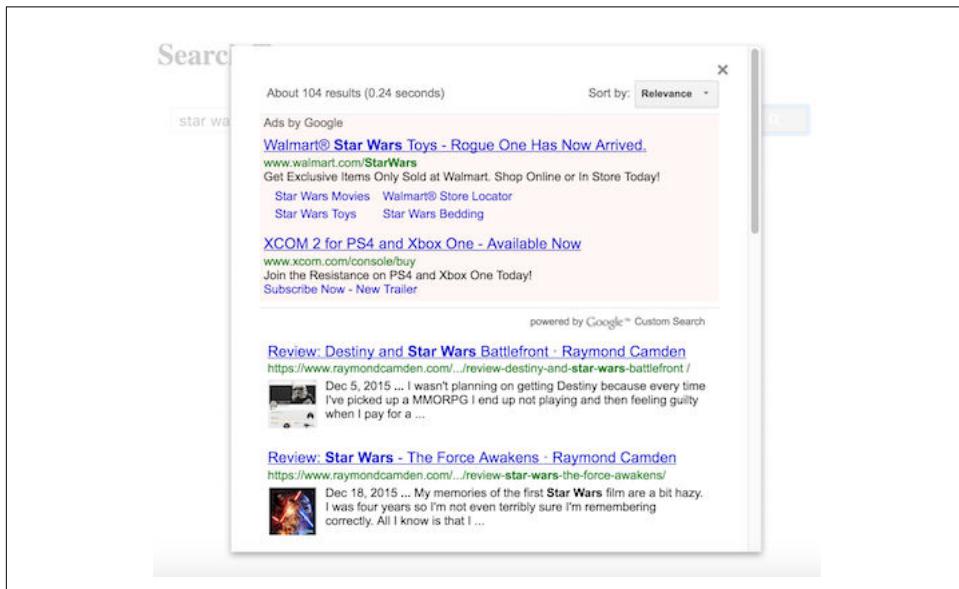


Figure 5-34. Search results from the CSE

You'll notice a few things right away. Yes, there are ads on top and no, you can't get rid of them unless you upgrade to a paid version of the CSE. On the screen shot above the ads take up a *lot* of the space available, but keep in mind the window used for the figure was shrunk quite a bit. On a "typical" web page (which you'll see an example of pretty soon) it won't be as overwhelming. Next, notice how Google creates a 'floating

window' type effect for the results. Finally, if you click one of the results, it opens in a new tab. All of that can be pretty annoying (although if you like that behavior, great!), but luckily you can change all of that pretty easily.

Back on the Google Custom Search Engine page, use the menu on the left and select "Look and feel". Under the "Layout" tab, you can see multiple different options to render the result. Clicking "Full width" will get rid of the overlay. Be sure to click the "Save" button, go back to "Setup" and click "Get code" to get an updated version of the code. Reload your test page and you can see the change.

star wars

About 104 results (0.28 seconds)

Sort by: Relevance

Ads by Google

[Walmart® Star Wars Toys - Rogue One Has Now Arrived.](#)
www.walmart.com/Sta...
Get Exclusive Items Only Sold at Walmart. Shop Online or In Store Today!

Star Wars Toys Star Wars Movies
Walmart® Store Locator Star Wars Clothing

[XCOM 2 for PS4 and Xbox One - Available Now](#)
www.xcom.com/console/buy
Join the Resistance on PS4 and Xbox One Today!
New Trailer - Subscribe Now

powered by Google™ Custom Search

[Review: Destiny and Star Wars Battlefront - Raymond Camden](#)
https://www.raymondcamden.com/.../review-destiny-and-star-wars-battlefront /
 Dec 5, 2015 ... I wasn't planning on getting Destiny because every time I've picked up a MMORPG I end up not playing and then feeling guilty when I pay for a ...

[Review: Star Wars - The Force Awakens - Raymond Camden](#)
https://www.raymondcamden.com/.../review-star-wars-the-force-awakens /
 Dec 18, 2015 ... My memories of the first Star Wars film are a bit hazy. I was four years so I'm not even terribly sure I'm remembering correctly. All I know is that I ...

[Finally! A good \(darn good\) Star Wars book - Raymond Camden](#)
https://www.raymondcamden.com/.../finally-a-good-dam-good-star-wars- book/

Figure 5-35. Nicer search results

Getting rid of the new window on search is a bit more of a pain. Unfortunately there is no simple toggle you can select in the CSE settings. If you look into the CSE docs you will find one specifically on customizing search results (<https://developers.google.com/custom-search/docs/element>). You'll see that Google provides multiple arguments you can add to the embed. Most services that give you embed code typically don't want you mucking with the code itself. They give it to you - you copy and paste it - and that's the end of it. But for Google's CSE, you are encouraged to modify the embed to fit your needs. This is the HTML portion of the embed as it comes, by default, from your new CSE:

```
<gcse:search></gcse:search>
```

From the docs Google provides, you can see an option, `linkTarget`, that defaults to `_blank`. That's the issue right there. Edit the `<gcse:search>` tag to change this to `_parent` (you can find this in `test2.html`):

```
<gcse:search linkTarget="_parent"></gcse:search>
```

Now when you search and click results, it will act like most search engines. If you return back to the dashboard, you'll find numerous other layout options to let you customize much more of the display of the results.

Adding a Custom Search Engine to a Real Site

In the previous sections of this chapter we demonstrated adding dynamic aspects back to sites by using earlier examples from the book. For this final example we'll use a site that's actually deployed live now, raymondcamden.com.

My site makes use of Hugo (covered in chapter 4) as a static site generator. Currently my blog has over five thousand and five hundred entries making a good search engine a requirement. I created my CSE much as described above (changing the layout results to get rid of the overlay and using the `linkTarget` attribute) and then added it to a new page, `search.html`.

```
+++
title = "Search"
+++

<div>
<script>
(function() {
var cx = '013262903309526573707:i2otogiya2g';
var gcse = document.createElement('script');
gcse.type = 'text/javascript';
gcse.async = true;
gcse.src = (document.location.protocol == 'https:' ? 'https:' : 'http:') +
'//cse.google.com/cse.js?cx=' + cx;
var s = document.getElementsByTagName('script')[0];
s.parentNode.insertBefore(gcse, s);
})();
</script>
<gcse:search linktarget="_parent"></gcse:search>
</div>
```

The first thing to note is that the embed is surrounded by a simple empty `div` tag. Why? Markdown tried to render parts of the `gcse` tag and broke the embed. By wrapping the code in a `div` Markdown will ignore the content inside.

One more small tweak was needed to integrate the CSE into my site. On the top right corner of my blog is a search field. This form submits to the search page. In order for

the CSE to “pick up” on the user input, I needed to name my search field “q”. Here is the form:

```
<form action="{{ .Site.BaseURL}}search/" method="get" accept-charset="UTF-8" class="search-form">
  <input type="search" name="q" results="0" class="search-form-input"
    placeholder="{{with .Site.Data.l10n.search.placeholder}}{{.}}{{end}}">
  <button type="submit" class="search-form-submit"></button>
</form>
```

If you don’t want to use the name “q” for any reason, you can also tell the CSE to use another field. As described in the documentation linked to earlier, if you add `query ParameterName="something"` to the `gcse` tag, you can tell the embed to look for a value in another field name. The choice is yours.

Even More Options

We’ve only scratched the surface in terms of the types of services you can add to your static sites. Here’s a few more examples to consider.

- Google isn’t the only option for search. You can even use completely client-side solutions like lunr.js (<http://lunrjs.com/>). I believe solutions like this are fundamentally dangerous as your content scales, but if you are sure about the size, and *potential* size, of your site, it may be an acceptable solution.
- If your organization has events, you may consider using a service like Eventbrite (<https://www.eventbrite.com/>). While Eventbrite can help you manage events and track attendees, it also lets you embed your events on your site.
- Another option for events is Google Calendar (<https://www.google.com/calendar>). Google Calendar has embed options as well, although in my experience the UI customization was a bit limited.
- If you remember that static site generators let you create non-HTML files too (like XML and JSON files), you could integrate with a plugin like FullCalendar (<https://fullcalendar.io/>). This jQuery plugin creates a nice interactive calendar that can integrate with a local data file (and Google Calendar too!). It gives you great control over the look and feel of how the calendar is used on the site.
- The “nuclear” option may be something like Firebase (<https://firebase.google.com>). This is a complete “data as a service” provider where anything at all could be stored. It provides a JavaScript client that could be used within a static site. This feels a bit... overkill to me, but it’s an option to consider nonetheless.

Finally, keep in mind that any external service is a dependency that is out of your control. You have to consider what will happen if that remote service goes down. How do you contact them to report an outage? How quickly do they respond? Also, some people automatically block certain embeds, like Disqus. Are you ok with that? These are all things you need to consider before employing these services.

CHAPTER 6

Adding a CMS

While using static site generators as an approach for both large and small-scale sites has been gaining more mainstream acceptance among developers, the experience for content contributors is generally sub-optimal. This is especially true for anyone who was used to editing content using the WYSIWYG editing tools found within popular blog engines like Wordpress or most content management systems. Transitioning these content editors to writing Markdown files in a text editor is probably a deal breaker for many companies. Solving this problem can often be more difficult than solving the technology problem, as Stefan Baumgartner describes in his article, [Using A Static Site Generator At Scale: Lessons Learned](#):

The biggest challenge in our journey to static sites was getting content editors to work with the new technology stack. You have to be hard as nails if you are willing to leave the comforts of your WYSIWYG editor — you know, those same comforts that drive web developers insane.

Stefan's solution was to build an editor with WYSIWYG featured for their less technical contributors, but this isn't necessarily a viable option for every project or team.

Even getting beyond the difficulties of editing content, publishing that content is non-trivial for your average non-technical content contributor. Depending on the setup, they either have to learn to test, build and FTP or how to commit and sync with something like GitHub to post their content. This can seem overwhelming to someone with a non-technical background who was used to just typing and hitting "publish."

The good news is that a number of tools have arisen in recent years to help make the process of creating and editing static sites for less technical users. While most of these are not free, they offer a range of features beyond simply providing a WYSIWYG content tool. In this chapter, we'll look at a handful of these including [CloudCannon](#), [Netlify CMS](#), and [Jekyll Admin](#).

CloudCannon

[CloudCannon](#) is a cloud-based CMS system specifically for Jekyll sites. While you can try it for free, it is a commercial product with various plans depending on your site's specific needs. CloudCannon is a hosted service, which means that if your business requires that your site be on-premise, it won't suit your needs, but as a hosted service it offers some additional benefits such as automatic optimizations including GZIP, minification and loading assets from a [CDN](#).

Another important consideration is that CloudCannon supports a limited set of Jekyll versions, currently 2.4.0 by default and 2.5.3 and 3.0.3 via configuration settings. This likely isn't an issue for a greenfield project, but could come into play if you are transitioning an existing Jekyll codebase. In addition, while CloudCannon does support Jekyll extensions, the feature is still currently in private beta as of this writing.

For full feature support and usage details, you can refer to their [documentation](#).

Creating a Site on CloudCannon

In order to create a site on CloudCannon, you'll need to sign up for an account. A 30 day trial is available for free, after which point it'll revert to a free account that has strict limitations on features and adds a splash page to the site.

Once you have an account and are logged in, click on "create site."

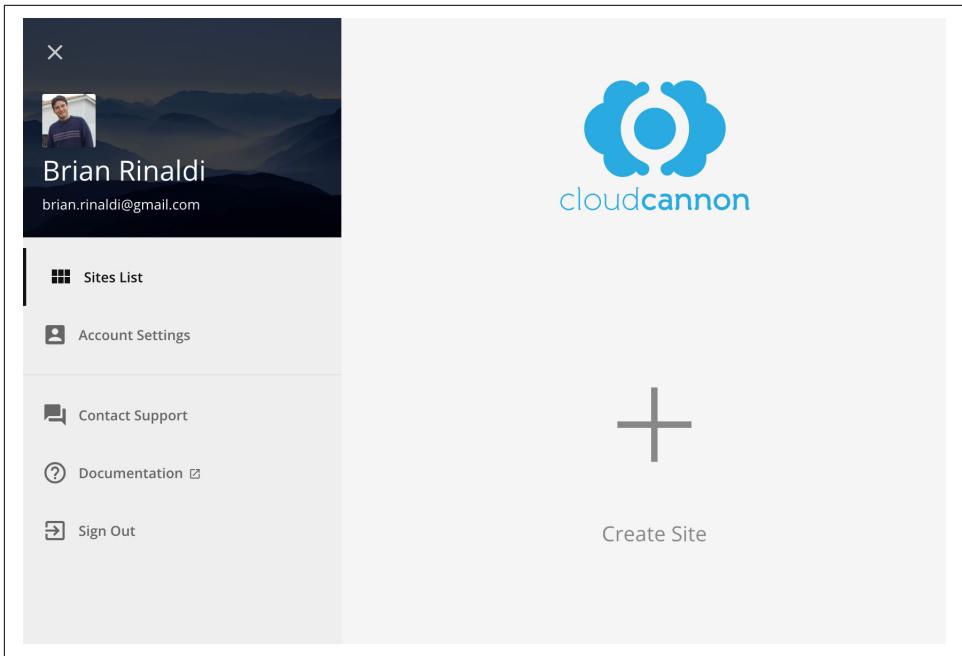


Figure 6-1. Creating a site in CloudCannon

cloudcannon_create.png

You should be prompted to give your site a name. For this tutorial, we'll be using a sample from a series of [samples that I have created for various static site engines](#) - obviously, in this case, we'll use the [Jekyll version](#). To make this process easier, however, I have placed a zip of the sample files in the [ch6 folder](#) of the GitHub repository for this book where you can [download it directly](#). Once you have it downloaded, unzip it. (Note that I have made some minor tweaks to the original source code to optimize it for some of CloudCannon's features.)

Since all of these samples are of an [Adventure Time!](#) fan site, let's name the site "Adventure Time"

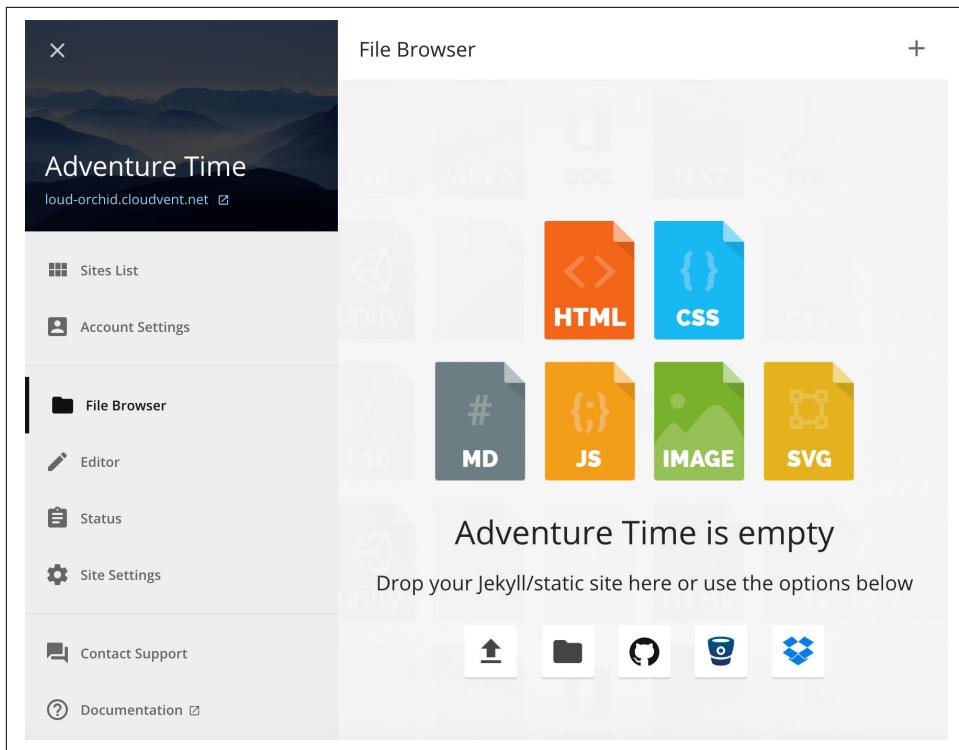


Figure 6-2. Creating a new site on CloudCannon

cloudcannon_newsite.png

One of the great things about CloudCannon is that it gives you a lot of options for pushing your Jekyll site onto its system. You can upload files individually, upload the contents of entire folders, connect to a GitHub or BitBucket repository or even pull files from DropBox.

To make things easy, for this example, we'll just use the folder option. Click the folder icon and locate and then select the folder containing the unzipped sample we downloaded earlier. After the upload is complete, your files should look like those in the screenshot below.

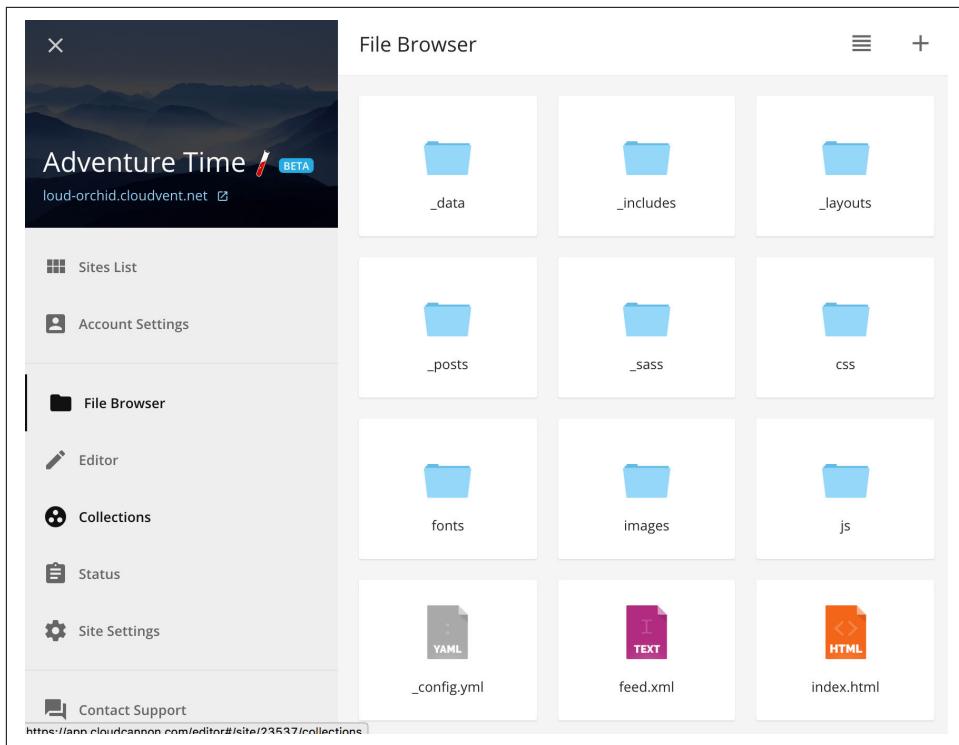


Figure 6-3. Uploaded source files on CloudCannon

cloudcannon_files.png

Editing a Site on CloudCannon

Now that we have the files uploaded, let's look at some of the editing options within CloudCannon.

The first option is that you can edit any file within the web-based code editor by clicking on it within the file browser.



```

1  ---
2  layout: default
3
4
5
6 <div class="row">
7   <div class="12u">
8
9     <!-- Characters -->
10    <section>
11      <header class="major">
12        <h2>Characters</h2>
13      </header>
14      <div class="row">
15        {% for character in site.data.characters %}
16          <div class="4u">
17            <section>
18              <span class="image featured"></span>
19              <header>
20                <h3>{{character.name}}</h3>
21              </header>
22              <p>{{character.description}}</p>
23            </section>
24          </div>
25        {% endfor %}
26      </div>
27    </section>
28
29  </div>
30 </div>
31 <div class="row">
32   <div class="12u">
33
34     <!-- Blog -->
35     <section>
36       <header class="major">
37         <h2>Recent Episodes</h2>

```

Figure 6-4.
The CloudCannon code editor

While the code editor works nicely, this doesn't solve our underlying problem of making the content easier to create and edit for non-technical authors — that is where the visual editor comes into play.

Click on the “Editor” option in the left-hand menu. This should open the site home page within the visual editor interface. You may notice that some of the text on the page is outlined in yellow — these text areas are allowed to be directly editable by the end-user.

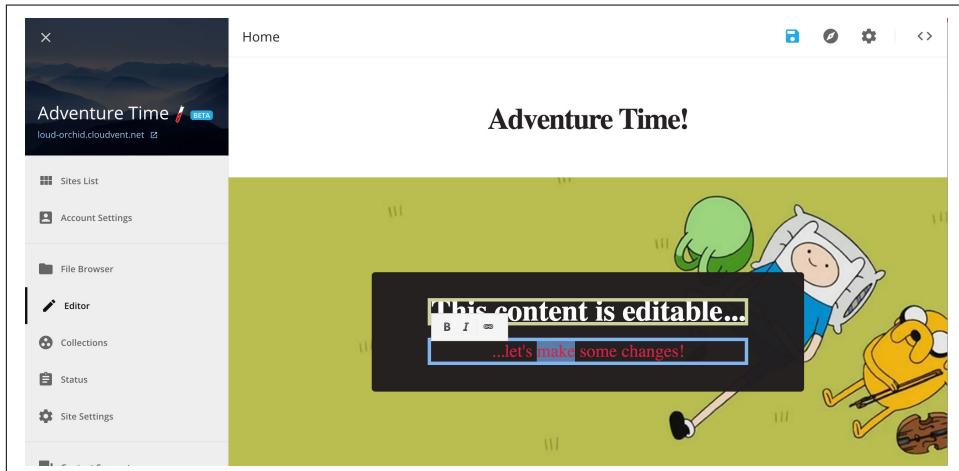


Figure 6-5. The outlined content is editable by the end user with basic formatting options provided.

cloudcannon_edit_text.png

Let's go ahead and click on the top editable item that reads "This content is editable..." You should not only be allowed to edit the text, but some basic formatting options will also be provided (i.e. bold, italics and link). In this case, formatting options are limited because we are editing a text element rather than a block element.

You can replace the text with whatever you like. I replaced the top with "Explore the Land of Ooo..." and the bottom with "...and its many kingdoms!"

So, how is the editing of text within the layout enabled? From within the template, you simply add a `class="editable"` to the tag. In this case, I added it to the `h2` and `p` tags within the header as you can see below.

```
<header>
  <h2 class="editable">This content is editable...</h2>
  <p class="editable">...let's make some changes!</p>
</header>
```

However, we can edit more than just plain text. Scroll down to the bottom of the home page within the editor and you should see an editable image and formatted text area within the footer.

Clicking on the image will give you the image editing options.

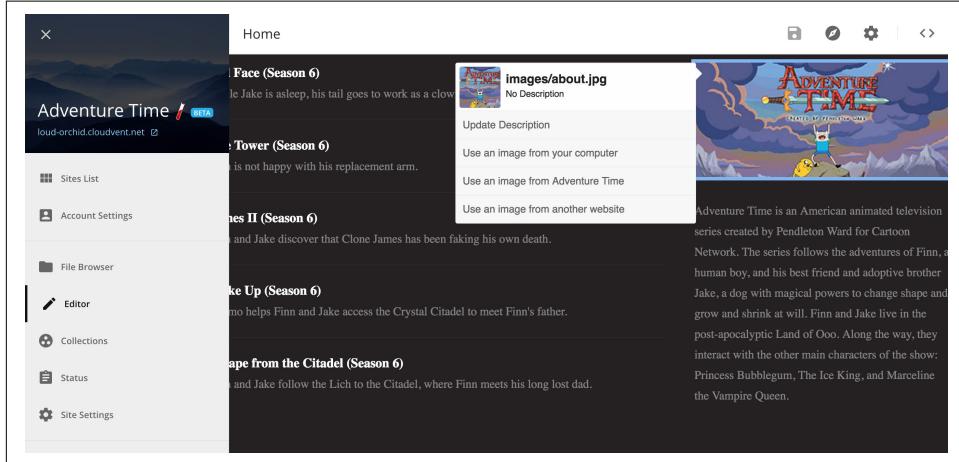


Figure 6-6. Editing an image within the CloudCannon visual editor

cloudcannon_edit_image.png

Clicking on the formatted text area gives you many more formatting options than the plain text area.

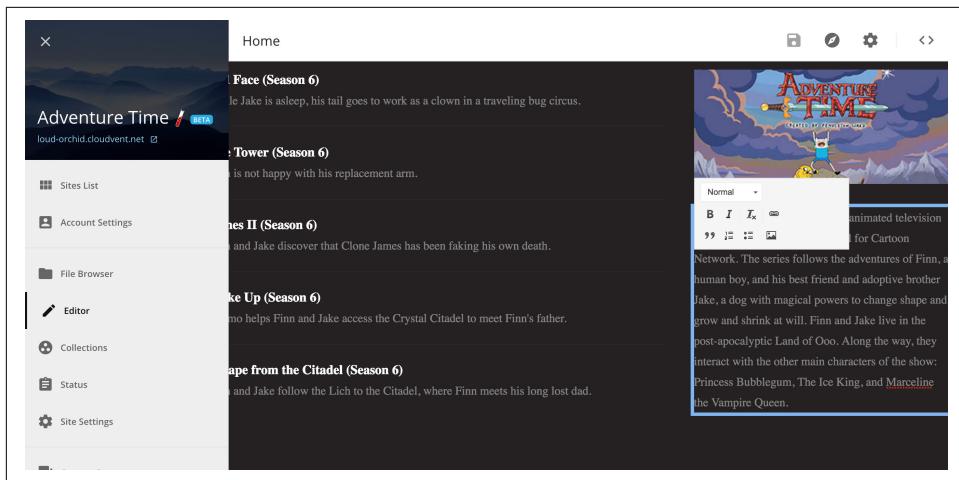


Figure 6-7. Editing a block text element within the CloudCannon visual editor

[cloudcannon_edit_block.png](#)

These were also made editable simply by adding `class="editable"` to the `img` and `div` tags respectively. Feel free to make any changes you like and then save them by clicking on the save icon in the upper-right hand corner of the editor.

Lastly, let's look at how to edit Markdown content within CloudCannon. Click on the “toggle pages” icon in the upper-right hand corner of the editor (it looks like a compass).

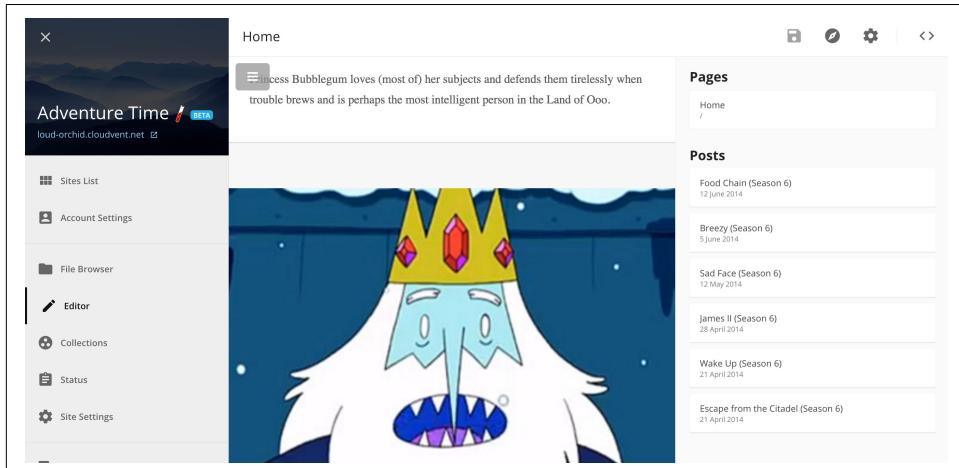


Figure 6-8. Pages and posts listed within the editor after clicking on “toggle pages”

[CloudCannon_pages_menu.png](#)

Next, let's choose the first post, "Food Chain (Season 6)." This will open the page in the visual editor (Note: you may want to toggle the pages menu off at this point to have a better view). At this point, you can view the page content within the editor and edit any of the editable elements on the page (i.e. the image and formatted text in the footer). However, what we'd like to do is edit the actual contents of the post.

To do this, click on the "switch to content editor" icon in the upper-right hand corner of the visual editor (it looks like a pencil). This will open the contents of the post within the content editor, where you can edit the body of the post using visual formatting tools (much like those from the formatted text area above). This means that your user doesn't need to know or understand Markdown to edit or create a page.

You may notice that you cannot edit the title of the page within the editor. Thankfully, the CloudCannon content editor provides a way to edit any of the settings for the page or post (these are the metadata items placed within the YAML front-matter on Jekyll posts and pages).

Click on the "toggle settings" icon on the upper-right hand corner of the page (it looks like a gear). This will open the settings editor.

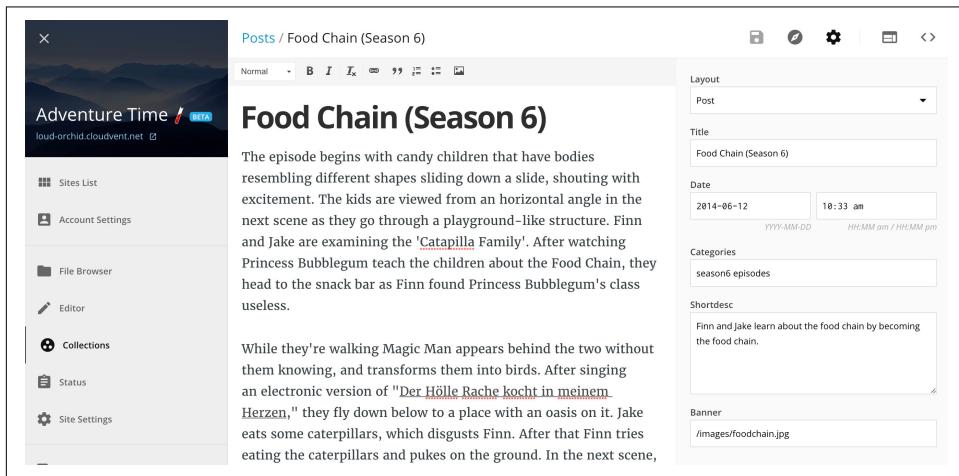


Figure 6-9. Editing post content and settings from within the CloudCannon content editor.

CloudCannon_settings.png

The editor not only allows you to edit standard metadata properties like title and date, but also custom properties like the `shortdesc` and `banner` properties within our sample site.

Where to Go From Here

Obviously, I've only touched on the basics of CloudCannon. Still, as you can see, it provides a lot of functionality for allowing user's who may not be comfortable editing HTML, Liquid, YAML and Markdown to make changes and maintain the content on the site. If you'd like to learn more about the features of the editors or of the hosting, check out the [CloudCannon site](#) or [documentation](#).

Netlify CMS

The [Netlify CMS](#) is not a hosted service like CloudCannon, but rather a free and open-source CMS-style admin that you can add to your static site project. The project is created and maintained by [Netlify](#), which does offer a popular service for continuous deployment and hosting of static sites, but this is not required to use the CMS project.

On top of being free, the Netlify CMS has the additional benefit that it is not tied to just Jekyll. As it just offers a backend to the content and data, it can, in theory, work with any static site generator. In fact, the project offers templates to work with Jekyll, Roots, Hexo and Pelican.



Versions of Netlify CMS

As of the writing of this, the original version of Netlify CMS, which was built using the Ember framework, was the most complete version. However, the tool is currently being rebuilt using the React framework, which will eventually replace the existing Ember version of the tool. The tutorial that follows here, uses the Ember version.

Setting up the Netlify CMS

So far, the two solutions we explored focused specifically on Jekyll. Let's try setting up the Netlify CMS for a static site using the [Hexo static site generator](#).

We'll start by getting a basic Hexo site running. I have placed a zip of the sample files in the [ch6 folder](#) of the GitHub, or you can [download the zip file here](#). Unzip it wherever you'd like to work on your project.



Sample Template

The example here is based upon the [Hexo template](#) provided as part of the Netlify CMS, which, in turn, was based upon my [static site samples project](#). The project includes some minor modifications to the original source to allow the tool to edit the character data.

Hexo is a Javascript-based generator that runs on Node.js. If you don't already have Node and npm installed, you'll need to [download and install it now](#). In addition, to make Netlify CMS run locally, you'll need to have [Git installed](#).

Once Node and Git are installed, let's start by getting our Hexo site running. Open a Terminal/Command Prompt in the folder containing the sample site. Once there, let's install Hexo and all of the necessary dependencies by simply doing an `npm install` via the command line.

Since we're installing things, there's one more piece that isn't necessary for Hexo, but will be necessary to run the Netlify CMS - the [Netlify Git API](#). It turns your local Git repository into a REST API that the Netlify CMS web application can call to do things like modify the content or add user access.

To install it, first grab the [appropriate version for your operating system](#). Next, unzip the file, making sure to note where you place the unzipped file (for instance, on Windows, you might simply place the file in a folder named `C://netlify-git-api` or on a Mac in `/Applications/netlify-git-api`). The final step is to add the folder to your [path variable](#). You can find instructions [for Windows](#) or [for Mac](#). If you have a Terminal/Command Prompt window open already, you'll need to close and reopen it for it to see changes to the path.

If you've done everything properly, you should be able to call `netlify-git-api` from the command line without error. For example, below you can see that the folder was added to my `$PATH` on my Mac and the results of calling the command.

```
[MCWFHBRRINALD:hexosite brinaldi$ echo $PATH  
/usr/local/bin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/Applications/netlify-git-api  
[MCWFHBRRINALD:hexosite brinaldi$ netlify-git-api  
usage: netlify-git-api [<flags>] <command> [<args> ...]  
  
Get a REST API for a Git repository  
  
Flags:  
  --help           Show help (also see --help-long and --help-man).  
  --db=".users.yml"  File path to the user db  
  
Commands:  
  help [<command>...]  
    Show help.  
  
  serve [<flags>]  
    Start a local Git API server  
  
  users list  
    List all users  
  
  users add [<flags>]  
    Add a new user  
  
  users del [<email>]  
    Remove a user  
  
MCWFHBRRINALD:hexosite brinaldi$ █
```

Figure 6-10. The `netlify-git-api` added to the `$PATH` on Mac and the results of calling the command.

`netlify-git-api.png`

If everything has finished installing, let's enter the following command to start the local Hexo server:

```
hexo serve -o
```

The `-o` option tells Hexo to open the page in your default browser automatically. You should see a version of the “Adventure Time!” fan page that I use as part of my [static site samples project](#).

The sample already includes the Netlify CMS admin template based upon the [Hexo template](#) they provide. However, it is important to understand how it was set up if you intend to build your own or modify the existing template to suit your needs.

Finally, in a separate Terminal/Command Prompt tab or window, let's go ahead and set up and start the Netlify Git API so that we can explore the CMS admin. First things first, you'll need to make sure you add the project to a local Git repository. Assuming that your command prompt is already open at the project's location, enter the following commands:

```
git init  
git add .  
git commit -m "First commit"
```

Now that we have a Git repository to work against, we can add a user to it via the Netlify Git API.

```
netlify-git-api users add
```

This command will start a series of prompts asking for the user's information. After it is complete, start the Netlify Git API by entering the command `netlify-git-api serve`.

Exploring the Netlify CMS Configuration

Within the sample project, open up the file `/source/admin/index.ejs` ([EJS](#) is the default templating language used by Hexo). This file is based upon the default HTML and configuration from the [Netlify CMS documentation](#). The key changes are in the configuration portion, which I am including below:

```
backend:
  name: netlify-api
  url: http://localhost:8080

media_folder: "source/assets/images" # Folder where user uploaded files should go
public_folder: "source"

collections: # A list of collections the CMS should be able to edit
  - name: "posts" # Used in routes, ie.: /admin/collections/:slug/edit
    label: "Post" # Used in the UI, ie.: "New Post"
    folder: "source/_posts" # The path to the folder where the documents are stored
    sort: "date:desc"
    create: true # Allow users to create new documents in this collection
    fields: # The fields each document in this collection have
      - {label: "Title", name: "title", widget: "string", tagname: "h1"}
      - {label: "Banner", name: "banner", widget: "image", class: "image featured"}
      - {label: "Short Description", name: "shortdesc", widget: "string"}
      - {label: "Body", name: "body", widget: "markdown"}
    meta: # Meta data fields. Just like fields, but without any preview element
      - {label: "Publish Date", name: "date", widget: "datetime"}
      - {label: "Categories", name: "categories", widget: "string"}
      - {label: "Layout", name: "layout", widget: "hidden", default: "post"}
  - name: "data"
    label: "Data"
    files:
      - name: "characters"
        label: "Characters"
        file: "source/_data/characters.yml"
        fields:
          - label: "Characters"
            name: "list"
            widget: "list"
            fields:
              - {label: "Name", name: "name", widget: "string"}
```

```
- {label: "Image", name: "image", widget: "image", media_folder: "assets/images"}
- {label: "Description", name: "description", widget: "string"}
```

We'll explore the backend a bit more in a moment, but first let's explore the pieces of this configuration.

The `backend` is what is essentially our REST API, which, in this case is the Netlify Git API that we set up earlier.

The `media_folder` designates where uploaded files, such as images for posts or pages, will be placed. The `public_folder` tells the CMS where the path to the source files so that the path to assets won't include it. For example, the `media_folder` here is `/source/assets/images`, but the path in the generated site will only be `/assets/images`.

The `collections` property defines the different types of content that the CMS can edit. In the case of our example site, the user will be able to add/edit posts as well as the character data. In order for this to work, the CMS needs to understand a little bit about the data these contain. The `fields` and `meta` properties under each define the data as well as the type of widget that will be used in the UI to edit it. The primary difference between them is that `fields` are generally elements of the page that will be displayed in the visual editor (i.e. they will be displayed in some manner as well on the page). Meanwhile `meta` are metadata elements that are not typically displayed and can be edited via the "settings" for each entry.

The benefit of this approach for defining the data is that we can define custom properties. For example, our posts contain properties like a "short description" (`short_desc`) and "banner" that are not part of a default Hexo post. It also means that we can allow users to edit data in YAML files, such as our site's character data. The one special thing to notice about the data structure for our characters is that we can define `fields` within `fields`, meaning we can manage even fairly complex data files.

However, this flexibility means that we also need to tell the CMS how to preview the output of our posts and data. This is also defined in `index.ejs`, just below the configuration from above.

```
<script type="text/x-handlebars" data-template-name='cms/preview/posts'>
<div id="main-wrapper">
  <div class="container">
    <!--Content-->
    <article class="box post">
      {{#if entry.banner }}
      {{cms-preview field='banner'}}
      {{/if}}
      <header>
        <h2>{{ entry.title }}</h2>
        <p>{{ entry.shortdesc }}</p>
      </header>
```

```
    {{cms-preview field='body'}}
  </article>
</div>
</div>
</script>

<script type="text/x-handlebars" data-template-name='cms/preview/characters'>
<div id="main-wrapper">
  <div class="container">
    {{#cms-preview field="list" as |character| }}
      <section class="box">
        <span class="image featured">{{cms-preview field="image" from=character}}</span>
        <header>
          <h3>{{character.name}}</h3>
        </header>
        <p>{{character.description}}</p>
      </section>
    {{/cms-preview}}
  </div>
</div>
</script>
```

What is defined here are **Handlebars** templates that tell the CMS how to display the output of our data. The first `script` block defines this for our posts and the second for our character data. This allows us not only to include our custom properties (like the short description), but also to match the preview to the actual page display as closely as possible.

Taking the Netlify CMS for a Test Run

Now that we've explored how to install and configure the Netlify CMS, let's quickly look at how it works for a user editing the content. Go to `http://localhost:4000/admin` in your browser (remember that your local Hexo server should be running on port 4000 by default).

We're first be prompted to log in. Use the user information that we created earlier when configuring the Netlify Git API.

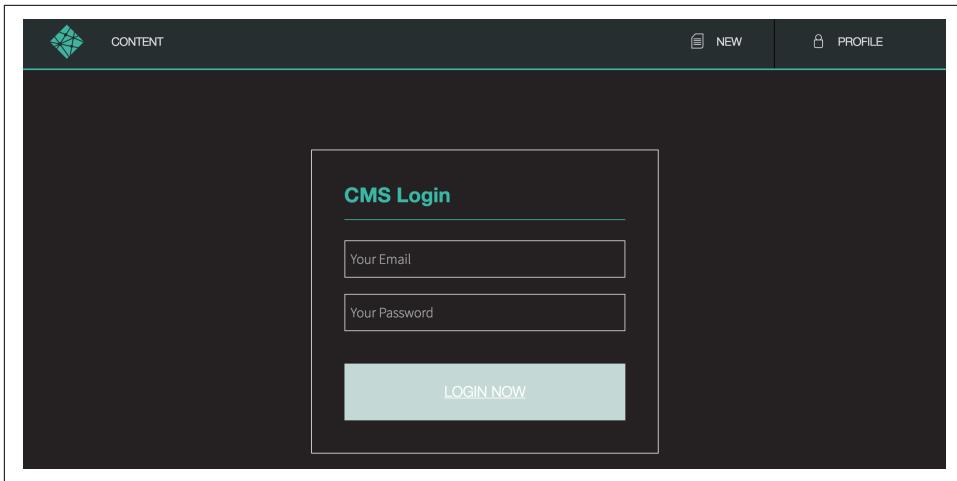


Figure 6-11. The login form for the Netlify CMS admin.

netlifycms_login.png

Once we log in, we are presented with a list of the existing posts from the site.

A screenshot of the Netlify CMS post list interface. The top navigation bar includes the 'CONTENT / POST LIST' logo, a 'NEW' button, and a 'PROFILE' button. On the left, there are filter buttons for 'Post' (which is selected) and 'Data'. The main content area is titled 'All Post entries' and displays two post cards. The first card is for 'Food Chain (Season 6)' and the second for 'Breezy (Season 6)'. Each card contains the title, a snippet of the post content, and a 'View' button.

Figure 6-12. All post entries in the Netlify CMS from our sample site.

netlifycms_posts.png

Clicking on a post will bring up the visual editor. Feel free to try it out and change some of the post's contents. Be sure to click on the “Settings” button in the lower-

right-hand corner to see how to edit the post's metadata. If you make any changes, click "Save" to save your changes.

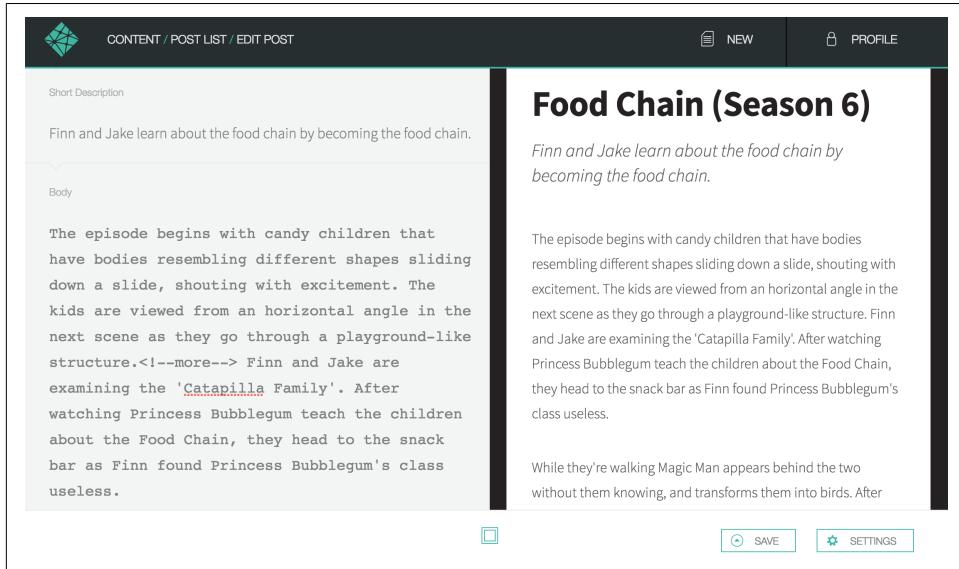


Figure 6-13. Editing a post using the Netlify CMS visual editor.

netlifycms_edit_post.png

To get back to the admin home, click the "Content" link at the upper-left hand corner of the page. Next click on "Data" in the left-side menu. This should bring up a list of the types of data that we have available to edit, which right now is only the characters.

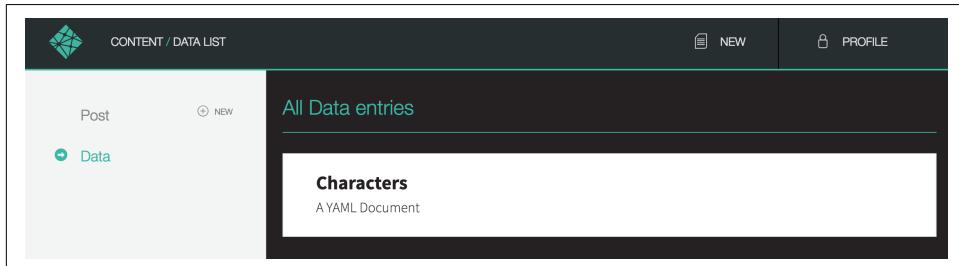


Figure 6-14. The available data types to edit via the Netlify CMS.

netlifycms_data.png

Click on "Characters." This will open up a visual editor for character data. Feel free to add or edit entries into our characters.

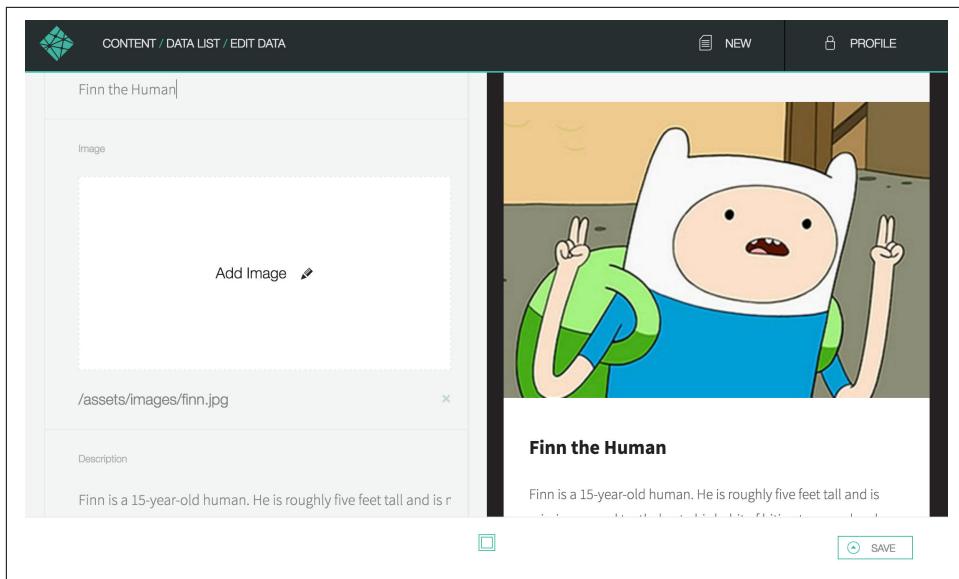


Figure 6-15. Editing character data via the Netlify CMS admin.

[netlifycms_edit_data.png](#)

Since we're editing static files on the file system, any changes we save should show up immediately on our site preview (you may need to refresh the browser to see them).

Where to Go From Here

Obviously, the Netlify CMS takes a decent amount of set up and configuration to get going, but it also offers a ton of flexibility. In this case, we are editing the key content elements of our site via an easy-to-use visual editor using the static site generator of our choice - in this case, Hexo. For anyone keen on using a particular static site generator, especially if they plan to customize it, this can be an ideal option for adding a CMS. If you would like to learn more about the Netlify CMS, check out the [full documentation on GitHub](#).

Jekyll Admin

Jekyll Admin is an extremely recent development, having only been released publicly in August of 2016 (literally, as I was writing this chapter). It isn't quite as full-featured as some of the other offerings we've looked at so far, but it has a couple of key things going for it.

First, it is the only officially supported CMS-like tool from any of the static site generator projects (at least that I am aware of). Not only that, but it is from the most popu-

lar and most widely supported static site generator available. (This probably also means that it sets an example other projects are likely to follow.)

Second, it is extremely easy to set up, while requiring no monthly subscription and no special configuration. Basically, if you have a Jekyll site, you can simply add the admin on and start adding and editing content using its graphical interface.

In my opinion, this kind of freely available and officially supported solution could potentially have a big impact on the adoption, not just of Jekyll, but of static site generators as a solution in general. That being said, the project is still very early, so it may not yet match the typical expectations of someone coming from a more traditional CMS solution. Let's take a look.

Setting Up Jekyll Admin

As I mentioned already, the setup and installation are incredibly simple. Let's walk through adding it to the [Jekyll version of the sample site](#) from my [static site samples GitHub repository](#). To make things a little easier, you can simply download a [zip of the Jekyll site files](#). Unzip them wherever you'd like to work on your project.

Assuming Jekyll is already installed (from earlier chapters in this book), the next thing we need to do is install Jekyll Admin. The easiest way is via [RubyGems](#).

```
gem install jekyll-admin
```

Inside the project files, open up `_config.yml` and add the following line:

```
gems: [jekyll-admin]
```

That's it! Now just open the Terminal/Command Prompt at the project folder and start Jekyll as we normally would.

```
jekyll serve
```

Our Jekyll admin is up and running. You can access it by going to `http://localhost:4000/admin`.

Editing a Site in Jekyll Admin

Opening the site admin makes it apparent that Jekyll Admin assumes that, if you have access to the admin, you are allowed to edit the entire site. This is an important consideration as one of the key features of Jekyll Admin is that it allows you to edit and, where applicable, create just about every aspect of the site from the configuration file, assets, pages and posts.

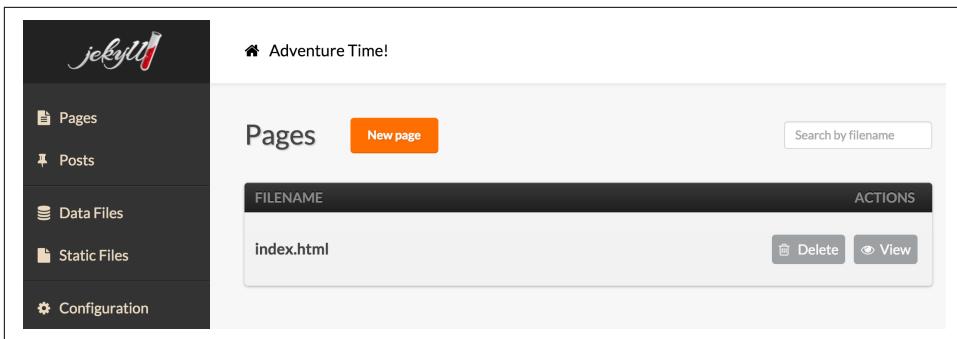


Figure 6-16. The main page of Jekyll Admin allows you to create and edit site pages.

jekyll_admin_home.png

Creating or editing pages (as well as posts) assumes that you are working in Markdown. So, while you can edit index.html, the formatting shortcuts on the editor add Markdown-specific formatting.

Let's take a quick look at the editing features.

The screenshot shows the Jekyll Admin interface for editing a post. The left sidebar contains navigation links: Pages, Posts (selected), Data Files, Static Files, and Configuration. The main content area has a header 'Adventure Time!' and a breadcrumb 'Posts / _posts/2014-06-12-season-6-food-chain.markdown'. The title 'Food Chain (Season 6)' is displayed with a 'Save' button. Below the title is a rich text editor toolbar. The post content starts with a paragraph about candy children. A large block of text follows, detailing a scene where Finn and Jake encounter a big bird and turn into bacteria. Another paragraph describes a scene where Finn and Jake turn into plants. The final paragraph is about Finn and Jake finding Erin. At the bottom of the content area, there are stats: 'lines: 17 words: 696 0:0'. The bottom section shows metadata fields: 'layout' (set to 'post'), 'date' (set to 'Jun 12, 2014 6:33 AM'), 'categories' ('season6 episodes'), 'shortdesc' ('Finn and Jake learn about the food chain by becoming the food chain.'), 'banner' ('/images/foodchain.jpg'), and a 'New metadata field' button.

Figure 6-17. Editing a post in Jekyll Admin

jekyll_admin_edit.png

The text editor is not the WYSIWYG (What You See Is What You Get) style editor that is common in popular CMS. Instead, you edit directly in Markdown, and can preview the output via a preview button. However, it is important to note that the preview is purely for text formatting, it does not offer a preview of the content within the site. Whether this is sufficient for the content contributors on your site depends on their skill level.

On the plus side, the editor automatically recognizes all the standard and custom metadata fields - and allows us to easily add more. On the minus side, editing metadata like the banner or layout use simple text editing, making it difficult to choose an image and potentially easy to typo the name of a layout.

Jekyll Admin also lets us edit any of the site's data files.

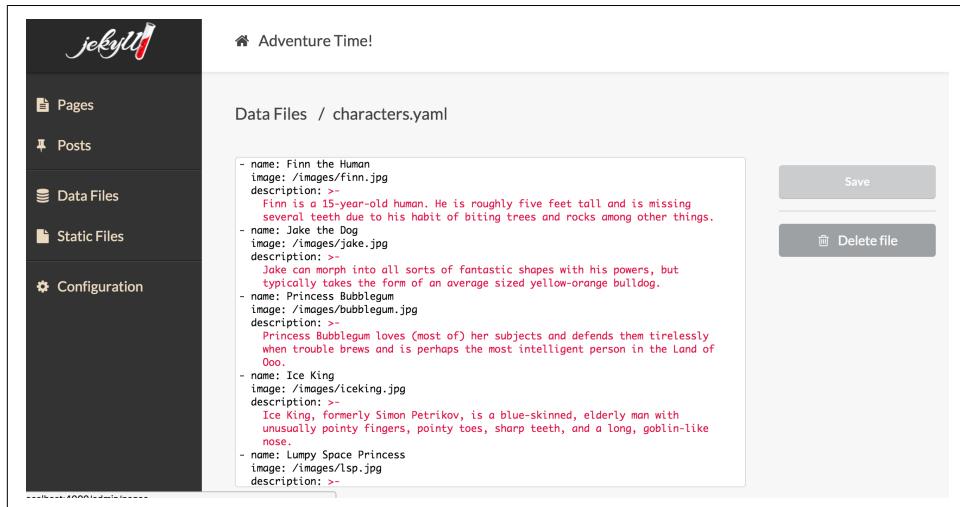


Figure 6-18. Editing a data file in Jekyll Admin

jekyll_admin_edit_data.png

However, the editing is directly in YAML via a simple text input. Again, this may not be suitable for your non-technical content contributors. The site configuration is edited in the same manner.

Jekyll Admin also makes it easy to add or delete static files for the site.

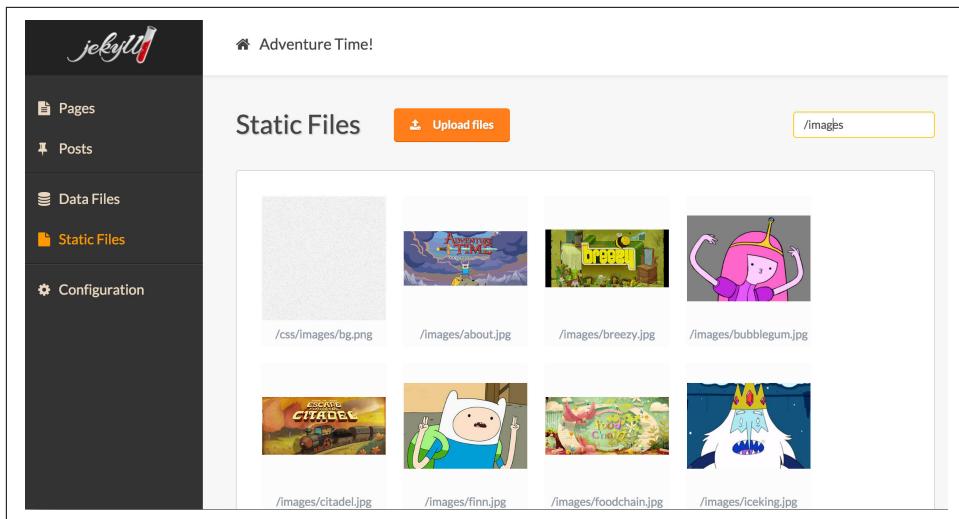


Figure 6-19. Viewing static site files in Jekyll Admin

jekyll_admin_staticfiles.png

However, this currently includes any static file on the site such as CSS files, and JavaScript files as well as images and assets. Someone with less experience could easily delete the wrong file by mistake. Also, there is currently no way to specify where an upload should be placed (they all end up in the site root), making this a far less useful feature for adding images to be used for things like pages or posts.

Where To Go From Here

Jekyll Admin is extremely easy to install and configure and allows you to add and edit pretty much any aspect of your site directly from within its visual editor. However, it is currently missing a lot of features and has some limitations that might make it more useful for your non-technical contributors to use. That being said, it is still an extremely new project and I certainly expect many of these issues to be addressed as it matures.

To learn more about the project, check out its [GitHub repository](#).

More Options

As static sites become a much more popular solution, there are an ever growing list of companies and projects that are stepping in to add in the CMS-like features that many users have come to expect when editing a site. Unfortunately, there is not enough space to cover them all here, but before we move on to the next chapter, I wanted to mention a couple of other significant options.

Forestry.io

Forestry.io offers a service similar to CloudCannon, but with some differentiating features such as support for both Jekyll and Hugo as well as the ability to easily deploy to a large number of destinations including AWS, GitHub and FTP. Editing a page or post allows you to easily modify standard or custom metadata attributes (and even choose the editor by which it should be edited - allowing you to edit the banner as an image but the description as text).

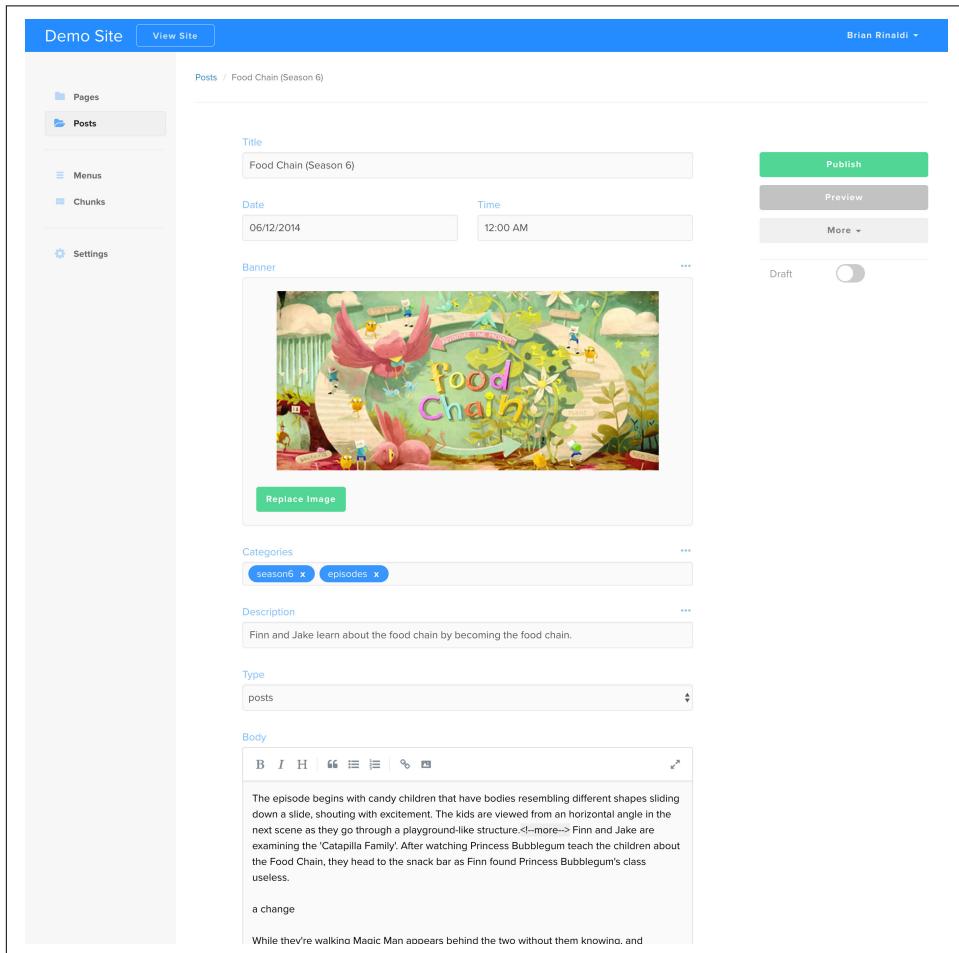


Figure 6-20. Editing a post in the Forestry.io demo site

[forestry_edit_post.png](#)

The editor is a Markdown editor (like Jekyll Admin) rather than a WYSIWYG editor (like CloudCannon), but it does have a shortcut that makes it easy to add images to a

page or post, and the form does also allow you to view an in-layout preview of the content. (On a side note, I need to share that I love that so many of these tools use my Static Site Samples site for their demo apps.) Of course, you can also edit data and site menus all using a simple visual editor.

Forestry.io is currently offered free for sites with up to 10 users without professional support. Paid plans are offered for more than 10 users or if you would like professional support.

Lektor

The final tool I want to mention is another very recent entry called [Lektor](#). Lektor is a static site generator built in Python. What makes it different from other Python-based tools like [Pelican](#), for instance, is that it has both a built-in admin as well as an installable desktop application (currently Mac OSX only).

The Lektor desktop application is designed to simplify running a local Lektor-based site that you already have on your computer, while also offering quick access to view and edit the site in a browser.

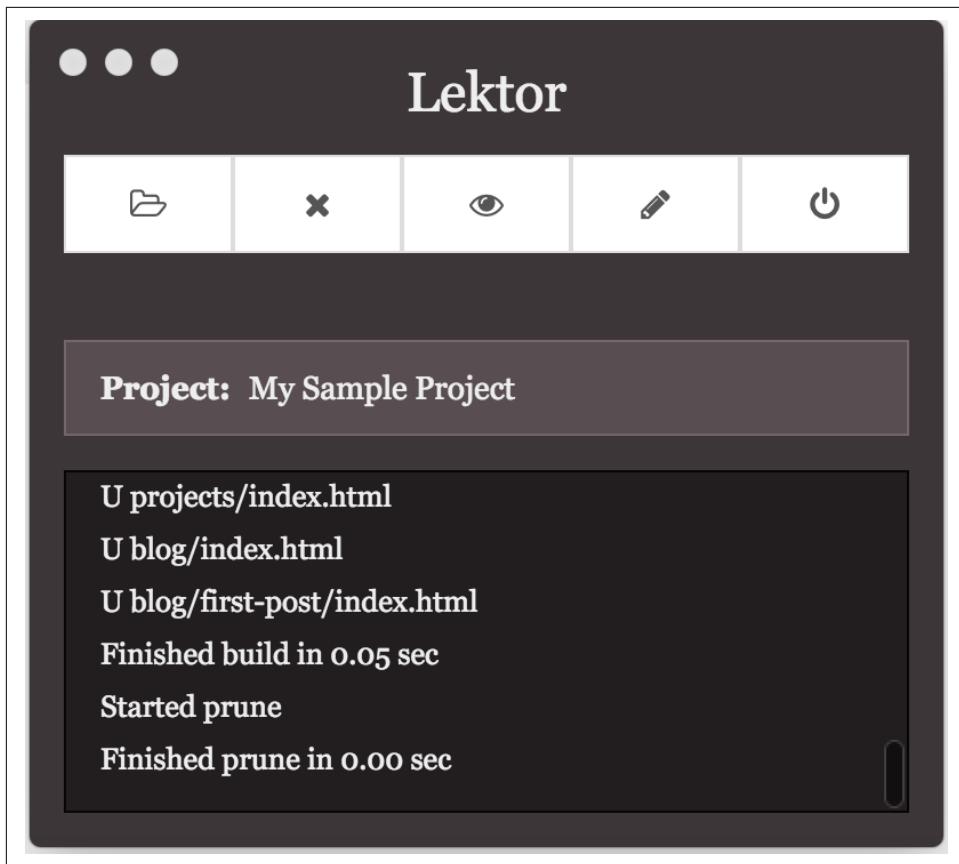


Figure 6-21. The Lektor desktop app running a sample site on Mac OSX

[lektor_desktop.png](#)

The admin for a Lektor site is pretty basic in terms of content editing - content is Markdown edited within a simple text area. However, it has a number of built-in features in terms of building out your site structure, uploading attachments and even deploying the site from the admin.

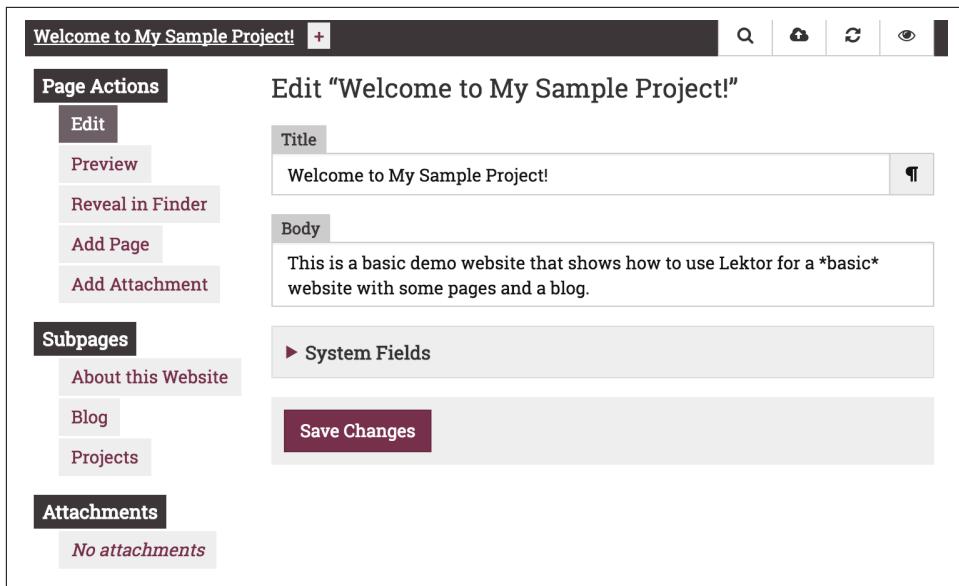


Figure 6-22. Editing a page in the Lektor site admin.

lektor_admin.png

As I mentioned, Lektor is also a relatively new project and it wouldn't be surprising to see it add any number of more robust editing features into either the admin or the desktop tooling.

Headless CMS

One last option I want to mention briefly is the concept of using a static site as the front-end to a [headless CMS](#). The term headless CMS is used because these services offer the back-end user interface of a traditional CMS for adding and editing content, without the front-end (i.e. the web site the content goes on). Instead, the service offers an API that you can connect to anything — a web site, a native mobile app or anything else that can consume this content.

In the case of a static site, rather than store your content as Markdown or YAML files on the filesystem, you'd store it within this CMS back-end. Then, you'd connect to the CMS backend via the API to generate your site.

A popular headless CMS service is [Contentful](#), which already offers official plugins for [Jekyll](#) and [Middleman](#). There is also an unofficial [Hugo tool for Contentful](#). Another option is [DatoCMS](#), which offers prebuilt integration with the Jekyll, Middleman, Hugo, Metalsmith, Hexo and Pelican static site generators.

The benefit of this option is that you are able to offer the CMS-style content editing that your authors may prefer and consume that content within your static site, but also consume it within any other type of application that can connect to the API.

CHAPTER 7

Deployment

One of the biggest selling points of static sites is that their production requirements are essentially nill. But even with the minimal requirements for supporting a static site, there's still multiple options developers can choose from.

In this chapter we'll discuss the various different ways you can take those simple, static files from your development machine and make them available to the world at large. Which you use is entirely dependent on your needs. As with all such choices, you will need to determine which option works best for you and your client.

Plain Ole' Web Servers

Probably the simplest, and most familiar, solution would be to make use of the same old web servers we've been using for the past twenty or so years. The two most popular options are the HTTP server from Apache (<http://httpd.apache.org/>) and IIS from Microsoft (<http://www.iis.net/>). Apache is available on multiple platforms while IIS is only available on Windows.

In both cases though if you have the server set up, then "deployment" is simply a matter of copying the output from your static site generator into the web root (or a relevant subdirectory) for your server.

Most likely you don't want to manually copy files every time, so you could look into various tools to make that process easier, like Grunt (<http://gruntjs.com/>) and Gulp (<http://gulpjs.com/>). Or you can simply use an old school shell script or BAT file as well.

There's nothing special about this setup and nothing else really to say, and that's a good thing!

Cloud File Storage Providers

Most developers are probably aware of cloud services provided by Amazon, Google, Microsoft, and others, that provide basic file storage. The idea being you can provision space and simply load as many files as you would like. You don't have to worry about the size of the disk - you simply treat it as an infinite hard drive that you can use as you see fit. (With costs, of course.)

What you may not know is that many of these providers also provide a way to turn their file storage system into a simple web server as well. In this section we'll take a look at how this is done with both Amazon's S3 service and Google's Cloud Storage service. Let's start with Amazon.

Hosting a Site on Amazon S3

Hosting with Amazon S3 requires you to have an AWS Account. You will have to provide a credit card during setup, but Amazon provides a large amount of disk space in their free tier. Obviously you should double check to ensure that the price is something you can afford, but in general, S3 storage is *incredibly* cheap. You can begin the signup process at <https://aws.amazon.com/s3>.



A Real Example

To give you a real example, I use S3 to host the media assets for my blog at <https://www.raymondcamden.com>. (The actual written content is stored in Surge, which we'll discuss later in the chapter.) I also host a few other small static sites there. My bills over the past year have hovered around the 10-15 cent mark.

After signing up, you will then go into the Amazon Web Services dashboard, which can be a bit overwhelming.

The screenshot shows the AWS Management Console dashboard. At the top, there are navigation tabs for AWS, Services, and Edit. On the right, user information (Raymond Camden, N. Virginia) and support links are displayed. Below the header, the main content area is organized into several sections:

- Amazon Web Services**: A large grid of service icons and names.
- Resource Groups**: A section for managing resource groups, with a "Create a Group" button and a "Tag Editor" link.
- Additional Resources**: A list of links to AWS mobile apps, marketplace, and re:Invent announcements.
- Service Health**: A status summary showing all services operating normally.

Some specific services listed in the main grid include EC2, Lambda, S3, CloudFront, CloudWatch, Cognito, Device Farm, Mobile Analytics, SNS, API Gateway, AppStream, CloudSearch, Elastic Transcoder, SES, SQS, and SWF.

Figure 7-1. The Amazon Web Services dashboard.

Simply look for (or in my case, 'CMD+F' for) S3 to open the specific dashboard for S3. S3 groups content into buckets. You can think of them as directories and generally you'll want a specific bucket for one specific web site. Any bucket can become a web site, but if you want to use a specific domain, like www.foo.com, then you must name the bucket with that domain. In case your curious, to support both foo.com and www.foo.com, you would create a bucket for foo.com that redirects to the www bucket. Redirects are supported by S3 but won't be covered here. Simply check the documentation for an example.

To begin, create a new bucket. You can name it what you will, but for the book we'll use orabook.raymondcamden.com. You will not be able to use the same name, so try using something that includes your own name, or if you have a domain at foo.com, try something like orabook.foo.com. For the region, just use US Standard. Depending on your location and your site's visitor's location, you may want to select a region that is closer.

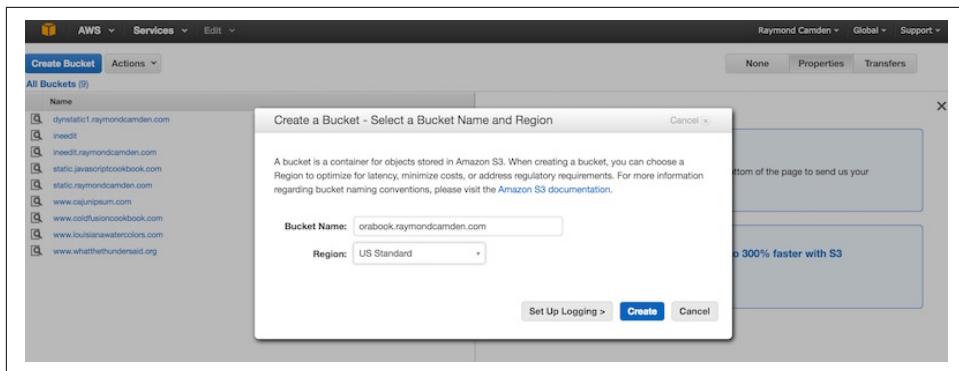


Figure 7-2. Creating a bucket for our site.

Immediately you'll see a properties panel open up and "Static Website Hosting" is one of the options. Click to open it, and you'll see how easy it is to enable web site hosting.

A screenshot of the AWS S3 bucket properties panel for 'orabook.raymondcamden.com'. The top section displays basic bucket details: Bucket: orabook.raymondcamden.com, Region: US Standard, Creation Date: Sat Aug 20 09:03:19 GMT-500 2016, Owner: cfjedimaster. Below this, the 'Static Website Hosting' section is expanded, showing the endpoint orabook.raymondcamden.com.s3-website-us-east-1.amazonaws.com. A note explains that each bucket serves a website namespace. Under the 'Enable website hosting' section, the 'Enable website hosting' radio button is selected. Fields for 'Index Document' and 'Error Document' are present but empty. A link to 'Edit Redirection Rules' is shown. At the bottom, a section for redirecting requests to another host name is visible.

Figure 7-3. Enabling static site hosting.

You will need to select an index document. This is simply the document that should be loaded when the site is requested. While you can enter anything here, the custom is to use index.html. You can enter an error document as well to be used when a page is requested that doesn't exist.

Once you click save, you can try hitting your site using the endpoint that S3 set up for you. In case you missed it, it is in the "Static Website Hosting" section in figure 7-3.

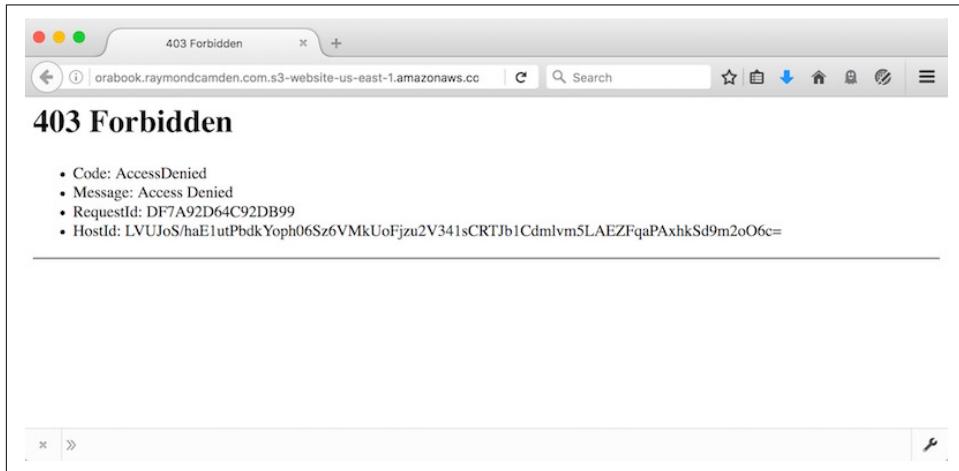


Figure 7-4. Your new site... broken.

Almost there! Obviously you need actual content in your bucket in order for a site to be displayed. How you get files into S3 is up to you. Most FTP clients support connecting to S3 buckets and that's probably the way you'll want to go, but you can also upload via the S3 web console. Simply use the **Upload** button and you can then drag and drop files and folders directly into the web page itself. For this demo, I've used the output from the Camden Grounds site created in Chapter 2. You can use any output from the earlier chapters, or simply create a new index.html file with some temporary content.

Before you hit upload, you need to modify permissions for the new assets so they are public. Click the **Set Details** button at the bottom, then click **Set Permissions**. On that page, click "Make everything public".

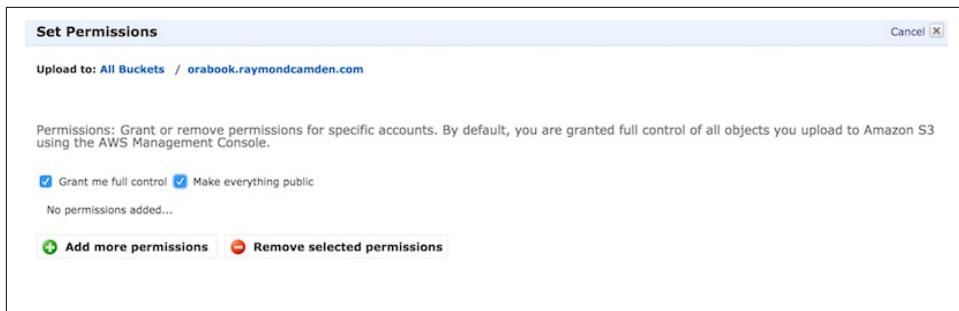


Figure 7-5. Making the assets public.

Finally, click `Start Upload` and the process will begin. Depending on how many files you're uploading and how large they are, it may take a few minutes. The website does a great job of providing feedback as assets are uploaded.

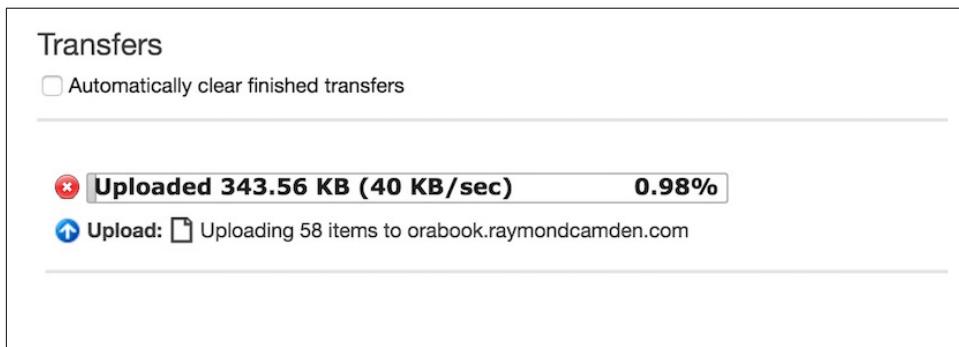


Figure 7-6. Upload progress.

When done, simply reload the browser with the endpoint URL and you will see your site!

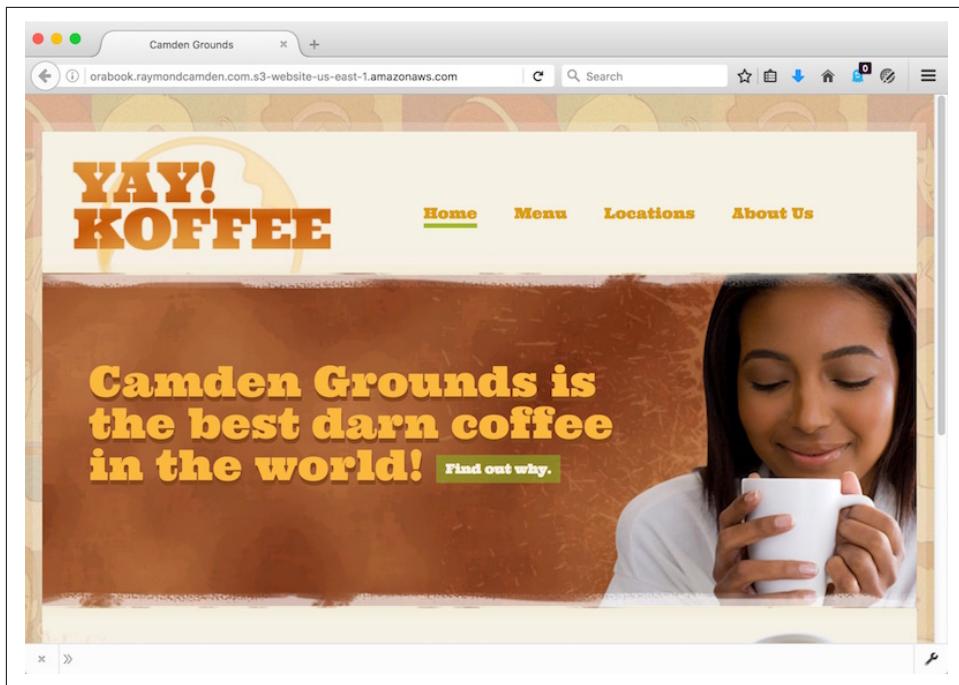


Figure 7-7. Your site is now live!

And that's basically it. The last step is to update the DNS for your domain to point to Amazon. Specifics on that can be found on the Amazon site. Amazon also supports basic redirects. If you migrated from a dynamic site using PHP for example, you could tell Amazon how to handle request for those old .php URLs so that visitors don't get errors requesting old URLs.

Hosting a Site on Google Cloud Storage

Working with Google's Cloud Storage system will follow a similar process to Amazon. As before, you'll have to sign up at the product page (<https://cloud.google.com/storage>) and provide credit card information again. As with S3, prices are very cheap, but I'll repeat the warning. Check the prices carefully so you don't get a surprise at the end of the first billing period. And again - just like S3, Google offers a generous free tier (currently 300 USD over 60 days of time).

After creating your account, the first thing you'll need to do (assuming you aren't a Google user already with existing services) is to create a project. You can open up the Project dashboard (<https://console.cloud.google.com/project>) and name it whatever you will. After creating the project you will then enable billing. Again, don't forget you've got 300 dollars of free service for 2 months, but also don't forget to close up the account later if you change your mind.

After you've done this initial work, you can simply open the storage browser (<https://console.cloud.google.com/storage/browser>) to create your bucket. (If this all seems familiar, that's good. It makes moving from S3 to Google, or vice versa, that much easier.)

Go ahead and click create a bucket, and let's give it the same name as we did for Amazon. Like with S3, you want your bucket name to match the domain you plan on using, and like S3, it has to be unique. Since we're creating this bucket in the book, you'll need to pick a different name.

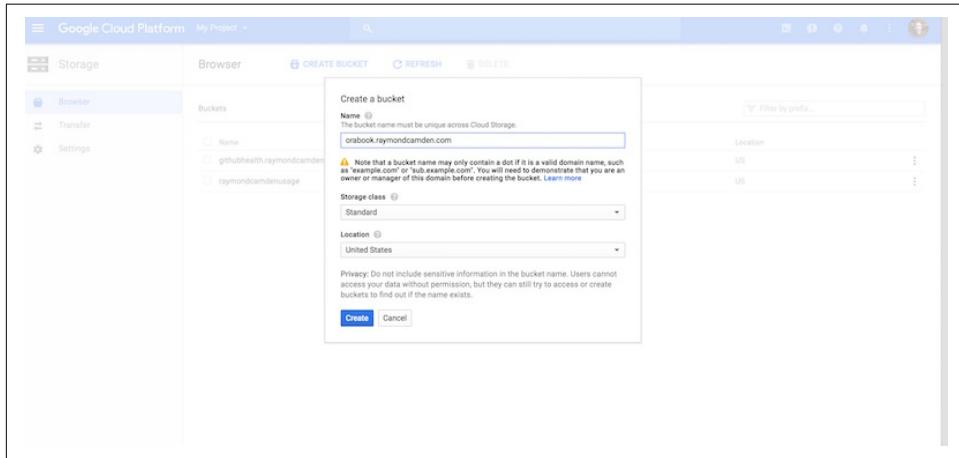


Figure 7-8. Creating the bucket on Google Cloud Storage

Once again, you have multiple ways of getting your files upload. You can use the web client, but Google also provides a command line program so let's try that for this platform. Go to the Google Cloud SDK page (<https://cloud.google.com/sdk/>) and select the Install link.

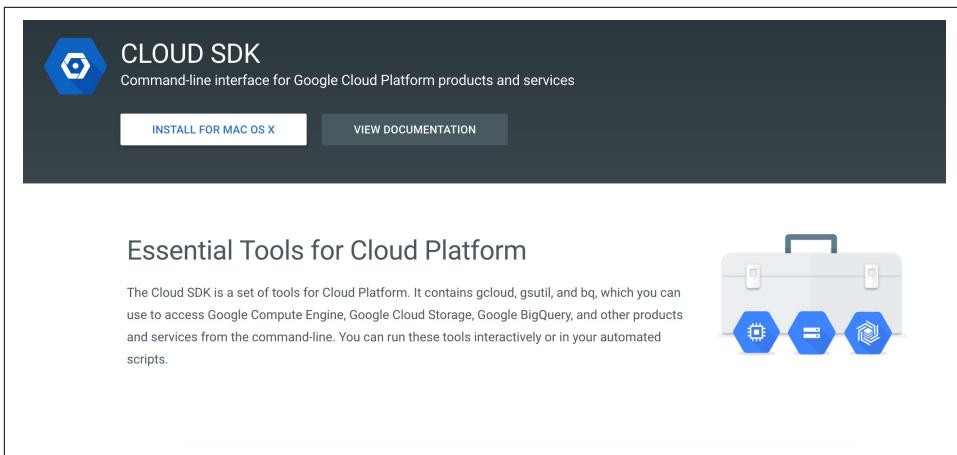


Figure 7-9. The Google Cloud SDK

The SDK provides a really rich set of tools related to the entirety of Google’s Cloud Platform. Everything we’ve done so far (making a project, creating a bucket) could have been done via the SDK as well. We’ll start off by setting permissions for the bucket. We want them to viewable by the Internet as a whole, so we’ll use the command line to assign everyone read permission to the bucket.

```
gsutil defacl ch -u AllUsers:R gs://orabook.raymondcamden.com
```

Let’s explain what this did.

- `defacl ch` is specifying a change to the default ACL (ACL stands for “access control list” and is a basic way of looking at security for resources).
- `-u AllUsers:R` specifies everyone (“AllUsers”) and the read permission.
- Finally we specify the bucket. Again, you must change this.

To upload our site, we’ll use `gsutil` and the `rsync` command. This will copy everything in one folder up to the bucket. Assuming you’re in the folder containing your site, use the following command:

```
gsutil -m rsync -R . gs://orabook.raymondcamden.com
```

Let’s break that down argument by argument.

- `-m` tells the command line to use multiple processes. This is especially handy for an operation moving a lot of files.
- `rsync` is the actual command we are using within the SDK.
- `-R` means to recursively sync all files and folders.
- The period simply means the current directory.

- `gs://orabook.raymondcamden.com` is the name of the bucket. You *must* change this as your bucket will have a different name.

If you refresh in the web client, you'll see your files. What you see here is - obviously - based on what you used for testing. As before, I used the output from the Camden Grounds web site.

The screenshot shows the Google Cloud Platform Storage browser interface. On the left, there's a sidebar with 'Storage' selected, followed by 'Browser', 'Transfer', and 'Settings'. The main area is titled 'Buckets / orabook.raymondcamden.com'. It lists several files and folders:

Name	Type	Last modified	Share publicly
DS_Store	application/octet-stream	8/20/16, 10:34 AM	<input type="checkbox"/>
about.html	text/html	8/20/16, 10:34 AM	<input type="checkbox"/>
coffeees/	Folder	—	⋮
coffeewebsitetemplate.psd	image/vnd.adobe.photoshop	8/20/16, 10:34 AM	<input type="checkbox"/>
cos/	Folder	—	⋮
fonts/	Folder	—	⋮
images/	Folder	—	⋮
index.html	text/html	8/20/16, 10:34 AM	<input type="checkbox"/>
locations.html	text/html	8/20/16, 10:34 AM	<input type="checkbox"/>
menu.html	text/html	8/20/16, 10:34 AM	<input type="checkbox"/>

Figure 7-10. The files now show up in the bucket.

To test your site, you can use this URL, with the bucket name in the middle changed of course: <https://storage.googleapis.com/orabook.raymondcamden.com/index.html>

There's one really important difference here when we compare S3 to Google. The S3 "temp/testing" URL was a root URL. The Google one is a subdirectory under the root domain storage.googleapis.com. If your static site uses root urls, for example: `href="/foo.html"`, then it will work on S3 and not on Google. This isn't a bug per se, and obviously when you have a real domain pointed to the bucket it won't matter.

Finally, if you want to get rid of the index.html in the URL, go to your bucket, click the little dots at the end to bring up the menu:

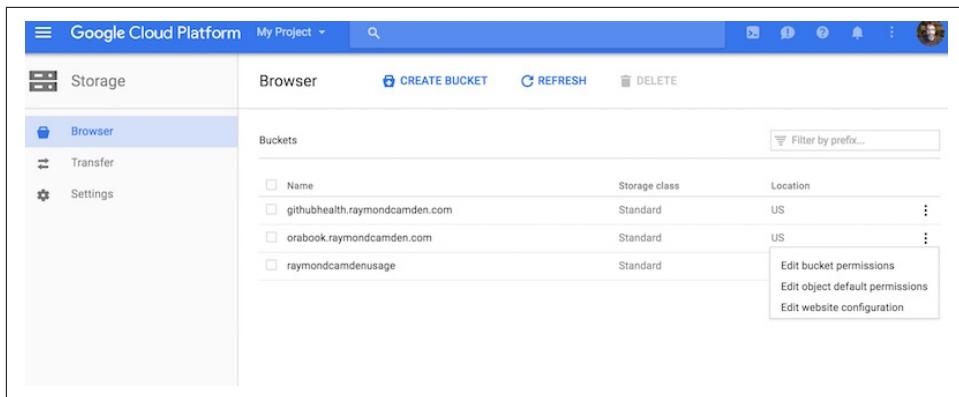


Figure 7-11. The Bucket menu.

Click “Edit website configuration” and in the dialog, simply enter the main page value. Note you can also specify an error page.

The dialog is titled "Configure orabook.raymondcamden.com website". It displays the message "You can access this website at orabook.raymondcamden.com". There are two input fields: "Main page" containing "index.html" and "404 (not found) page" containing "404-page.html". At the bottom are "Save" and "Cancel" buttons.

Figure 7-12. Setting website configuration values for the bucket.

Deploying with Surge

Cloud-based file storage systems like those from Google and Amazon work great for static web sites, but multiple products now offer tools specifically tailored for static web sites and hostings. The list of such products is growing rapidly and in this chapter we'll look at two of them. The first is Surge (<https://surge.sh/>).

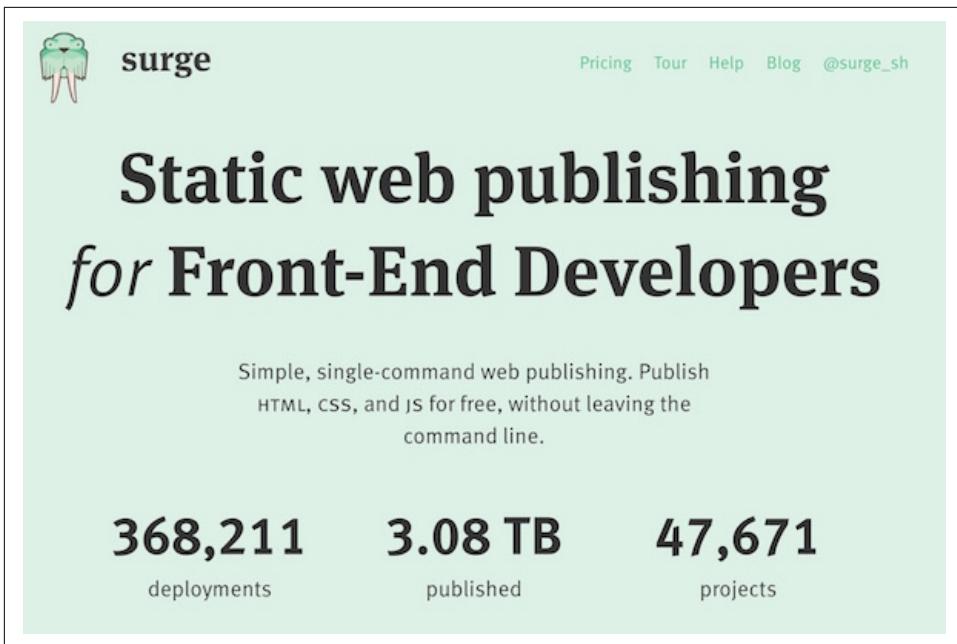


Figure 7-13. The Surge web site.

Surge is a command-line only tool for deploying static web sites. It offers a free tier which is great for testing and a paid tier that offers additional features. At the time this book was written, the paid tier was thirteen dollars a month and included custom SSL and redirects as part of the feature set. You can see the rest of the features and current price on their pricing page: <https://surge.sh/pricing>.

Installing Surge is trivial. If you have npm installed, simply run this in your terminal:

```
npm install -g surge
```

Once installed, you use the `surge` to deploy a static site. Change directories to any of the outputs from previous chapters, or simply make a new folder with an index.html file in it, and then type `surge`. On your *very first* usage of Surge, it will prompt you to login or create an account.

```
→ camdengrounds surge
Welcome to Surge! (surge.sh)
Please login or create an account by entering your email and password:
email: [REDACTED]
```

Figure 7-14. The Surge Login/Registration.

After that, it will prompt for the directory to deploy. It defaults to the current directory.

```
+ camdengrounds surge  
Welcome to Surge! (surge.sh)  
Please login or create an account by entering your email and password:  
email: rcamden@gmail.com  
password:  
project path: /Users/raymondcamden/Desktop/camdengrounds/
```

Figure 7-15. The Surge command line.

Simply hit enter to accept the default. Next, it will prompt for a domain. Notice that it gives you a random domain by default. This is great for testing as you don't have to worry about DNS settings. You can, of course, use a "real" domain name, but for now, just accept the temporary domain provided by the command line.

```
project path: /Users/raymondcamden/Desktop/camdengrounds/  
size: 57 files, 34.1 MB  
domain: earthy-room.surge.sh
```

Figure 7-16. The Surge command line.

Hit enter again, and Surge will begin deploying your site. The command line will provide a simple progress report, and when done, let you know if it succeeded.

```
project path: /Users/raymondcamden/Desktop/camdengrounds/  
size: 57 files, 34.1 MB  
domain: earthy-room.surge.sh  
upload: [=====] 100%, eta: 0.0s  
propagate on CDN: [ ] 0% /fonts/opensans-regular-webfont.s  
propagate on CDN: [=====] 100%  
plan: Free  
users: rcamden@gmail.com  
IP Address: 45.55.110.124  
  
Success! Project is published and running at earthy-room.surge.sh
```

Figure 7-17. Surge has deployed your site.

And that's it! Literally seconds after installing Surge, you can have your site up and running for testing purposes right away. Open your browser up to the domain used in the command line and your site will be there. For updates, you'll want to use the same domain name as before. You can either specify the domain at the command line with the -d flag, or provide the directory to deploy and domain as arguments, like so:

```
surge ./ earthy-room.surge.sh
```

Yet another option is to create a file called `CNAME` that contains the domain name. This file should be in the same directory as your site and it will *not* be deployed. Be sure to *just* the domain name, like `earthy-room.surge.sh`, and not the URL. In other words, don't include `http://`.

Another cool feature of Surge is that it automatically supports “clean” URLs. Any URL that ends in HTML can have the extension left off. For our testing we used the “Camden Grounds” site from chapter two. One of the URLs is the menu. You can see it at <http://earthy-room.surge.sh/menu.html>. But you can also view it at <http://earthy-room.surge.sh/menu>. You’ll want to update the links in your HTML of course to leave off the HTML.

Surge provides support for custom 404 files by simply looking for a file called `404.html`. If you don’t provide one, and request a file that doesn’t exist, you’ll get a Surge branded 404 page:

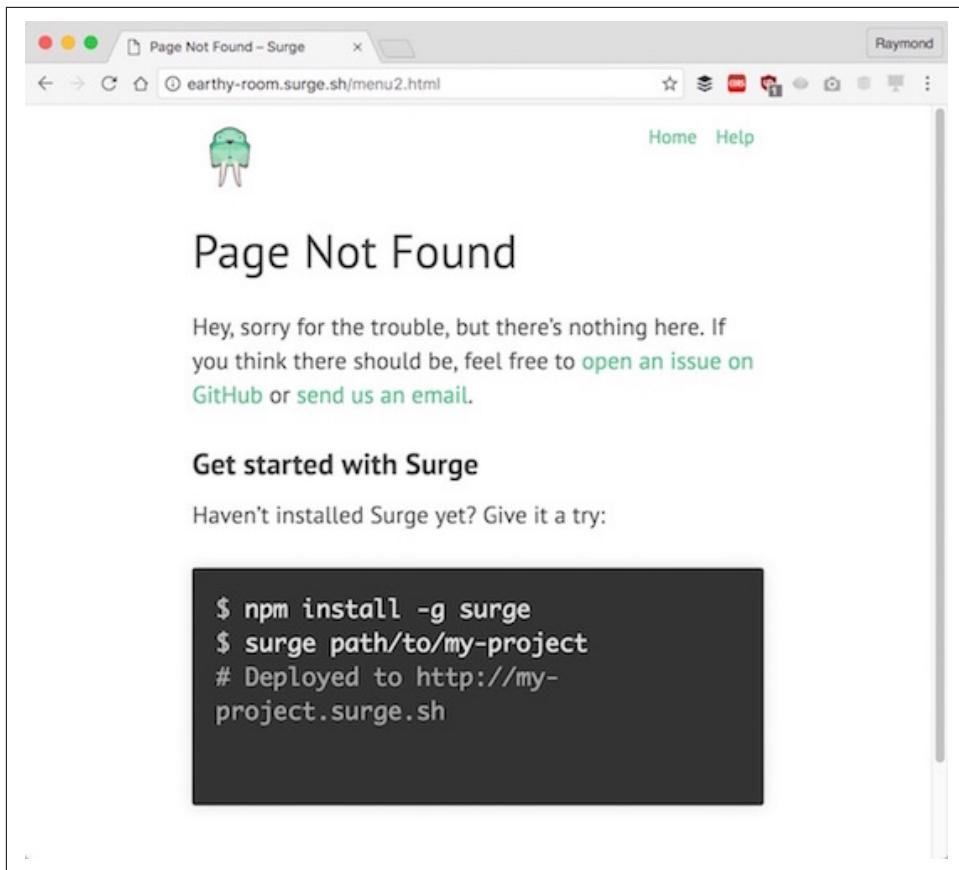


Figure 7-18. The default 404 page.

Once you've built and deployed a file called 404.html, Surge will automatically use it when it can't find a requested file.

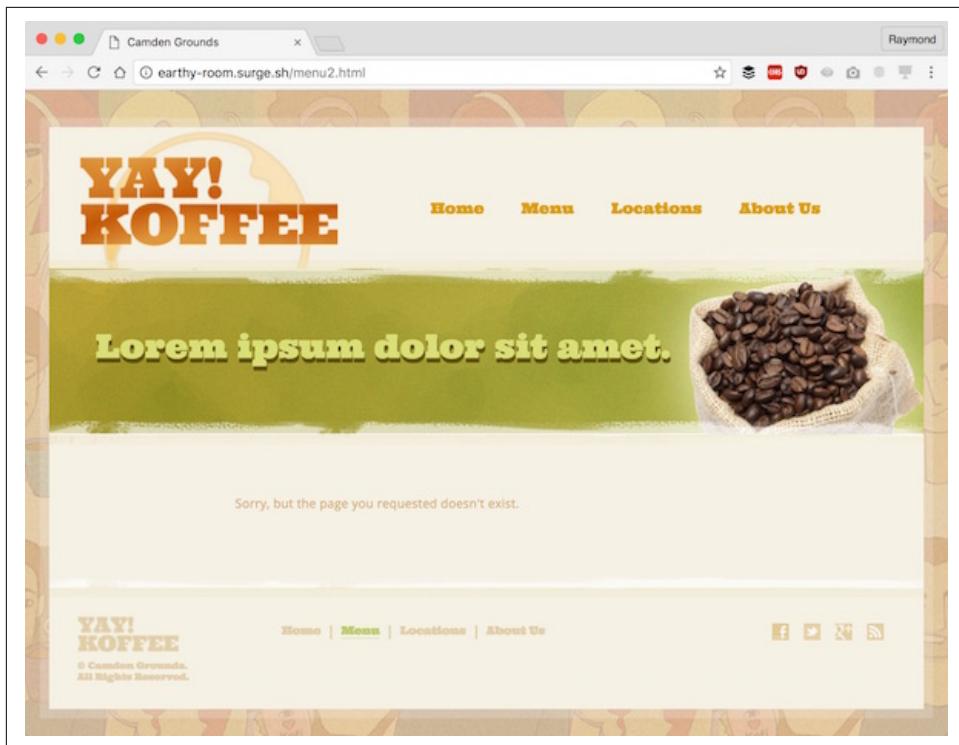


Figure 7-19. A custom 404 page.

As you can probably tell, the entirety of Surge is run at the command line. So for example, if you've forgotten what you've deployed, you can use `surge list` to see your current sites:

```
→ camdengrounds surge list

Surge - surge.sh

email: rcamden@gmail.com
token: ****

earthy-room.surge.sh
```

Figure 7-20. Your current list of Surge sites.

As another example, if you need to delete a site, you do so with the `teardown` command:

```
surge teardown earthy-room.surge.sh
```

So how do premium features work? One of the features you have to pay for is custom URL redirects. Custom redirects allows you to define rules for mapping one URL to another. So for example, imagine you deployed a file called `menyou.html`. Now imagine you miss this error for months. If you just rename it, then people who bookmarked the old URL will get the 404 page instead. With custom redirects, you can make it so that any request for the old URL is automatically sent to the new one.

Surge supports this by using a special file called `ROUTER`. (Note that it is all caps like the special CNAME file.) Within this file, every line represents one mapping. A mapping is defined with a status code, an old URL pattern, and a new URL. Status codes can be either 301 (“Moved Permanently”) or 307 (“Moved Temporarily”). Here’s an example that would fix the problem described above.

```
301 /menyou.html /menu.html
```

Note that you only supply the portion of the URL after the domain name. You can also match a “general” format. So if your site had a subdirectory called `blog` with URLs like so:

```
http://mysite.com/blog/welcome-to-our-blog.html
```

You’ve decided you want to rename the directory from `blog` to `news`. In order for the old URLs to work, you could either create one entry in `ROUTER` for every single blog entry, or use a generic pattern:

```
301 /blog/:title /news/:title
```

In this example, `:title` will match any string and be used in the new URL as well.

Once you add a `ROUTER` file to your project, Surge will recognize this as a premium feature and prompt you to upgrade when you deploy.

```
+ camdengrounds surge
Surge - surge.sh

    email: rcamden@gmail.com
    token: *****
project path: /Users/raymondcamden/Desktop/camdengrounds/
    size: 59 files, 1.2 MB
    domain: earthy-room.surge.sh
    upload: [=====] 100%, eta: 0.0s
    upload: [=====] 96%, eta: 0.0s
Project requires the Plus plan. $13/mo (cancel anytime). This plan provides...
- Platform setup at *.earthy-room.surge.sh
- Custom SSL
- Force http to https
- ROUTER for custom redirects
- CORS for resource sharing
- AUTH support for private deploys

Please enter your payment info... [all payment transfers are PCI compliant]
card number: 
```

Figure 7-21. Upgrading your Surge site.

This is a one time process. After entering this information, Surge will remember it and won't prompt you again.

Surge is a great, and simple to use, deployment tool for static sites. Be sure to check the site for a complete list of premium features.

Deploying with Netlify

For our next static site publishing and hosting service, we'll take a look at Netlify (<https://www.netlify.com>). Netlify is my favorite static site hosting service and powers my own blog (<https://www.raymondcamden.com> which has near six thousand blog posts).

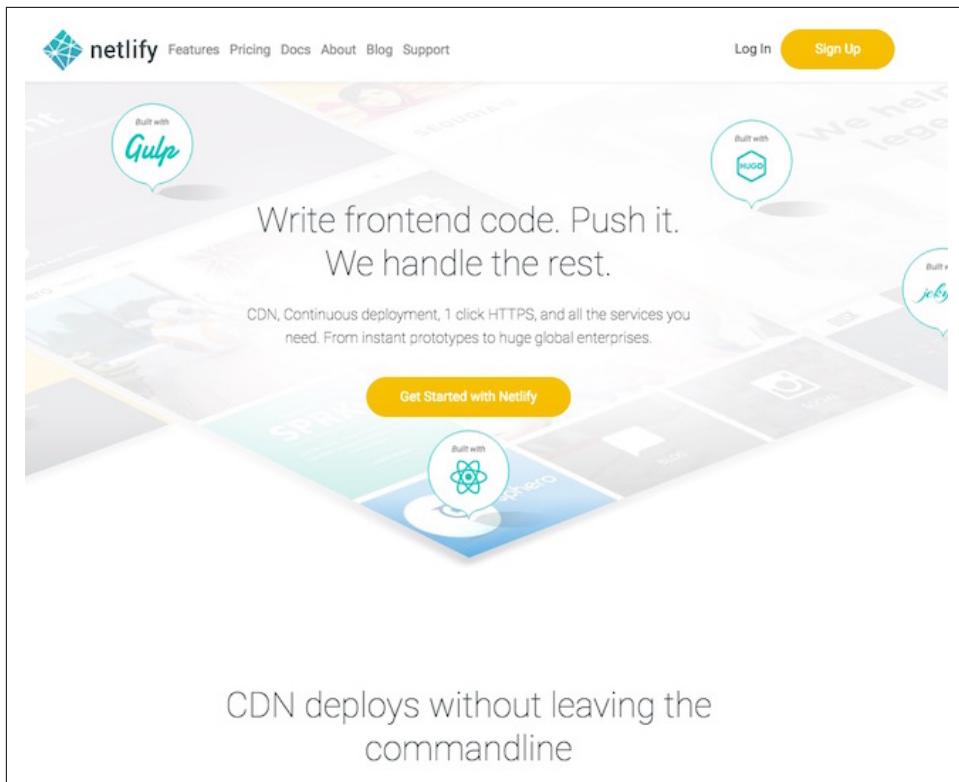


Figure 7-22. The Netlify web site.

Like Surge, Netlify is command-line driven, but comes with a web-based dashboard as well. Currently Netlify comes with five different pricing levels (<https://www.netlify.com/pricing/>). The lowest level is free and should be more than adequate for testing. The “Pro” level (currently \$49) can be used for free by open source projects.

Netlify really shines in terms of performance. They use a CDN with multiple endpoints around the planet as well as a speedy DNS and heavy caching. You can also easily add (for free) SSL to your site. There’s many more features as well, but my favorite is probably one of the simplest - form handling. We covered this in chapter 5 using external services, but Netlify has a generic forms handler built into the service itself. We’ll see an example of this later in the chapter.

For a full list of their features, see the web site (<https://www.netlify.com/features/>), but for now, begin by installing the command line:

```
npm install -g netlify-cli
```

This will install the `netlify` command line program. Go into your static site (for this example we'll use the “Cat Blog” from chapter 4) and run:

```
netlify deploy
```



The first time you do this you'll be prompted to authenticate with your browser. This is a one time process and you can even manage multiple logins if you have to work with different Netlify accounts.

The command line will notice that the site is unknown to the platform and will prompt you to create a new one:

```
➔ Desktop cd catblog
➔ catblog netlify deploy
? No site id specified, create a new site (Y/n)
```

Figure 7-23. Deploying a new site.

It will then prompt for the directory to deploy (and as it defaults to the current directory, you can just hit enter).

Netlify will then do its thing (with a nice little progress bar) and when complete, will give you the URL for your new site. Like Surge, it defaults to selecting a random domain for you, and obviously you can use a “real” domain before going live.

```
➔ catblog netlify deploy
? No site id specified, create a new site Yes
? Path to deploy? (current dir)
Deploying folder: /Users/raymondcamden/Desktop/catblog

Deploy is live (permalink):
http://57c870f56686743d5a35f56c.robber-mousedeer-87783.netlify.com

Last build is always accessible on http://robber-mousedeer-87783.netlify.com
```

Figure 7-24. The site is now deployed.

Notice right away something incredibly cool with Netlify. It deployed your site as well as a unique one just for this particular version of the site. Netlify automatically gives you a historical list of views for your content. That means if something goes wrong, you can examine the previous versions to try to nail down when things went haywire. You can even roll back (via the web admin) to quickly correct the issue. To see that web administrator in action, simply run `netlify open`:

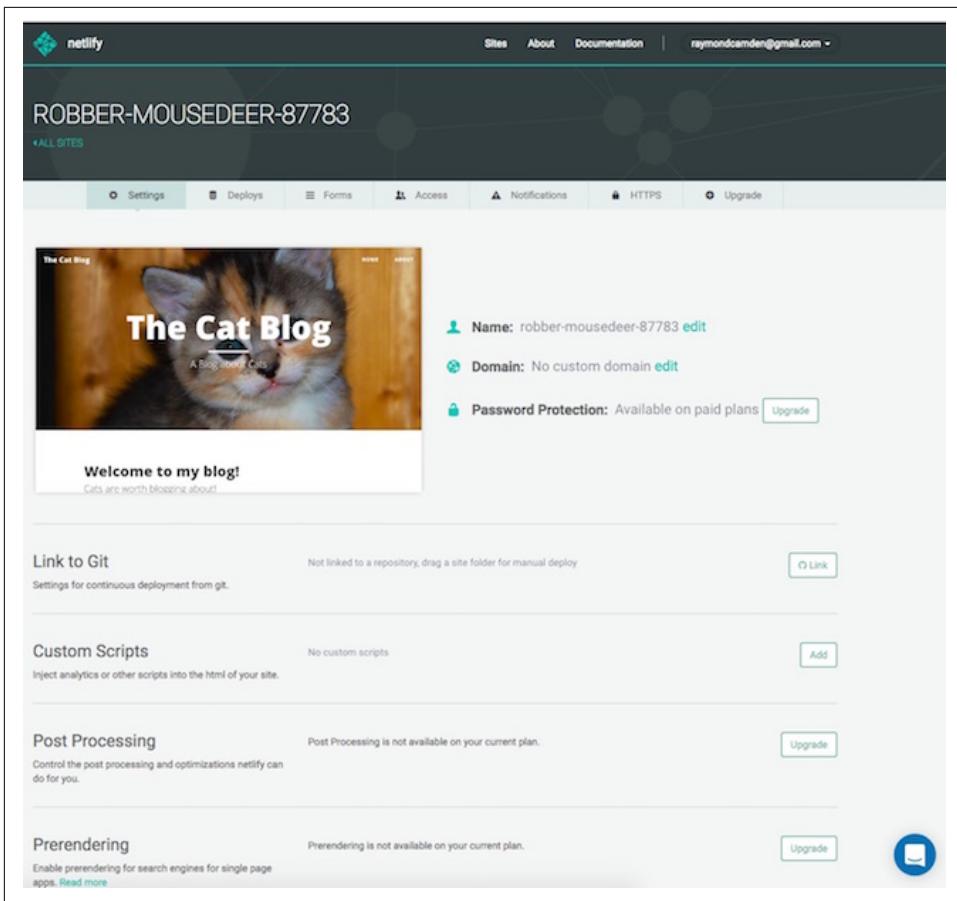


Figure 7-25. Netlify's administrator for your site.

Right away you can see there's *quite* a bit of information here. Many of the options will require you to upgrade, but you can at least see what's available and decide if they are worth the price. You can see the version history by clicking on **Deploys**. Here is what it looks like after a second run of the command line.

The screenshot shows the deployment history for the site 'ROBBER-MOUSEDEER-87783'. At the top, there are navigation links: 'ALL SITES', 'Settings' (selected), 'Deploys' (highlighted in blue), 'Forms', 'Access', and 'Notifications'. Below this, the heading 'Currently Published' is displayed. A card shows a deployment from 'PUBLISHED' on September 1, 2016, at 1:39 PM, with a 'VIEW LOG' button. Another card shows a deployment from 'READY' on September 1, 2016, at 1:39 PM, with a 'VIEW LOG' button. A third card shows a deployment from 'READY' on September 1, 2016, at 1:18 PM, with 'VIEW LOG' and 'PUBLISH' buttons.

Figure 7-26. Version history for the site.

One feature available on the free plans is hooks and notifications. It would be nice if your client could know everytime the site was updated. Click on “Notifications”, “Add Notification”, and then “Email notification”:

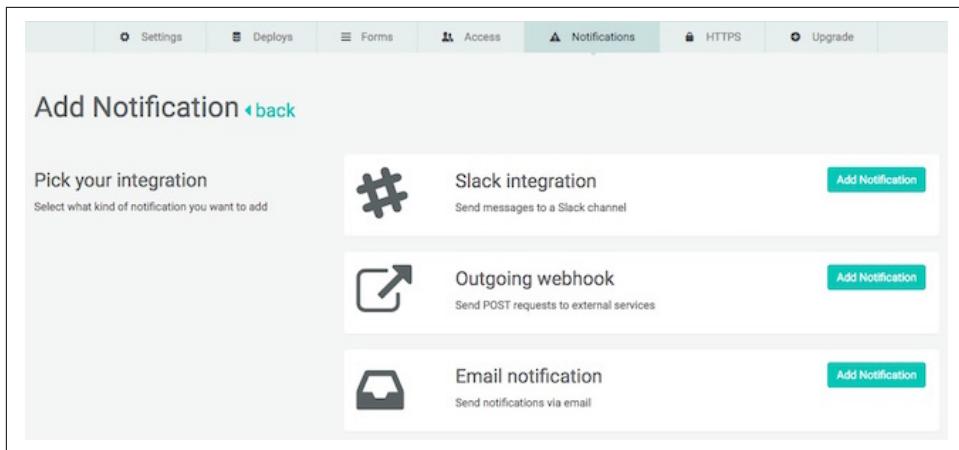


Figure 7-27. Adding a Notification.

On the next screen, select “Deploy succeeded” for the event and then enter your email address.

A screenshot of a configuration dialog for an 'Email notification'. The title is 'Email notification'. Under 'Event to listen for', 'Deploy succeeded' is selected. Under 'Email to notify', 'raymondcamden@gmail.com' is entered. At the bottom are 'Save' and 'Cancel' buttons.

Figure 7-28. Setting up the “Deploy succeeded” Notification.

Click save and the notification will be created. Back at the command line, run `netlify deploy` again. When the deploy is complete, you’ll get an email notification that includes both the main domain as well as the versioned URL.

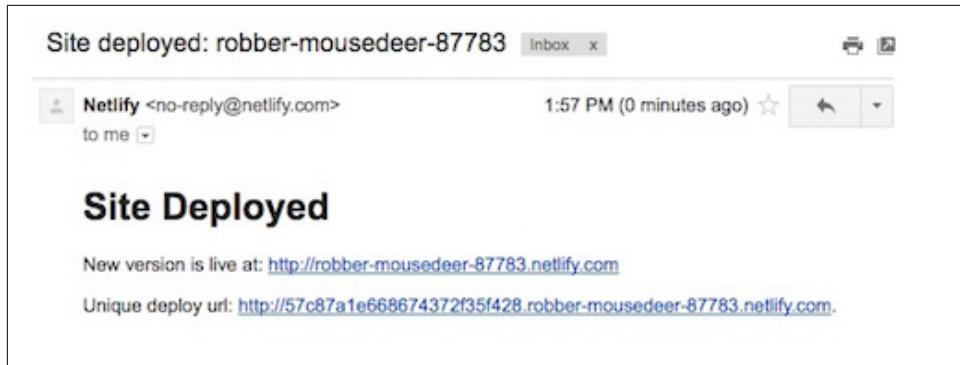


Figure 7-29. The email notification from Netlify.

Another free tier feature, redirects, are done by using a particularly named file in the root of your site: `_redirects`.

The file should contain one line per redirect with the format being:

old path new path

So a real example could look like so:

/kats /cats

Your file can also include comments by using a hashmark as the first character:

```
#We thought this spelling would be cute. It wasn't.  
/kats                          /cats
```

The default HTTP status code is 301, but you can specify another one by adding it to the end of the line.

```
#We thought this spelling would be cute. It wasn't.  
/kats                          /cats                          302
```

And then finally, you can also use “slugs” in the URLs to match specific patterns of URLs.

```
#Old URL for cat adoptions  
/adoptions/:year/:breed                          /catsneedinghomes/:year/:breed
```

Now that we've looked at some of the free features, let's look at a few of the features available on the paid tier. The easiest one to demonstrate, and one of the most useful features, is automatic form processing. Let's begin by adding a simple contact form to the cat blog we uploaded earlier. The cat blog was created in chapter 3, “Building a

Blog”, and makes use of the Jekyll static site generator. If you skipped that chapter, you may want to quickly read it to get an idea of how it works.

The cat blog made use of a template (“Clean Blog”) that had a contact form done already. I literally just took their form and added it to our local Jekyll blog.

```
---
layout: page
title: "Contact"
description: "Contact page."
header-img: "img/about-bg.jpg"
---

<p>Please send us your feedback. We care a lot.</p>

<form method="post">
<div class="row control-group">
    <div class="form-group col-xs-12 floating-label-form-group controls">
        <label>Name</label>
        <input type="text" class="form-control" placeholder="Name" id="name"
               required data-validation-required-message="Please enter your name.">
        <p class="help-block text-danger"></p>
    </div>
</div>
<div class="row control-group">
    <div class="form-group col-xs-12 floating-label-form-group controls">
        <label>Email Address</label>
        <input type="email" class="form-control" placeholder="Email Address" id="email"
               required data-validation-required-message="Please enter your email address.">
        <p class="help-block text-danger"></p>
    </div>
</div>
<!-- parts deleted -->
<br>
<div id="success"></div>
<div class="row">
    <div class="form-group col-xs-12">
        <button type="submit" class="btn btn-default">Send</button>
    </div>
</div>
</form>
```

To keep the code listing a bit shorter, a few fields were removed. You can find the complete source code in the book’s GitHub repository. Note that the form has no action. After generating the static version and deploying to Netlify, you can see the form in action.

If you submit the form, however, you’ll get an error:

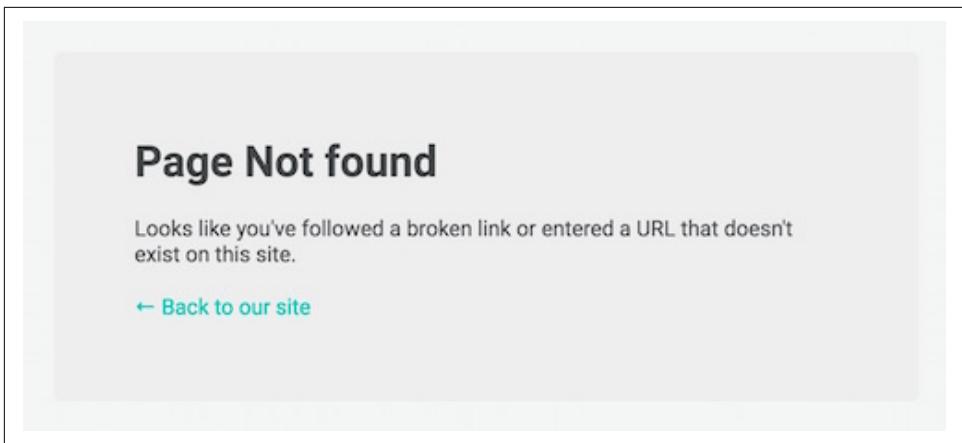


Figure 7-30. The form doesn't work quite well yet.

Correcting this is incredibly easy. First, add a `netlify` attribute to your form tag:

```
<form method="post" netlify>
```

Then - add a proper action to the tag. In this case, we're going to point to a new "thank you" page. (Again, the source of this page may be found on GitHub, but it's just a quick 'thank you' message.)

```
<form method="post" action="thankyou" netlify>
```

Don't forget to redeploy the site, and now when you submit the form, you'll be automatically redirected to the thank you page. So how do you see your submissions? You've got a couple of options. On the Netlify dashboard for your site, click on the Forms tab:

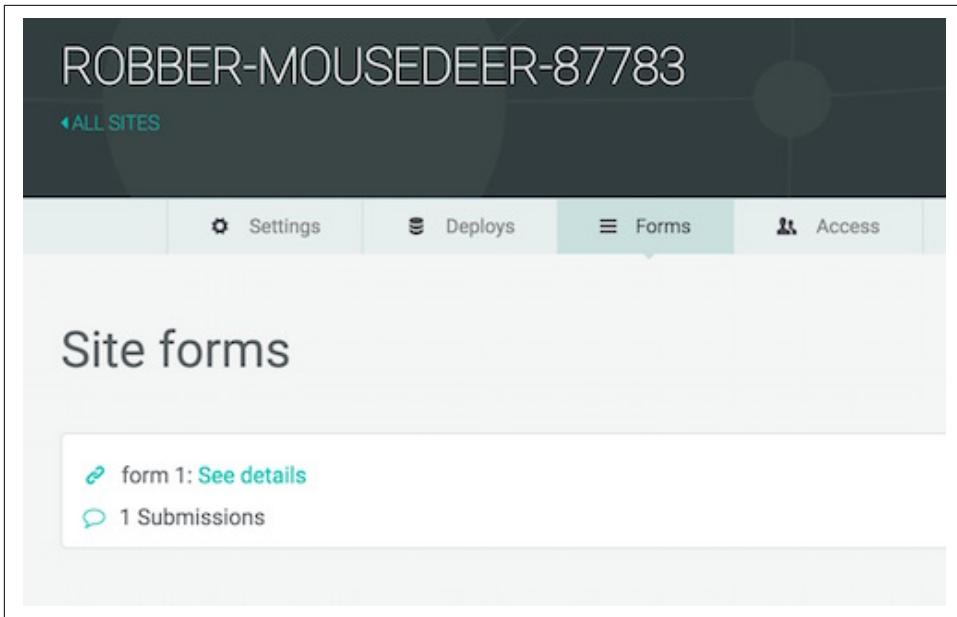


Figure 7-31. Netlify's Form Dashboard

Notice how Netlify refers to the form as “form 1”? You can correct this by adding a `name` attribute to the form tag:

```
<form method="post" name="Contact Form" action="thankyou" netlify>
```

Now submissions will be labelled correctly:

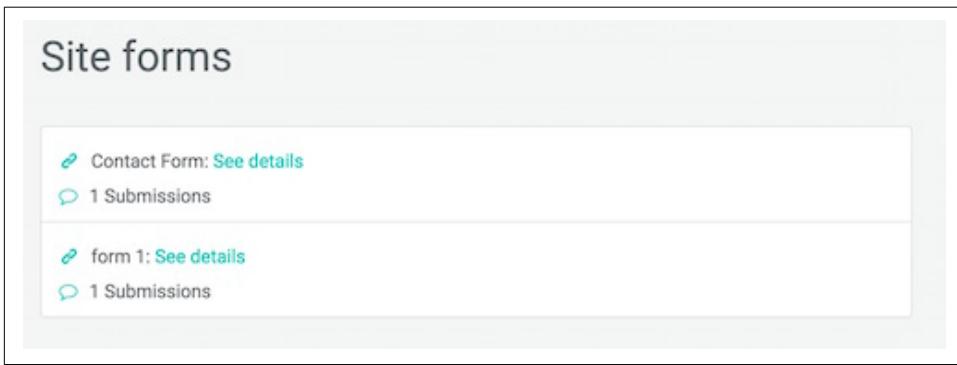


Figure 7-32. How Netlify can identify your form submissions

You can view your form submissions by clicking the “Details” link in the dashboard, and Netlify provides an API to fetch form results as well, but most likely you simply want to have the submissions sent to you. In your dashboard, go to Notifications and

click to add a new Email notification. For the event, select “New Form Submission”. Notice how the dashboard recognizes what form’s you’ve used. Simply add in your email address and save the notification.



Figure 7-33. Setting up email notifications for your form

Submit your form again, and in a few moments you’ll get an email with the contents of your form.

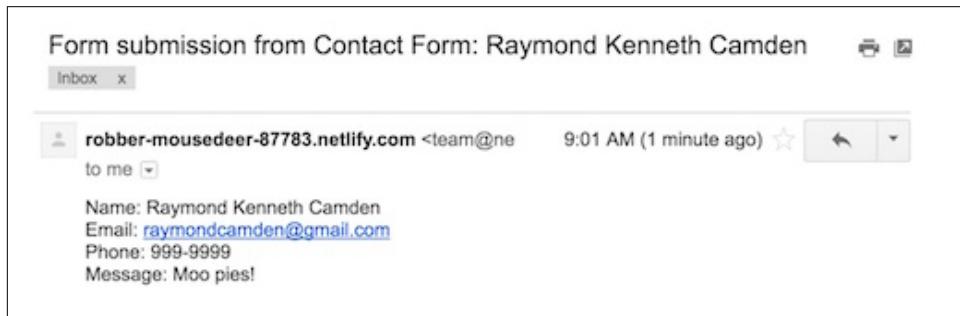


Figure 7-34. Form submissions are now emailed to your account.

Finally, let’s look at one of the most impressive features of Netlify - automatic processing. On your dashboard, you may have noticed a “Post Processing” section. If you’ve upgraded to a paid plan, you can enable multiple different processing options for your site:



Figure 7-35. Setting up post processing for your site

Netlify can automatically perform multiple optimizations for your site's CSS, JavaScript, and images. To be clear - these optimizations are 100% automatic. You literally click a checkbox and when you deploy, Netlify will optimize what you've asked it to and your users immediately get the benefit. While these are all things you could do yourself, let's be honest, having it done by Netlify while keeping your own code nice and simple, is an incredible feature.

How well this works will depend of course. In my specific case, the template used by the "Cat Blog" had optimized images and mostly optimized CSS and JavaScript already. In my testing I saw a savings of about 18k. This was approximately 5% of the total page load for the home page. Not a huge savings, but for twenty seconds of work on my part, that's a big win. Going forward, as new images and content are added, I can rest assured that Netlify will handle the optimizations for me.

There's quite a bit more to Netlify than covered here and you're encouraged to peruse the docs (<https://www.netlify.com/docs/>) for a full list of features. One feature in particular will be of interest to sites with a large or complex review process - deploy contexts. This allows you to deploy different versions of your site with unique settings to allow for previews, QA, and reviews before updating your main site. By connecting your Netlify site to a Git repository, the *entire* publication process can be easily tied to your source control for completely automated updates.

Summary

Hopefully you've seen that moving your static site from a local environment to a live web site is a fairly simple process. As static site generators become more popular, you'll see even more options for deployments in the future.

Migrating to a Static Site

By this point, hopefully you are thoroughly convinced that creating a static site with one of the many static site generators is a worthwhile and viable option. Perhaps you are even thinking that it would be the perfect option for an existing CMS site you already maintain, but you're concerned about the complexity of moving all of the site's content from the CMS to static files.

The good news is that there are lot of solutions that exist for this problem. While the migration may still be quite a bit of work, depending on the size and complexity of your existing site, these migration tools can massively decrease the effort and difficulty.

Of course, all of this depends on two factors - what CMS engine you are coming from and what static site generator you intend to use. Most of the major static site generators have tools to import content from various CMS, but they don't cover every available option, and some static site generators have many more available importers than others.

Since there are a lot of factors involved in this sort of project, in this chapter, I'll aim to offer an overview of what might be involved, how some of these importers work and what available options are out there. First, let's start off by examining one common scenario.

Migrating from Wordpress to Jekyll

Since its introduction in 2003, Wordpress has become perhaps the most dominant CMS. According to W3Techs, Wordpress is used by 26.8% of the top 10 million websites as determined by Alexa — essentially one-quarter of the entire web runs Wordpress. For comparison, the next largest CMS would be Joomla at 2.8%.

Suffice it to say, there's a reasonable chance that, if you are migrating to a static site engine, you are migrating from Wordpress. This explains why many static site generators offer a Wordpress importer (in fact, many *only* offer a Wordpress importer). Let's look at how you might use the importer provided by Jekyll.

Jekyll provides a very comprehensive [list of available importers](#) (24 by my count) from a wide array of CMS and other data formats. All of these run off of [Jekyll Import](#) Ruby gem. If you've already installed Jekyll, you'll have had to have Ruby Gems installed on your machine, so you can easily install the Jekyll Import gem via the Terminal/Command Prompt.



Installing Jekyll

If you haven't installed Jekyll yet, please refer to Chapter 3 on Building a Blog for more detailed instructions.

```
gem install jekyll-import
```

Each importer also has their own set of additional dependencies. For instance, the Wordpress importer also requires the gems for [unidecode](#), [sequel](#), [mysql2](#) and [htmlentities](#). At this point, you should have Ruby Gems installed, so you should be able to install all of these with one simple command line statement.

```
gem install unidecode sequel mysql2 htmlentities
```



Wordpress running on MAMP

A lot of developers (myself included) use [MAMP](#) to run their local Wordpress development. As someone who is not really a PHP developer, it's option of a one-click install of all of the necessary pieces (MySQL, Apache, PHP) is compellingly simple.

However, MAMP does not come with the full MySQL installation that is necessary to run the Wordpress importer for Jekyll. This will cause the gem installation above to fail.

After a lot of research and trial and error, I found that the simplest solution was to download and install the full [MySQL](#). This will not interfere with your MAMP copy of MySQL. Once installed, the gems should install without error.

In order to make the importer use the MAMP copy of MySQL rather than the full MySQL installation, in the configuration (which we will look at in more detail below) you need to point it at the socket for MAMP. For example, on my Mac running OS X El Capitan, the socket was located at `/Applications/MAMP/tmp/mysql/mysql.sock`.

Once everything is installed, you can run the importer, providing the configuration via the command line. The easiest way to do this (in my opinion) is to first copy the [configuration from the documentation](#) into a text editor and edit it there before pasting it into the Terminal/Command Prompt.

Here is the documentation's configuration.

```
ruby -rubygems -e 'require "jekyll-import";
JekyllImport::Importers::WordPress.run({
  "dbname"    => "",
  "user"      => "",
  "password"  => "",
  "host"       => "localhost",
  "socket"     => "",
  "table_prefix" => "wp_",
  "site_prefix"  => "",
  "clean_entities" => true,
  "comments"   => true,
  "categories" => true,
  "tags"        => true,
  "more_excerpt"  => true,
  "more_anchor"   => true,
  "extension"   => "html",
  "status"      => ["publish"]
})'
```

It's important to note that each of the configuration settings above includes the default value. Thus, if the default value is sufficient for you (for example "host" => "localhost"), then you can feel free to remove it.

For example purposes, I went ahead and set up a Wordpress variation of my [static site samples](#) example site, which is a Adventure Time! fan page. For this example, I did not implement a full copy of the design in Wordpress, I simply focused on adding the blog posts.



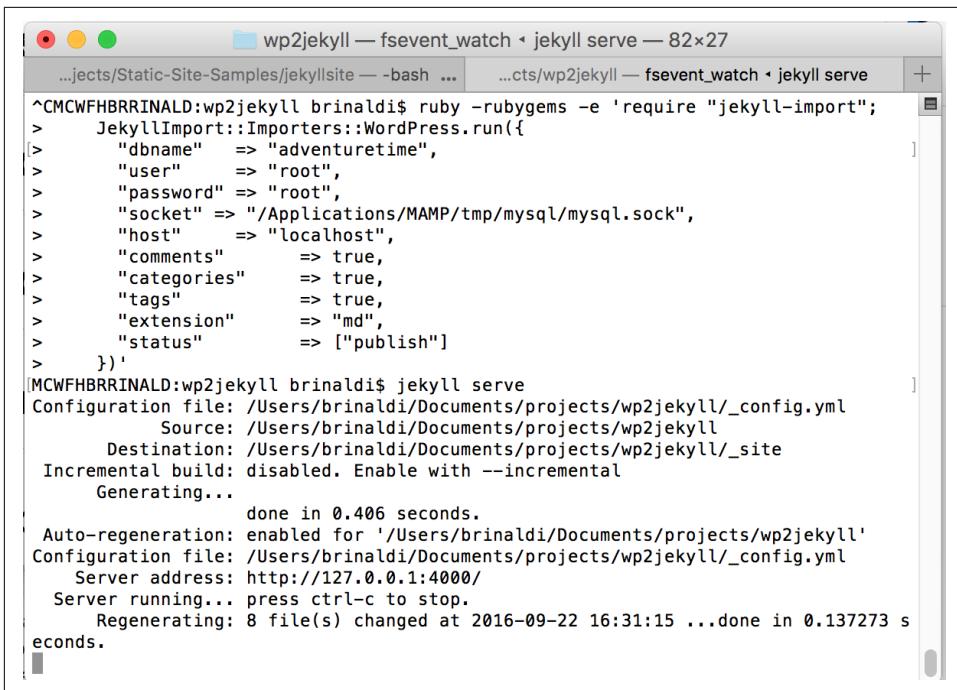
Figure 8-1. A simple Wordpress example site

wordpress_site.png

After creating a new Jekyll site, I then opened the Terminal/Command Prompt within the Jekyll site folder and ran the importer with the configuration as follows:

```
ruby -rubygems -e 'require "jekyll-import";  
JekyllImport::Importers::WordPress.run({  
  "dbname"    => "adventuretime",  
  "user"      => "root",  
  "password"  => "root",  
  "socket"    => "/Applications/MAMP/tmp/mysql/mysql.sock",  
  "host"      => "localhost",  
  "categories" => true,  
  "tags"       => true,  
  "extension"  => "html",  
  "status"     => ["publish"]  
})'
```

The only thing “out of the ordinary” here is the socket setting, which, as I explained above, was necessary in order to connect to the MySQL database running on MAMP.



```
wp2jekyll — fsevent_watch ↵ jekyll serve — 82x27
...jects/Static-Site-Samples/jekyllsite — -bash ... ...cts/wp2jekyll — fsevent_watch ↵ jekyll serve
^CMCWFHBRRINALD:wp2jekyll brinaldi$ ruby -rubygems -e 'require "jekyll-import";
>   JekyllImport::Importers::WordPress.run({
|>     "dbname"    => "adventuretime",
|>     "user"      => "root",
|>     "password"  => "root",
|>     "socket"    => "/Applications/MAMP/tmp/mysql/mysql.sock",
|>     "host"      => "localhost",
|>     "comments"  => true,
|>     "categories"=> true,
|>     "tags"       => true,
|>     "extension" => "md",
|>     "status"     => ["publish"]
|>   })'
|MCWFHBRRINALD:wp2jekyll brinaldi$ jekyll serve
Configuration file: /Users/brinaldi/Documents/projects/wp2jekyll/_config.yml
  Source: /Users/brinaldi/Documents/projects/wp2jekyll
  Destination: /Users/brinaldi/Documents/projects/wp2jekyll/_site
Incremental build: disabled. Enable with --incremental
  Generating...
    done in 0.406 seconds.
  Auto-regeneration: enabled for '/Users/brinaldi/Documents/projects/wp2jekyll'
Configuration file: /Users/brinaldi/Documents/projects/wp2jekyll/_config.yml
  Server address: http://127.0.0.1:4000/
  Server running... press ctrl-c to stop.
  Regenerating: 8 file(s) changed at 2016-09-22 16:31:15 ...done in 0.137273 s
econds.
```

Figure 8-2. Running the Wordpress importer

wp2jekyll_cli.png

As you can see from the image above, all 8 posts from my sample site were imported as .html posts into the _posts folder of my Jekyll site.

Posts

Sep 12, 2016

[Escape from the Citadel \(Season 6\)](#)

Sep 12, 2016

[Wake Up \(Season 6\)](#)

Sep 12, 2016

[James II \(Season 6\)](#)

Sep 12, 2016

[The Tower \(Season 6\)](#)

Sep 12, 2016

[Sad Face \(Season 6\)](#)

Sep 12, 2016

[Breezy \(Season 6\)](#)

Figure 8-3. A default Jekyll site with the imported Wordpress posts

imported_wordpress_site.png

It's important to note that the imported posts were HTML (regardless of the extension setting). They also include a ton of metadata added to the Jekyll front matter that was pulled from the Wordpress site including author details and even the Wordpress ID. All of these values could be useful when trying to recreate the site as static.

I did run into some formatting issues that needed to be cleaned up. For example, Wordpress' editor allowed for (or generated) some HTML that was technically invalid (for example, `<p></p>`) that the importer didn't handle. Nonetheless, even accounting for these sorts of issues, using the importer could potentially save you a lot of time when moving from Wordpress to Jekyll.

Other Migration Options

Jekyll has the most exhaustive list of import tools, but other static site engines also offer some.

Hugo

Hugo offers importers from CMS such as Wordpress and Drupal, blog hosting services such as Blogger and Tumblr and even other static site generators like Jekyll and Octopress (a Jekyll-based tool).

Let's take a really quick look at how to use the Wordpress migration tools for Hugo. The nice part about this tool is that it is actually built as a plugin for Wordpress, making the installation relatively simple.

First, download the code from the [GitHub repository](#) by clicking the “Clone or Download” button and choosing “Download Zip.” Unzip the file and place the unzipped folder in your Wordpress /wp-content/plugins directory. If you go to your site’s admin, you should now see the plugin listed under plugins — go ahead and activate it.

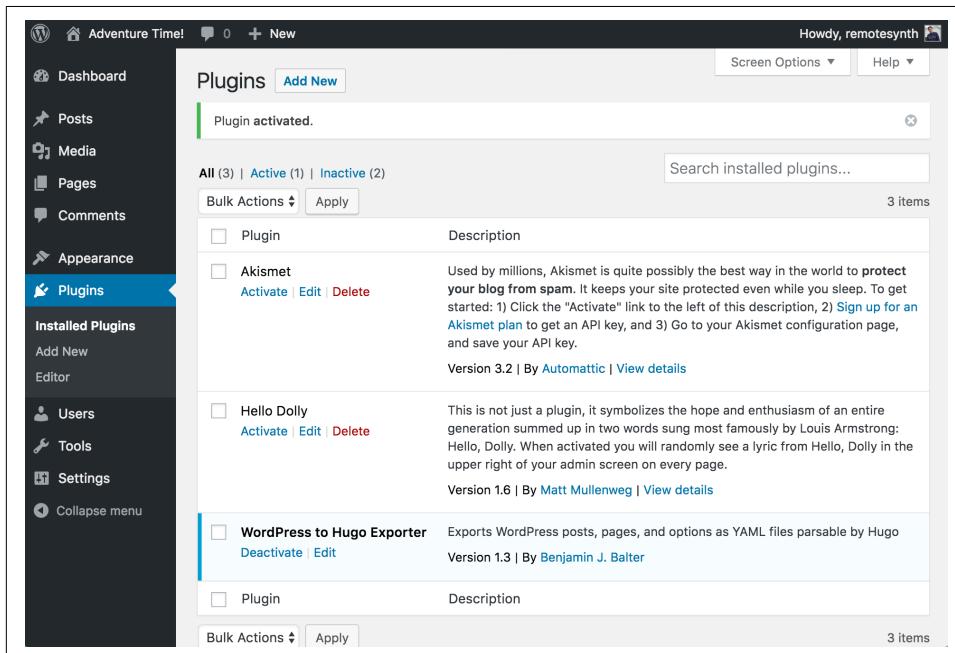


Figure 8-4. The WordPress to Hugo exporter plugin activated.

Once activated, you should now have an “Export to Hugo” option under your Tools menu. Selecting that will download a zip file containing the posts, pages and files you’ll need to copy into Hugo.

Assuming that you have a new Hugo site already set up, you can simply copy the contents of the export post folder into your Hugo site’s content folder. All your posts should be there converted nicely into Markdown.

Another nice feature of the plugin is that it also includes the images and uploads from your Wordpress site's `wp-content/uploads` folder. You can place this inside the `static` folder under your Hugo site to access any images and uploads from your new content.

This allowed me to easily add in the thumbnail metadata for the **Robust** Hugo theme to the front matter for my posts as:

```
thumbnail: "wp-content/uploads/2016/09/wakeup.jpg"
```

...which added back in my header images for my sample Adventure Time! site.

The screenshot shows a Hugo-powered website for 'ADVENTURE TIME!' with the following layout:

- Header:** ADVENTURE TIME!
- Left Column:**
 - Thumbnail:** *Escape from the Citadel* (Season 6)
 - Title:** Escape from the Citadel (Season 6)
 - Date:** Sep 12, 2016
 - Description:** Finn and Jake follow the Lich to the Citadel, where Finn meets his long lost dad. Finn and Jake hitch a ride on a Citadel Guardian after he imprisons and transports The Lich to the Citadel.
- Center Column:**
 - Thumbnail:** *Wake Up* (Season 6)
 - Title:** Wake Up (Season 6)
 - Date:** Sep 12, 2016
 - Description:** Prismo helps Finn and Jake access the Crystal Citadel to meet Finn's father. The episode begins with Jake partying at Time Room with Prismo and The Cosmic Owl, Death, Peppermint.
- Right Column:**
 - Section:** LATESTS
 - Thumbnail:** *Escape from the Citadel* (Season 6)
 - Title:** Escape from the Citadel (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *Wake Up* (Season 6)
 - Title:** Wake Up (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *James II* (Season 6)
 - Title:** James II (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *The Tower* (Season 6)
 - Title:** The Tower (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *Sad Face* (Season 6)
 - Title:** Sad Face (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *Breezy* (Season 6)
 - Title:** Breezy (Season 6)
 - Date:** Sep 12, 2016
 - Section:** LATESTS
 - Thumbnail:** *Food Chain* (Season 6)
 - Title:** Food Chain (Season 6)
 - Date:** Sep 12, 2016
- Bottom:** CATEGORY episodes

Figure 8-5. My sample site after importing it into Hugo displayed using the **Robust** theme

`hugo_wordpress_sample.png`



Building a Complete Hugo Site

For more detailed instructions on working with Hugo, please refer back to Chapter 4 which details building a complete documentation site using Hugo.

Middleman

Middleman is another very popular Ruby-based static site generator. Unfortunately, Middleman does not maintain an official list of migration tools, though there is a

Wordpress to Middleman conversion tool (there are guides for other CMS, but I was unable to locate any other tools).

The importer is available as a Ruby gem (as is Middleman). So, first, install the tool:

```
gem install wp2middleman
```

From within your Wordpress site, go to Tools > Export to perform a standard export of all content into a Wordpress XML file. Next, open the Terminal/Command Prompt in the folder where you have the export XML and run the tool. For example, with an export file named `adventuretime.wordpress.2016-10-11.xml`, I ran:

```
wp2mm adventuretime.wordpress.2016-10-11.xml
```

This generated a folder named `export` that contained all of the Wordpress posts as Markdown with front-matter ready for Middleman.

Hexo

Hexo is a popular JavaScript-based static site generator that runs on Node.js. Hexo does list a number of migration tools within its [plugins directory](#). These include migrators for **Wordpress**, **Blogger** and **Joomla** as well as from a [standard RSS feed](#), meaning you can import from most any CMS that has an RSS output.

Harp

Harp is another JavaScript-based static site generator. It also does not maintain an official list of migration tools, but the community has created some including for importing from **Wordpress** or importing from **Jekyll**, although it is worth noting that the Wordpress importer requires being configured in Python (assuming you are comfortable with that).



Building a Site with Harp

If you're interested in learning how to build a basic static site using the Harp engine, please refer back to Chapter 2 of this book.

Many More...

Obviously, we cannot cover every available option here. The important thing to take away from this is that these tools do exist, though, for the most part, they are built and maintained by user communities of whichever static site generator you choose. Assuming one exists that suits your needs, it's certainly worth trying as it has the potential to save hours of work doing manual migration, and makes the move to a static site that much more compelling.

APPENDIX A

Appendix Title

This Is an A-Head

An appendix is generally used for extra material that supplements your main book content.

Index